

A Novel Approach to Detecting Covert DNS Tunnels Using Throughput Estimation

by

Michael Himbeault

Context:

Motivation

- DNS tunnels uses:
 - Malware and botnet command-and-control channels
 - In/ex-filtration channels for transporting data through corporate security layers
- The ability to monitor the existence of these channels is important when securing a network.
 - Currently quite difficult

Context:

Goals and Objectives

- Be able to identify DNS tunnels that do not violate DNS RFCs or specifications in near real time with high accuracy on commodity hardware.
 - Low false-positive rate

Background:

Domain Name System (DNS)

- Translates text references to other forms of records.
For example:
 - IP or IPv6 address (A or AAAA)
 - Another domain name (CNAME)
 - IP address to name (PTR)
 - Name to bulk data (TXT)
- Provides the human-interaction layer for the Internet
- Offers a great deal of flexibility for deploying automated services over existing infrastructure
 - For example: spam, malware, and address blacklists.

Background:

Covert Channels

- Utilize standard means of transportation in non-standard ways
 - Often transporting unintended data types over existing protocols
 - Occasionally involves new custom protocols built on existing ones
- Intention is rarely benign, often circumventing existing security layers

Background:

Covert Channels

- May or may not modify the standard protocols in ways that are conforming to specifications.
- May sacrifice 'common' features such as bidirectionality
- Examples:
 - IP timing channels
 - May use third party services such as Twitter, Facebook, or image hosting providers
 - Encoding information in JPEG headers
 - DNS tunnels

Background:

Entropy

- Rough measure of the amount of information contained in a collection, C :

$$C = \{c_1, \dots, c_n\}, P(c_i) = p_i$$

$$H(C) = -\sum_{i=1}^n p_i \log p_i$$

Background: DNS Tunnels

- Existing implementations are commonplace
 - Iodine
 - OzymanDNS
 - Dns2tcp
 - DNScat
 - DeNiSe
 - PSUDP
- A custom implementation was built to simulate a next-generation tunnel

Background:

DNS Tunnels

- Raw DNS tunnels
 - Utilize UDP/TCP port 53 for transmitting arbitrary data without respect for DNS protocol specifications
 - Not difficult to block
- Conforming DNS tunnels
 - Makes use of DNS packets that do not violate the protocol specifications to transmit arbitrary data
 - Can be very difficult to identify and block
 - The focus of this work

Custom DNS Tunnel Application

- Prototype proof-of-concept implementation
- Implements encoding to match character frequencies to circumvent detections approaches that tune to that statistic.
- Limited to client-to-server transfer only

State of the Art: DNS Tunnel Detection

- Fall into several categories
 - Signature based
 - Domain hash/blacklist
 - Flow data based
 - Character frequency analysis on queries
 - Behaviour of DNS queries on a per-domain basis
- The proposed approach falls into the last category

Proposed Approach:

Assumptions

- DNS tunnels move more data than benign traffic
- Attempt to detect this increase in data transmission volume

Proposed Approach: Theory

- Collect DNS queries into temporal buckets
 - Ten-second windows were used in the analysis
- Further group queries by top-level domain (TLD)
 - google.com
 - cbc.ca
 - Etc...
- For each TLD (in the current window), compute a measure of how much data was transmitted

Proposed Approach:

Measuring Data Volume

- Since common domains may appear more than uncommon, simple character count is insufficient
 - Modulo caching effects as described in section 5.1.3
- Average character count is similarly uninformative
 - Tunnels can use any length queries, including modelling a length distribution of legitimate traffic

Proposed Approach:

Measuring Data Volume

- DNS tunnels can be expected to have very few queries that appear more than once
 - Since they are transmitting arbitrary data
- Benign domains can have many queries that appear a great number of times
 - Such as www.google.com

Proposed Approach: Domain Length-Weighted Entropy (DLWE)

- Consider the collection of queries to a TLD in an interval
- Treat each query as a symbol, and compute the entropy of the collection.
- Multiply the result by the average query length for the TLD in the interval.
- Expect large values for tunnel domains, small values for benign domains.

Evaluation:

Literature Candidates and Test Data

- Proposed approach was tested against candidates from the literature
 - *N*-gram detection proposed by Born
 - *Gzip* compression detection proposed by Paxson
 - Naïve counting of characters
- All approaches were implemented on a common Python framework
 - Analysis was done on approximately one billion UDP port 53 packets from a live ISP network as well as intentionally generated tunnel traffic

Evaluation Criteria

- Packet processing performance
 - A relative comparison, due to lack of optimized implementations no absolute target is chosen
- False positive rate
 - A relative ranking is employed
 - Intuitively, a false-positive rate of 1% will result in up to fifty alerts per second during average daytime traffic of the captured sample.


Sample Data

- May contain malicious traffic
 - In particular, DNS tunnels which will affect the false-positive rates of the detection methods.

Sample Data: Packets per second

UDP Port 53 Throughput Over Time – Packets

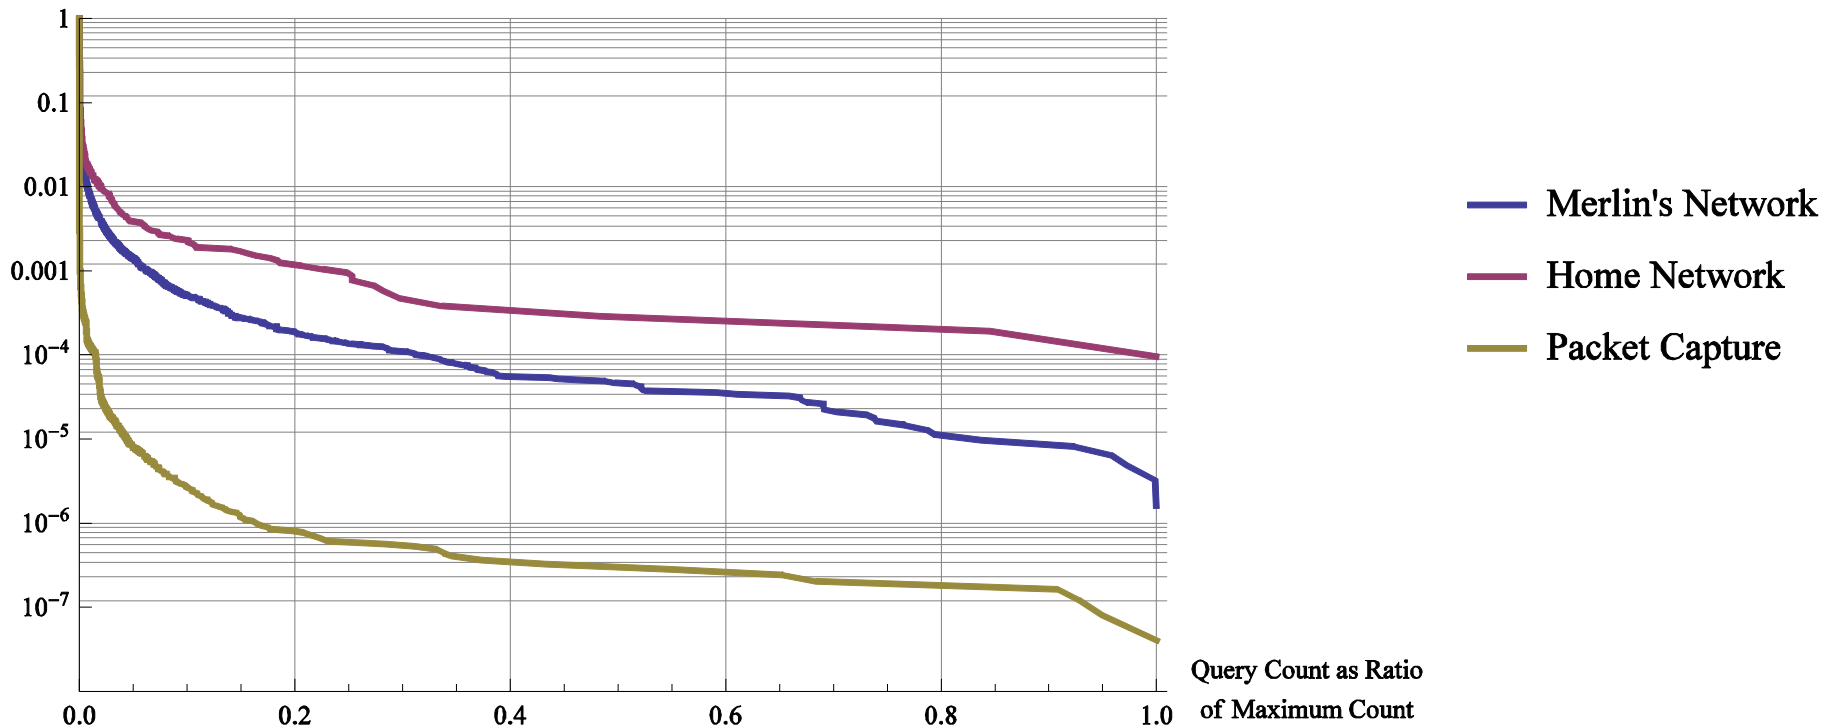
Packets per Second



Confounding Factor: Effect of DNS Caching

Effect of DNS Caching
on Query Repetition Counts


Percent of Samples With
Count Greater than x



Processing Performance

Performance of Analysis Method on Real World Data
Processing Speed by Python Interpreter (Cython, PyPy)


Processing Rate
(packets / second)




Processing Performance

Performance of Analysis Method on Real World Data
Processing Speed by Python Interpreter (Cython, PyPy)

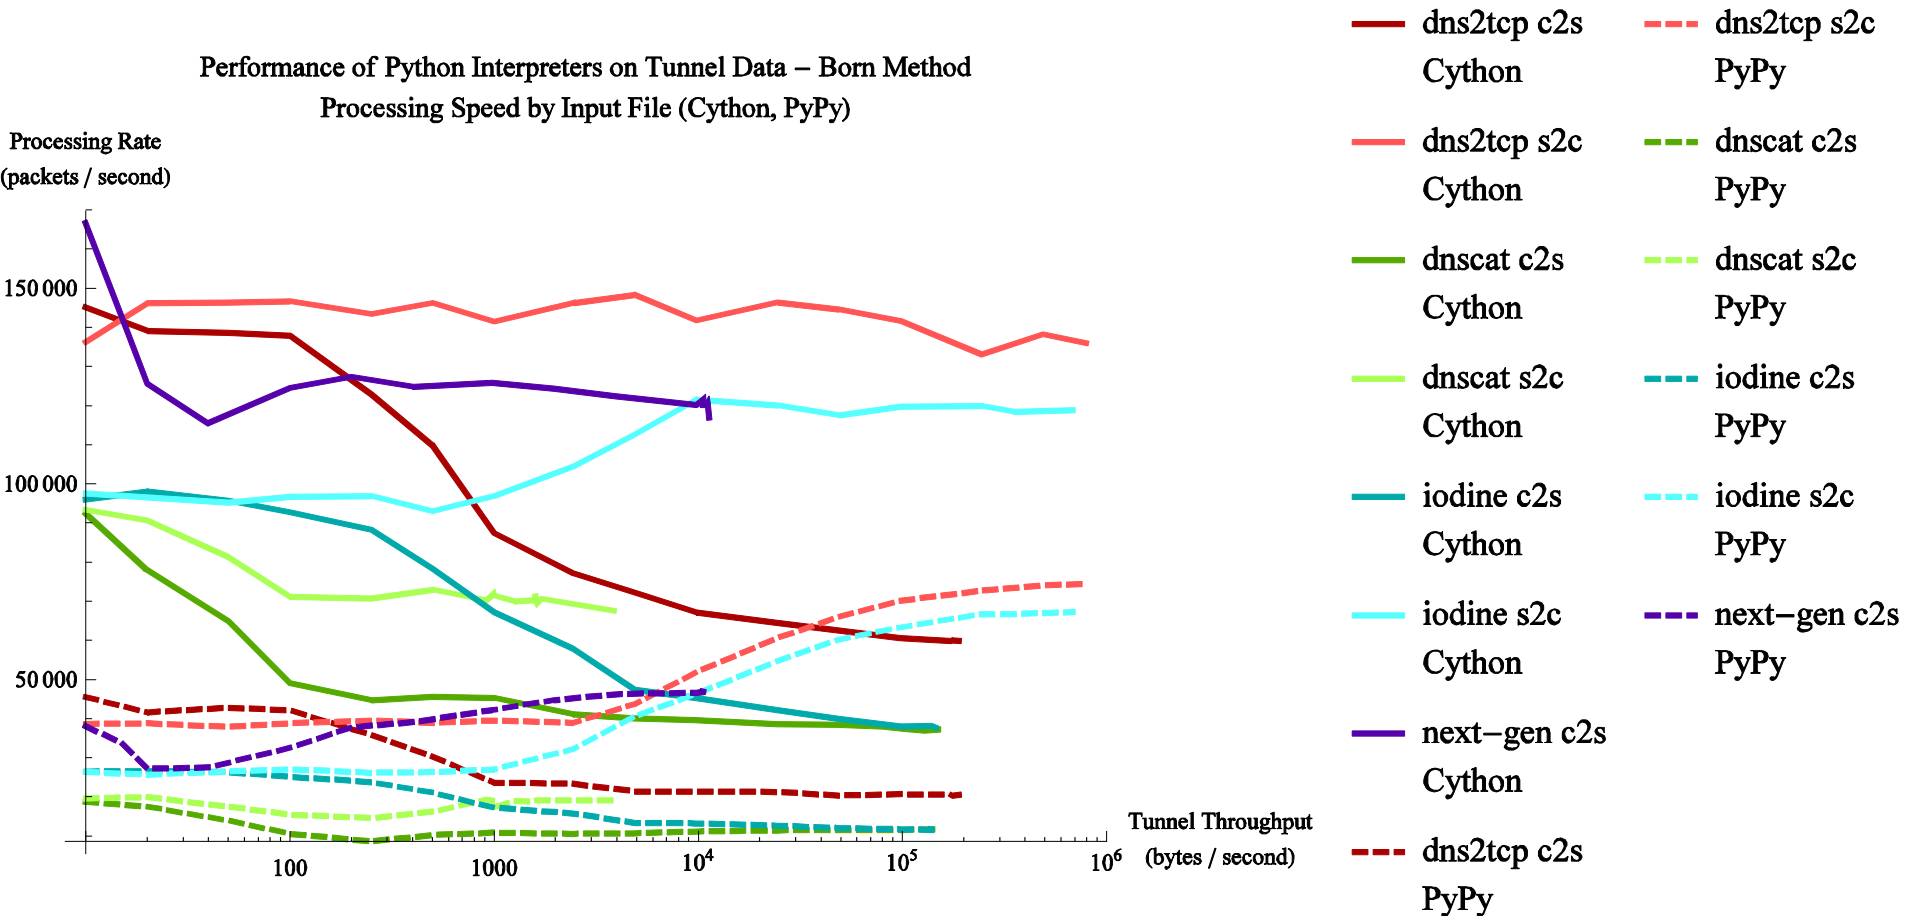
Processing Rate
(packets / second)




Processing Performance: Naïve Method



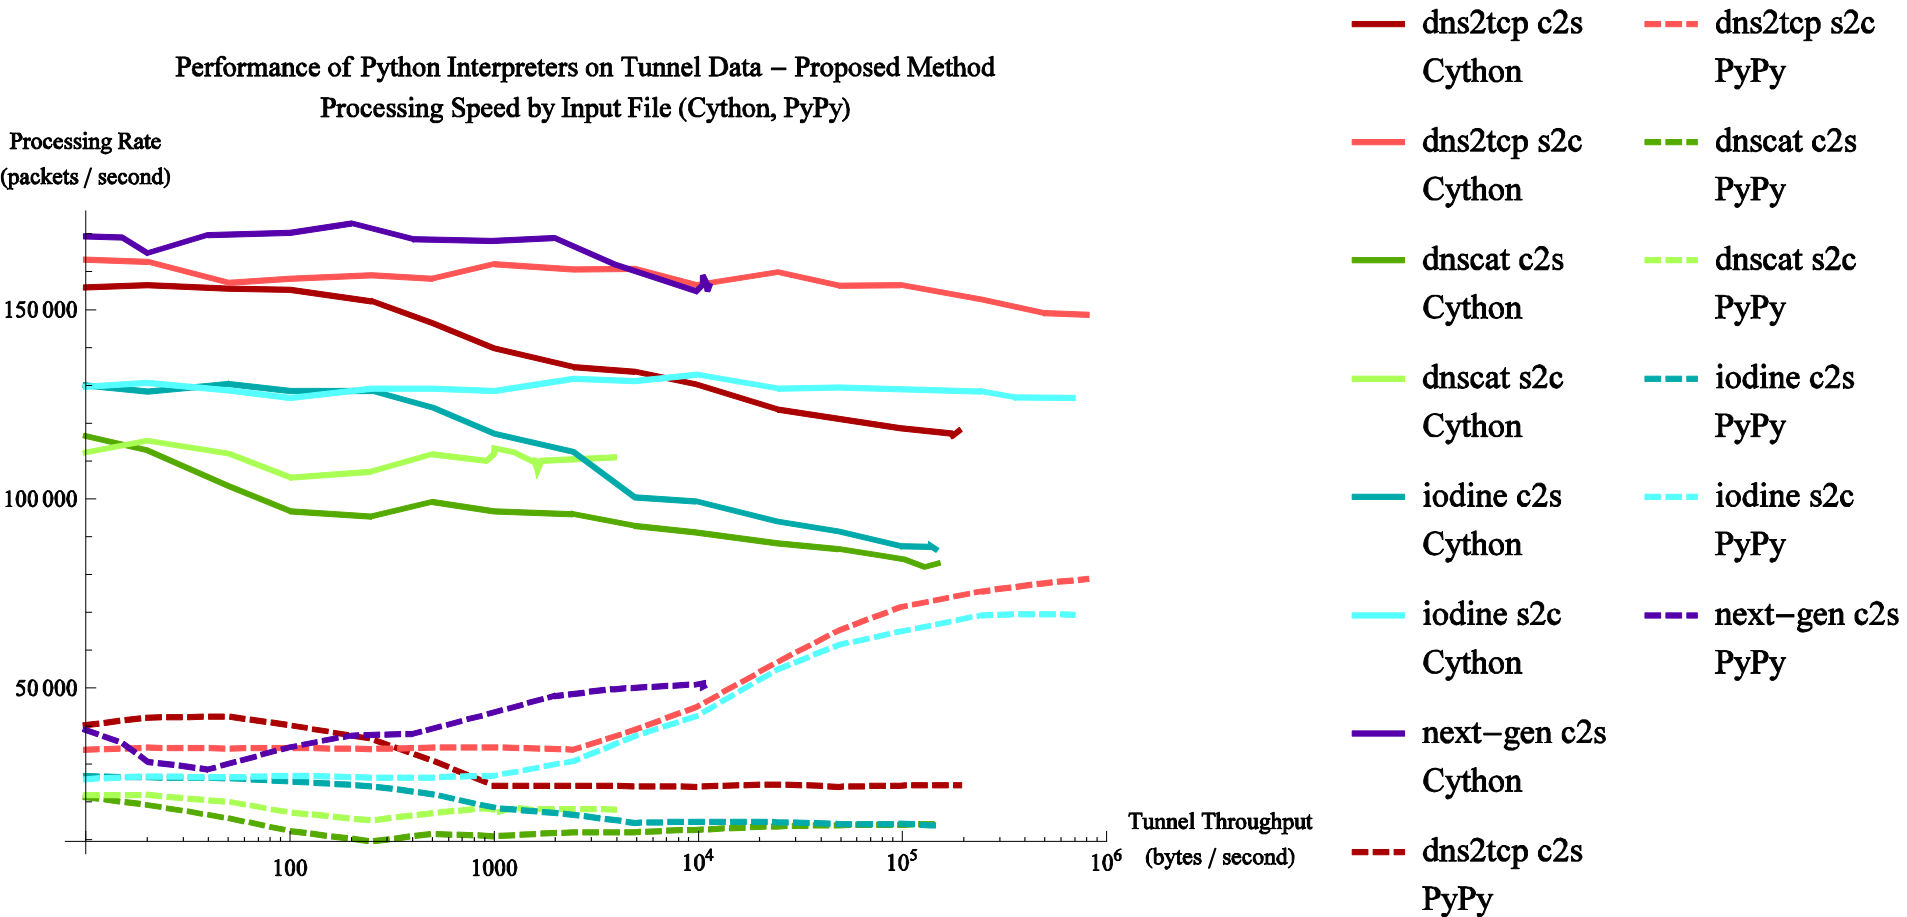
Processing Performance : Born Method



Processing Performance : Paxson Method



Processing Performance : Proposed Method




Processing Performance:

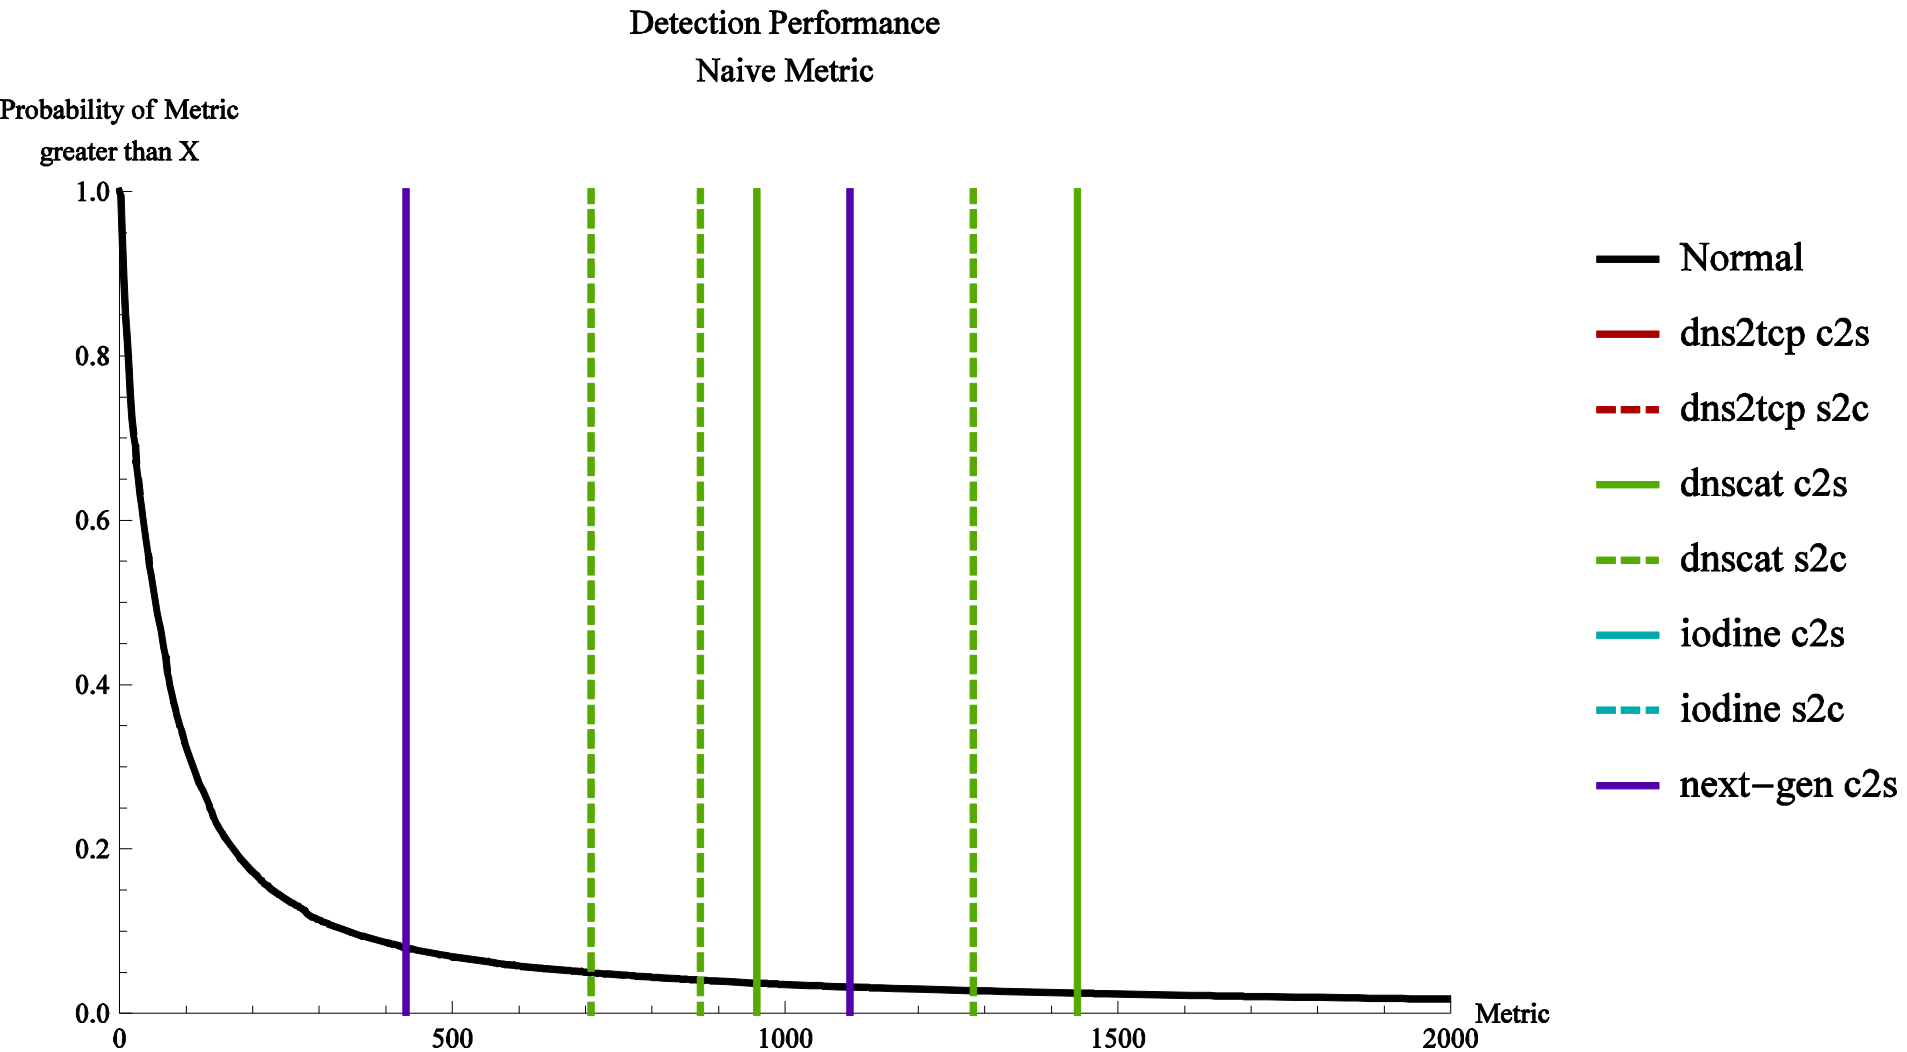
Conclusions

- The naïve and proposed methods far outperform the other methods
- As throughput increases, both Paxson and Born approaches suffer severe degradation in performance.

Trend of False Positive Rate – Naive Metric




Detection Performance : Naïve Method



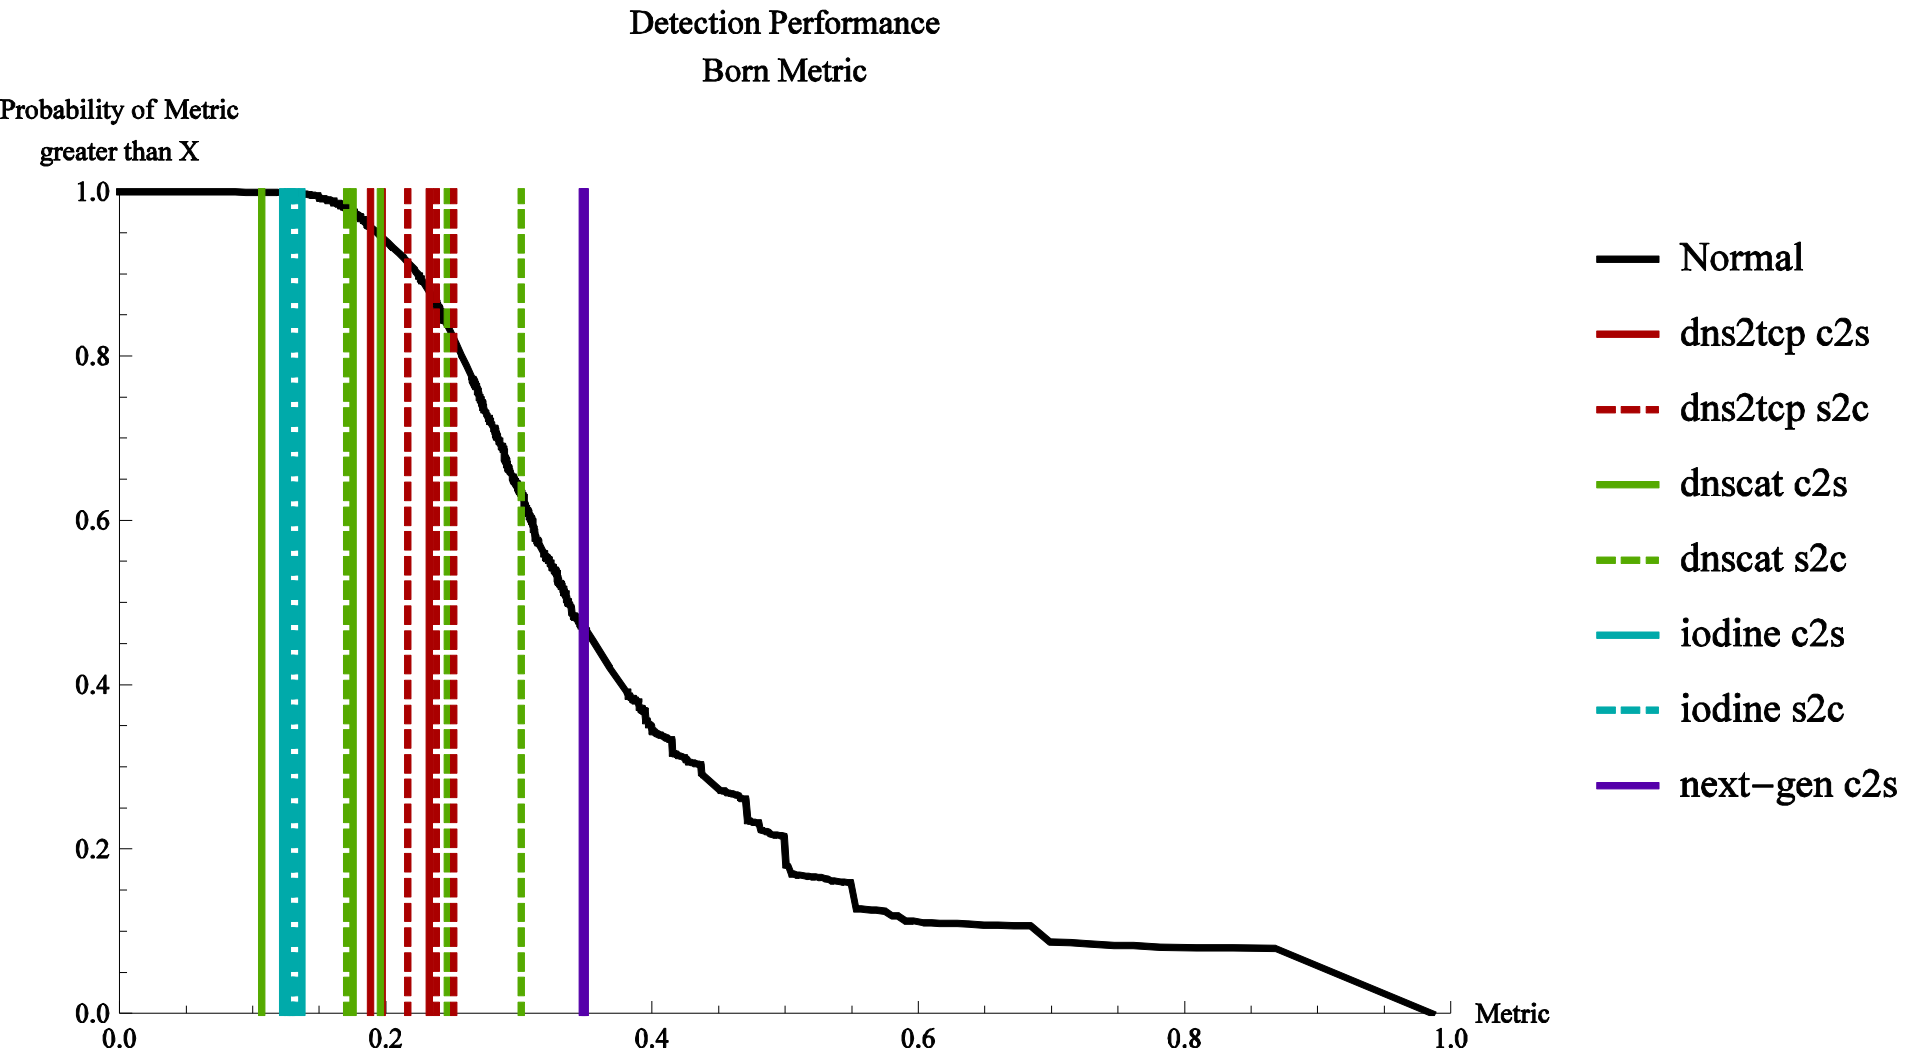
Detection Performance : Born Method

Trend of False Positive Rate – Born Metric

Percent of Normal Traffic
With Metric Less
Than Tunnel




Detection Performance : Born Method



Trend of False Positive Rate – Paxson Metric


1



Detection Performance : Paxson Method

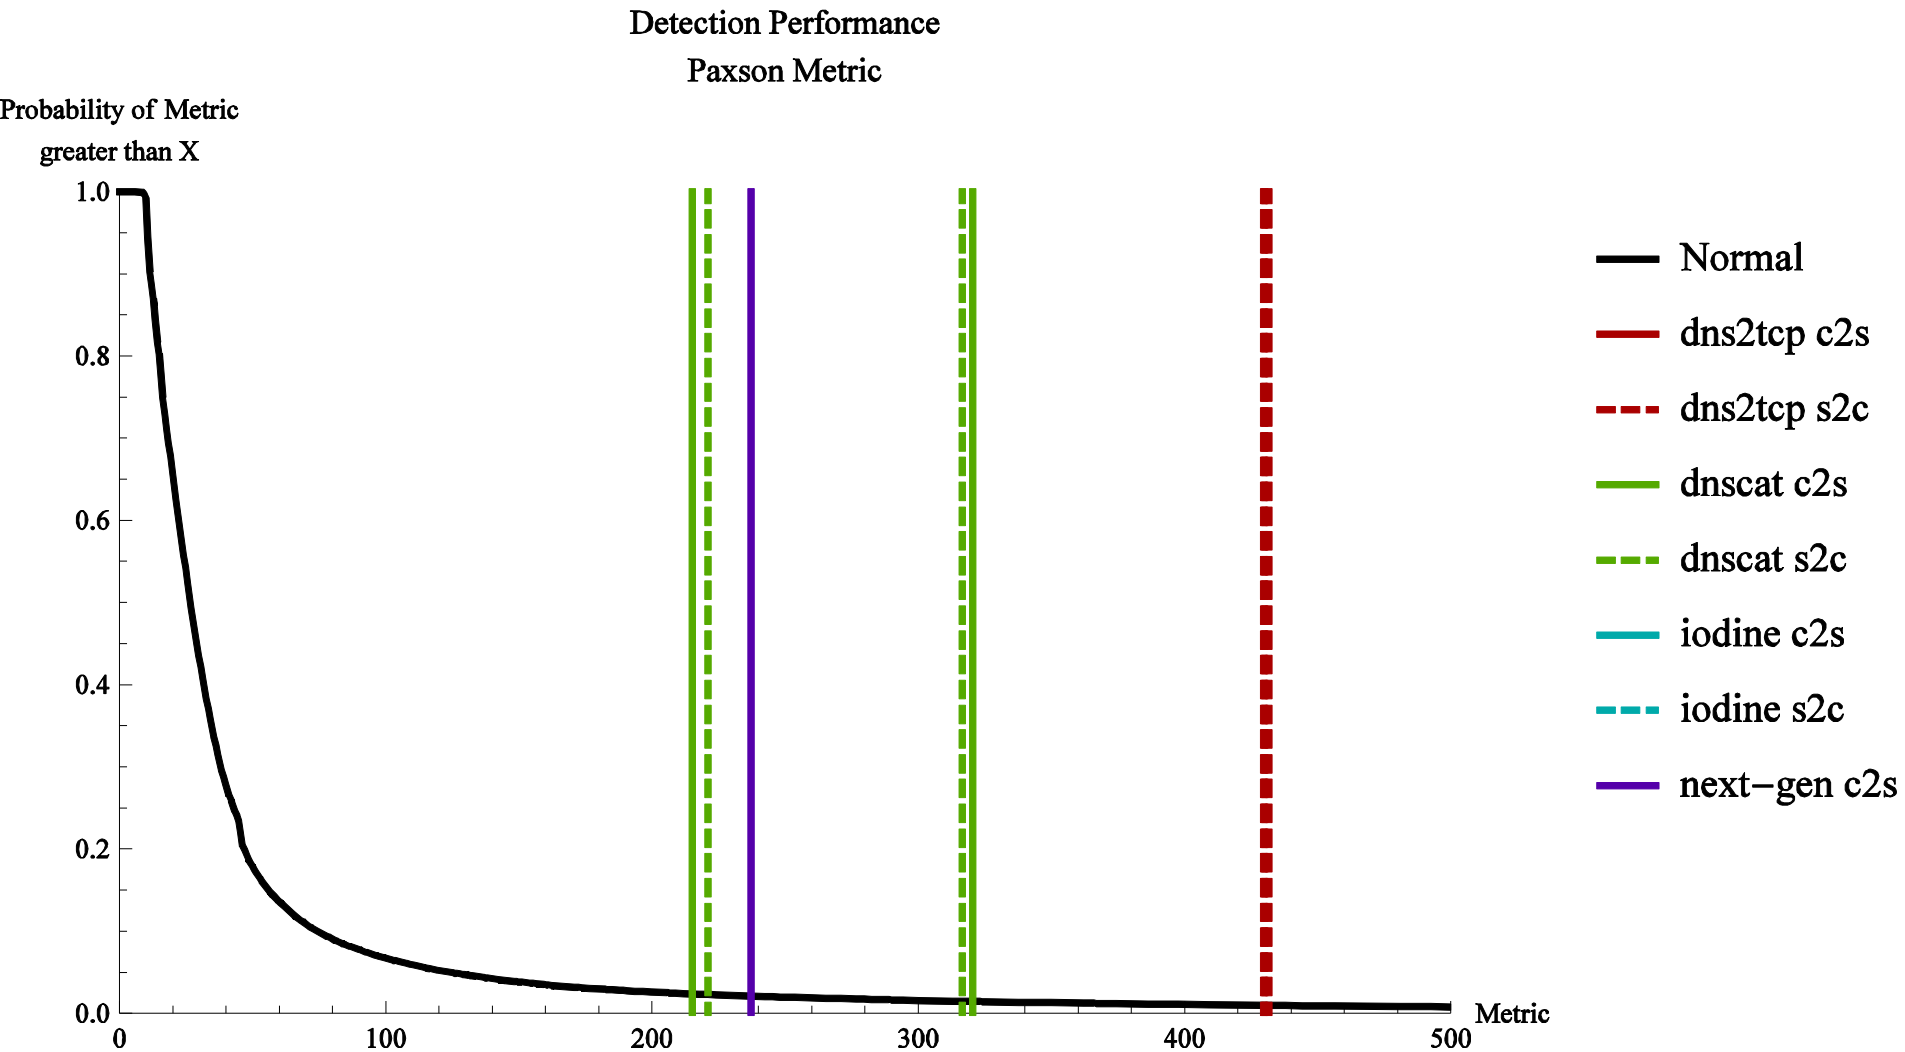
Detection Performance
Paxson Metric

Probability of Metric
greater than X



- Normal
- dns2tcp c2s
- dns2tcp s2c
- dnscat c2s
- dnscat s2c
- iodine c2s
- iodine s2c
- next-gen c2s


Detection Performance : Paxson Method



Detection Performance : Proposed Method

Trend of False Positive Rate – Proposed Metric


Percent of Normal Traffic
With Metric Greater
Than Tunnel



Detection Performance : ProposedMethod

Detection Performance
Proposed Metric

Probability of Metric
greater than X




- Normal
- dns2tcp c2s
- dns2tcp s2c
- dnscat c2s
- dnscat s2c
- iodine c2s
- iodine s2c
- next-gen c2s

Detection Performance : Proposed Method


Detection Performance
Proposed Metric

Probability of Metric
greater than X




- Normal
- dns2tcp c2s
- dns2tcp s2c
- dnscat c2s
- dnscat s2c
- iodine c2s
- iodine s2c
- next-gen c2s

Detection Performance: Comparison of False Positive Rates




Detection Performance: Comparison of False Positive Rates

Certainty of Tunnel Detection



Detection Performance: Comparison of False Positive Rates

Certainty of Tunnel Detection



Detection Performance:

Conclusions

- The proposed approach achieves categorically lower false positive rates than all other approaches.
- The prototype next-gen tunnel is the most difficult tunnel to detect by far.
 - Born's approach has a false-positive rate little better than random chance.

Conclusions

- Proposed method:
 - Achieves the best detection performance, and nearly the best processing performance.
 - Represents a notable and novel contribution to the field.
 - *Is already implemented in high-performance C/C++, making deployment possible.*

Potential Future Work

- Test on more strictly curated data sets to remove confounding factors.
- Identify ways of improving false positive rate
 - Potentially with a more tailored metric
 - Potentially with more temporal knowledge and correlation

Thank you

- Questions

