

# A Novel Approach to Detecting Covert DNS Tunnels Using Throughput Estimation

by

Michael Himbeault

# Context:

## Motivation

- DNS tunnels uses:
  - Malware and botnet command-and-control channels
  - In/ex-filtration channels for transporting data through corporate security layers
- The ability to monitor the existence of these channels is important when securing a network.
  - Currently quite difficult

# Context:

## Goals and Objectives

- Be able to identify DNS tunnels that do not violate DNS RFCs or specifications in near real time with high accuracy.
  - Low false-positive rate

# Background:

## Domain Name System (DNS)

- Translates text references to other forms of records.  
For example:
  - IP or IPv6 address (A or AAAA)
  - Another domain name (CNAME)
  - IP address to name (PTR)
  - Name to bulk data (TXT)
- Provides the human-interaction layer for the Internet
- Offers a great deal of flexibility for deploying automated services over existing infrastructure
  - For example: spam, malware, and address blacklists.

# Background:

## Covert Channels

- Utilize standard means of transportation in non-standard ways
  - Often transporting unintended data types over existing protocols
  - Occasionally involves new custom protocols built on existing ones
- Intention is rarely benign, often circumventing existing security layers

# Background:

## Covert Channels

- May or may not modify the standard protocols in ways that are conforming to specifications.
- May sacrifice 'common' features such as bidirectionality
- Examples:
  - IP timing channels
  - May use third party services such as Twitter, Facebook, or image hosting providers
    - Encoding information in JPEG headers
  - DNS tunnels

# Background:

## Entropy

- Rough measure of the amount of information contained in a collection,  $C$ :

$$C = \{c_1, \dots, c_n\}, P(c_i) = p_i$$

$$H(C) = -\sum_{i=1}^n p_i \log p_i$$

# Background: DNS Tunnels

- Existing implementations are commonplace
  - Iodine
  - OzymanDNS
  - Dns2tcp
  - DNScat
  - DeNiSe
  - PSUDP
- A custom implementation was built to simulate a next-generation tunnel



# Background: DNS Tunnels

- Raw DNS tunnels
  - Utilize UDP/TCP port 53 for transmitting arbitrary data without respect for DNS protocol specifications
  - Not difficult to block
- Conforming DNS tunnels
  - Makes use of DNS packets that do not violate the protocol specifications to transmit arbitrary data
  - Can be very difficult to identify and block
  - The focus of this work

# Custom DNS Tunnel Application

- Prototype proof-of-concept implementation
- Implements encoding to match character frequencies to circumvent detections approaches that tune to that statistic.
- Limited to client-to-server transfer only

# State of the Art: DNS Tunnel Detection

- Fall into several categories
  - Signature based
  - Domain hash/blacklist
  - Flow data based
  - Character frequency analysis on queries
  - Behaviour of DNS queries on a per-domain basis
- The proposed approach falls into the last category

# Proposed Approach:

## Assumptions

- DNS tunnels move more data than benign traffic
- Attempt to detect this increase in data transmission volume

# Proposed Approach: Theory

- Collect DNS queries into temporal buckets
  - Ten-second windows were used in the analysis
- Further group queries by top-level domain (TLD)
  - google.com
  - cbc.ca
  - Etc...
- For each TLD (in the current window), compute a measure of how much data was transmitted

# Proposed Approach:

## Measuring Data Volume

- Since common domains may appear more than uncommon, simple character count is insufficient
  - Modulo caching effects as described in section 5.1.3
- Average character count is similarly uninformative
  - Tunnels can use any length queries, including modelling a length distribution of legitimate traffic

# Proposed Approach:

## Measuring Data Volume

- DNS tunnels can be expected to have very few queries that appear more than once
  - Since they are transmitting arbitrary data
- Benign domains can have many queries that appear a great number of times
  - Such as [www.google.com](http://www.google.com)

## Proposed Approach: Domain Length-Weighted Entropy (DLWE)

- Consider the collection of queries to a TLD in an interval
- Treat each query as a symbol, and compute the entropy of the collection.
- Multiply the result by the average query length for the TLD in the interval.
- Expect large values for tunnel domains, small values for benign domains.



# Evaluation:

## Literature Candidates and Test Data

- Proposed approach was tested against candidates from the literature
  - *N*-gram detection proposed by Born
  - *Gzip* compression detection proposed by Paxson
  - Naïve counting of characters
- All approaches were implemented on a common Python framework
  - Analysis was done on approximately one billion UDP port 53 packets from a live ISP network as well as intentionally generated tunnel traffic

# Evaluation Criteria

- Packet processing performance
  - A relative comparison, due to lack of optimized implementations no absolute target is chosen
- False positive rate
  - A relative ranking is employed
  - Intuitively, a false-positive rate of 1% will result in up to fifty alerts per second during average daytime traffic of the captured sample.


# Sample Data

- May contain malicious traffic
  - In particular, DNS tunnels which will affect the false-positive rates of the detection methods.

# Sample Data: Packets per second

UDP Port 53 Throughput Over Time – Packets

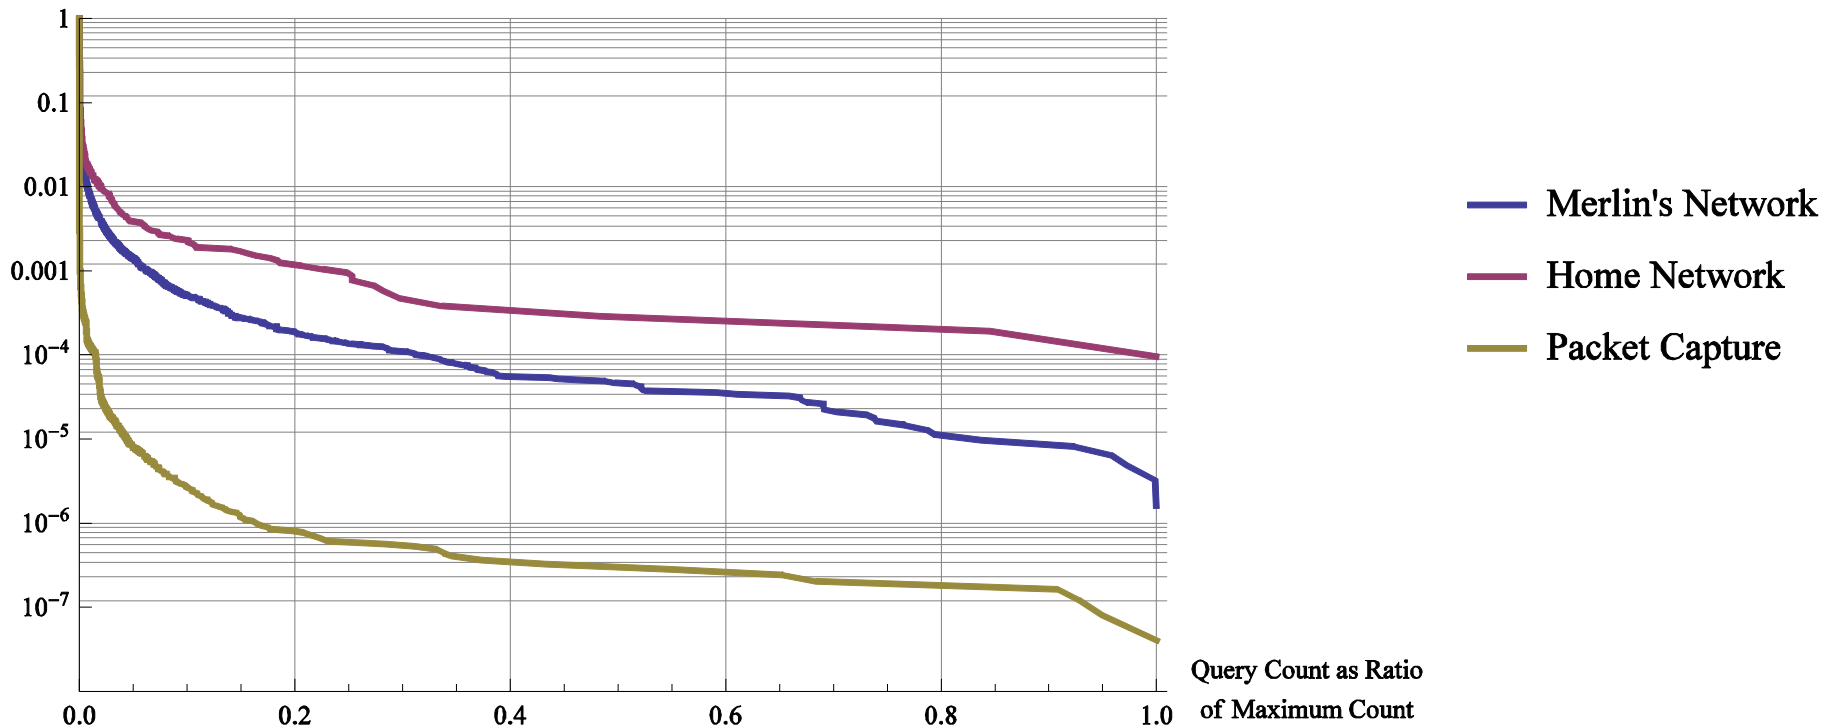
Packets per Second



# Confounding Factor: Effect of DNS Caching

Effect of DNS Caching  
on Query Repetition Counts


Percent of Samples With  
Count Greater than x



# Processing Performance

Performance of Analysis Method on Real World Data  
Processing Speed by Python Interpreter (Cython, PyPy)


Processing Rate  
(packets / second)




# Processing Performance

Performance of Analysis Method on Real World Data  
Processing Speed by Python Interpreter (Cython, PyPy)

Processing Rate  
(packets / second)




# Processing Performance: Naïve Method

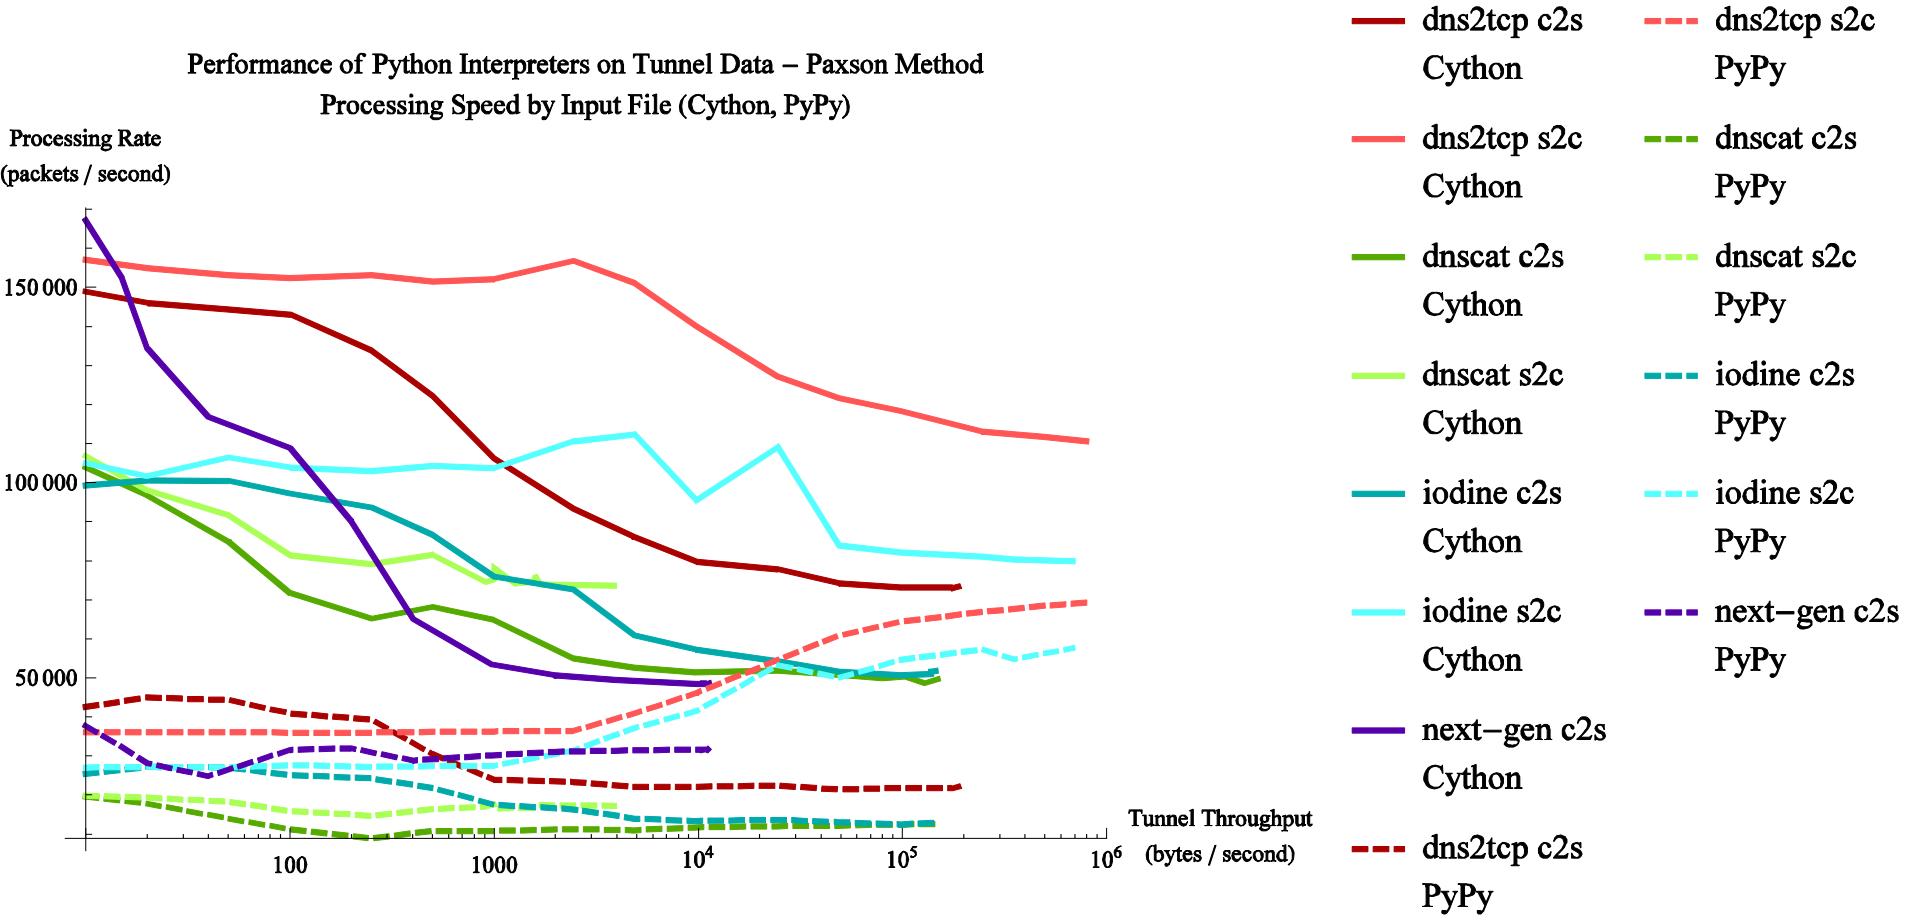




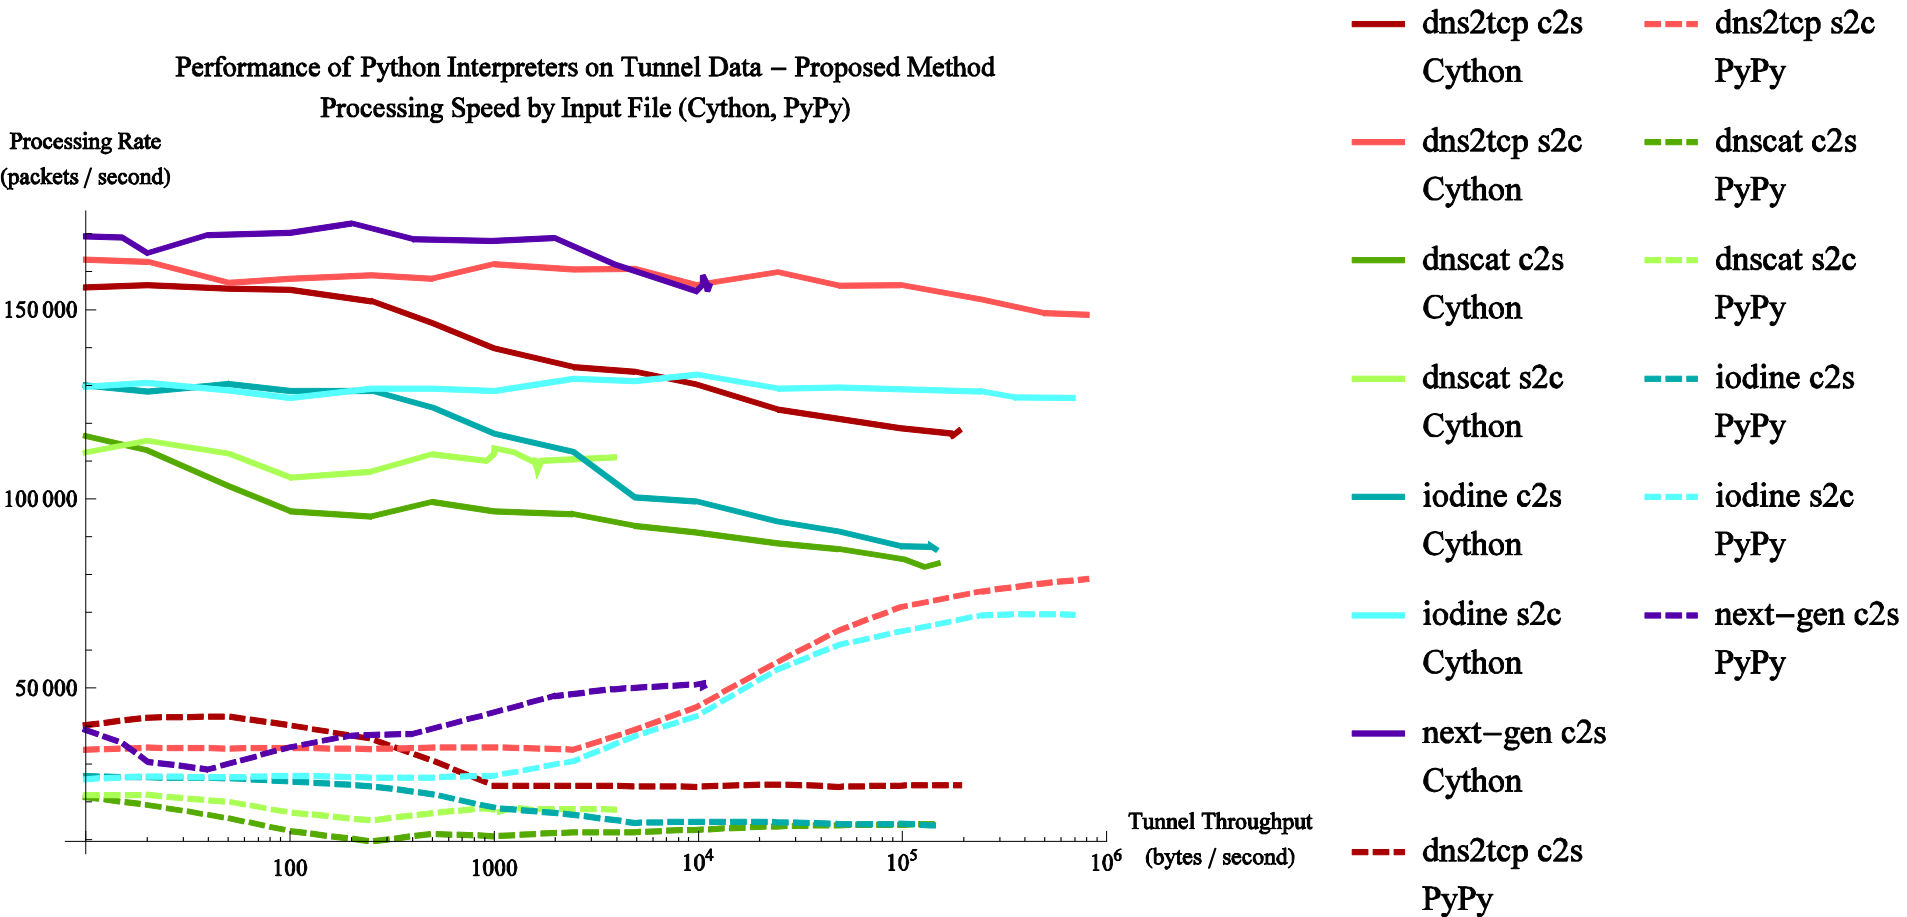
# Processing Performance : Born Method



# Processing Performance : Paxson Method



# Processing Performance : Proposed Method




# Processing Performance: Conclusions

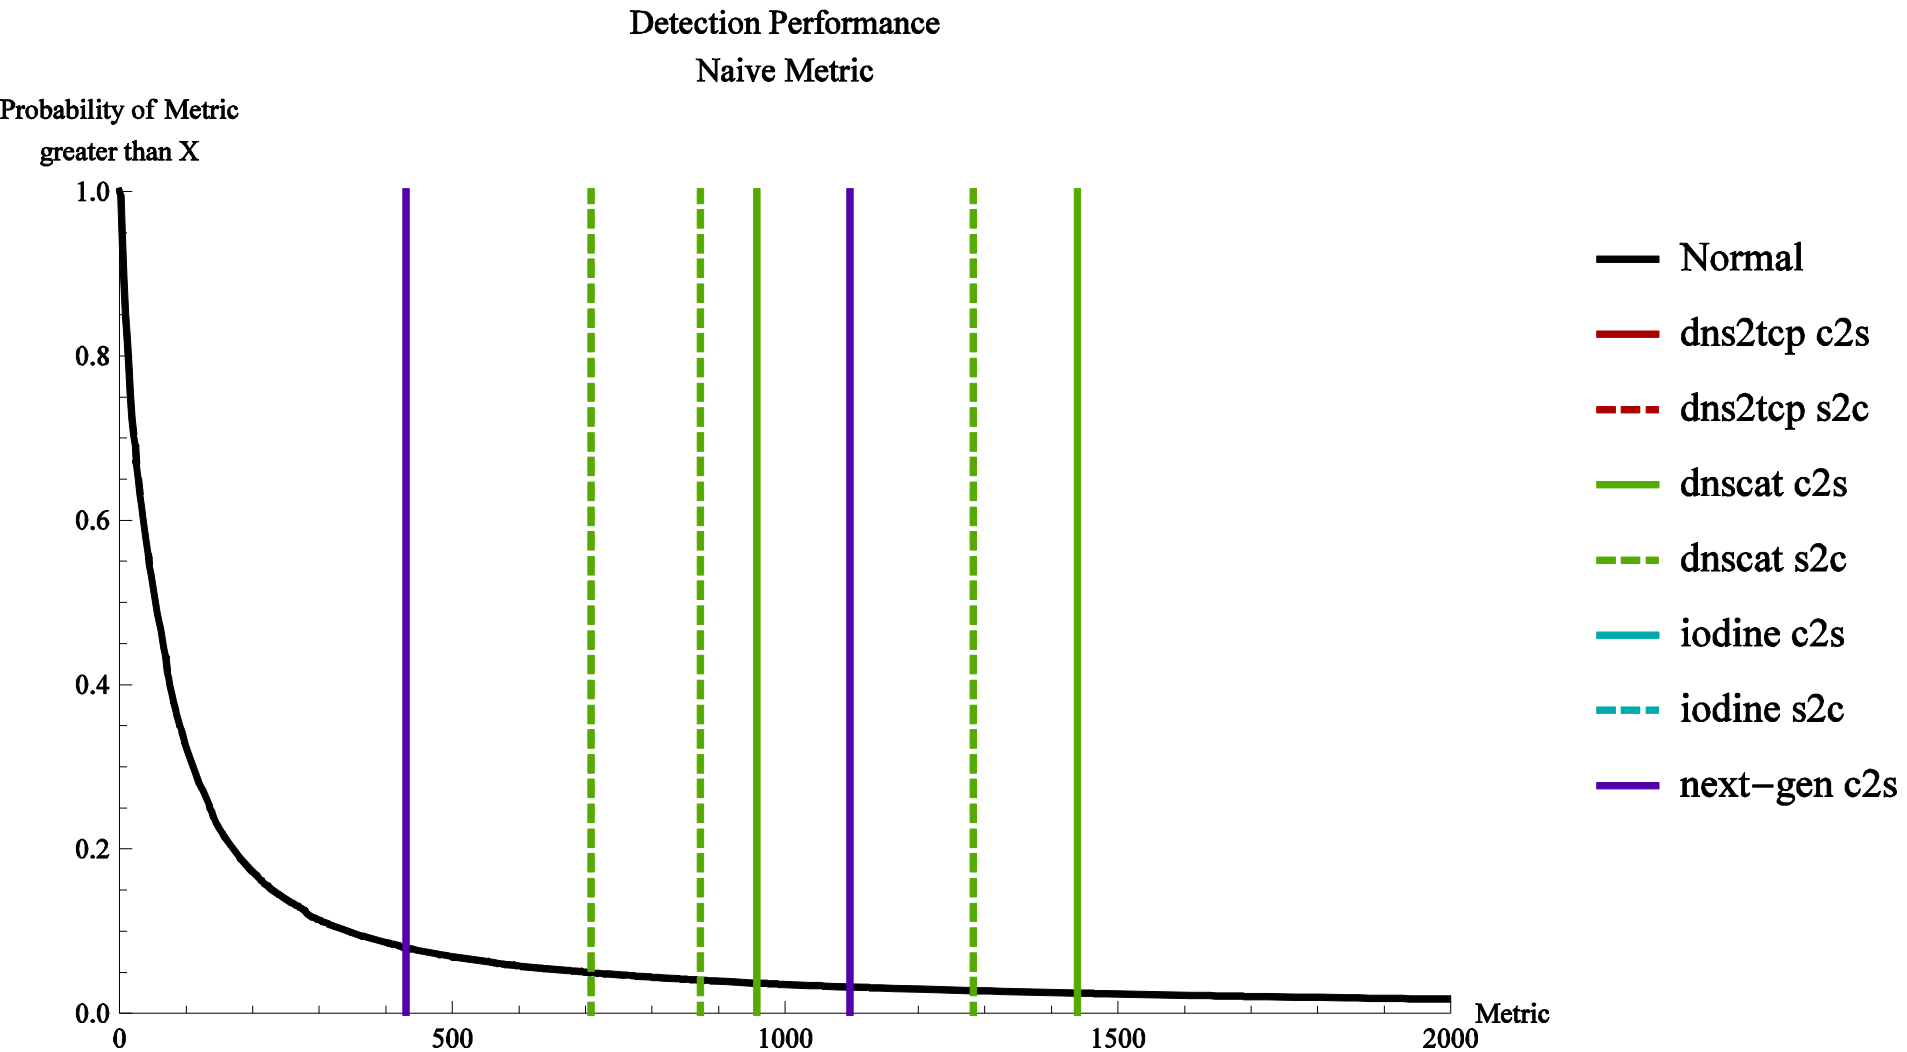
- The naïve and proposed methods far outperform the other methods
- As throughput increases, both Paxson and Born approaches suffer severe degradation in performance.

### Trend of False Positive Rate – Naive Metric

**1**




# Detection Performance : Naïve Method



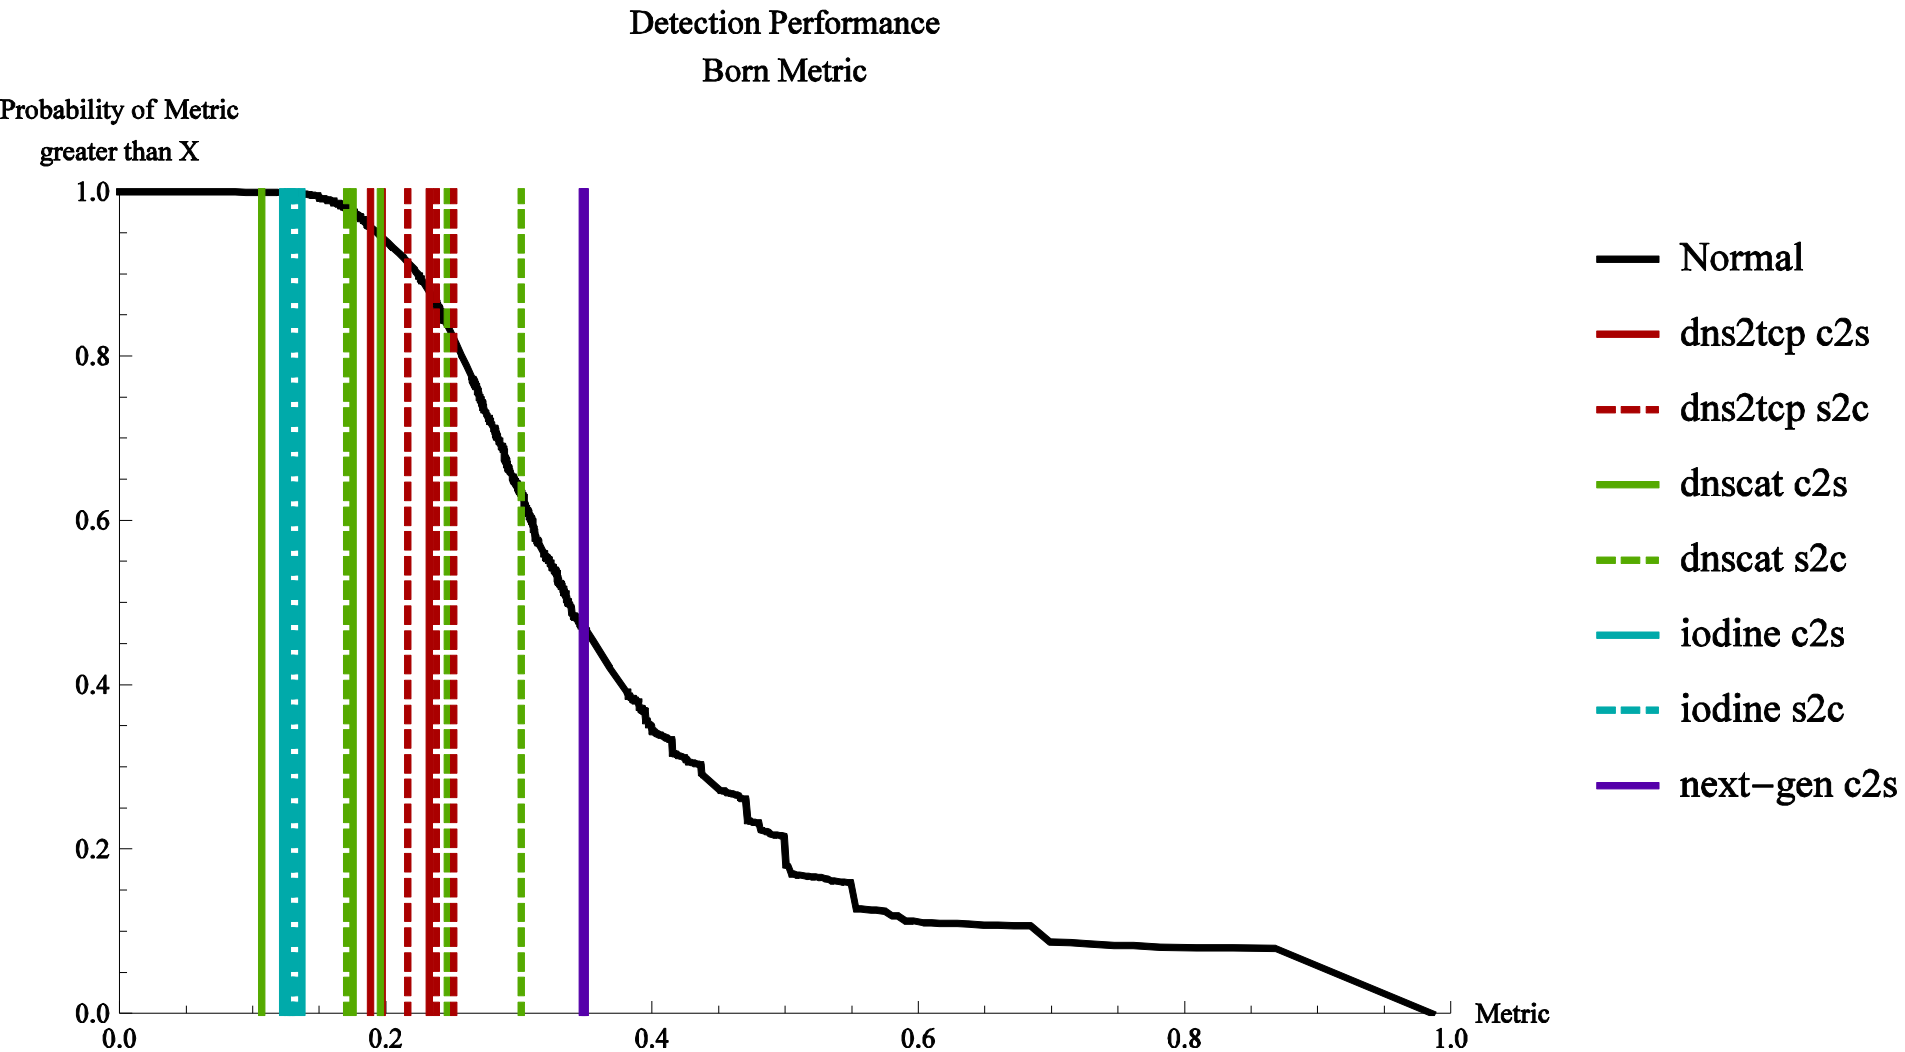
# Detection Performance : Born Method

Trend of False Positive Rate – Born Metric

Percent of Normal Traffic  
With Metric Less  
Than Tunnel




# Detection Performance : Born Method





### Trend of False Positive Rate – Paxson Metric


**1**



# Detection Performance : Paxson Method


Detection Performance  
Paxson Metric

Probability of Metric  
greater than X



- Normal
- dns2tcp c2s
- dns2tcp s2c
- dnscat c2s
- dnscat s2c
- iodine c2s
- iodine s2c
- next-gen c2s


# Detection Performance : Paxson Method



# Detection Performance : Proposed Method

Trend of False Positive Rate – Proposed Metric


Percent of Normal Traffic  
With Metric Greater  
Than Tunnel



# Detection Performance : ProposedMethod

Detection Performance  
Proposed Metric

Probability of Metric  
greater than X




- Normal
- dns2tcp c2s
- dns2tcp s2c
- dnscat c2s
- dnscat s2c
- iodine c2s
- iodine s2c
- next-gen c2s

# Detection Performance : Proposed Method


Detection Performance  
Proposed Metric

Probability of Metric  
greater than X

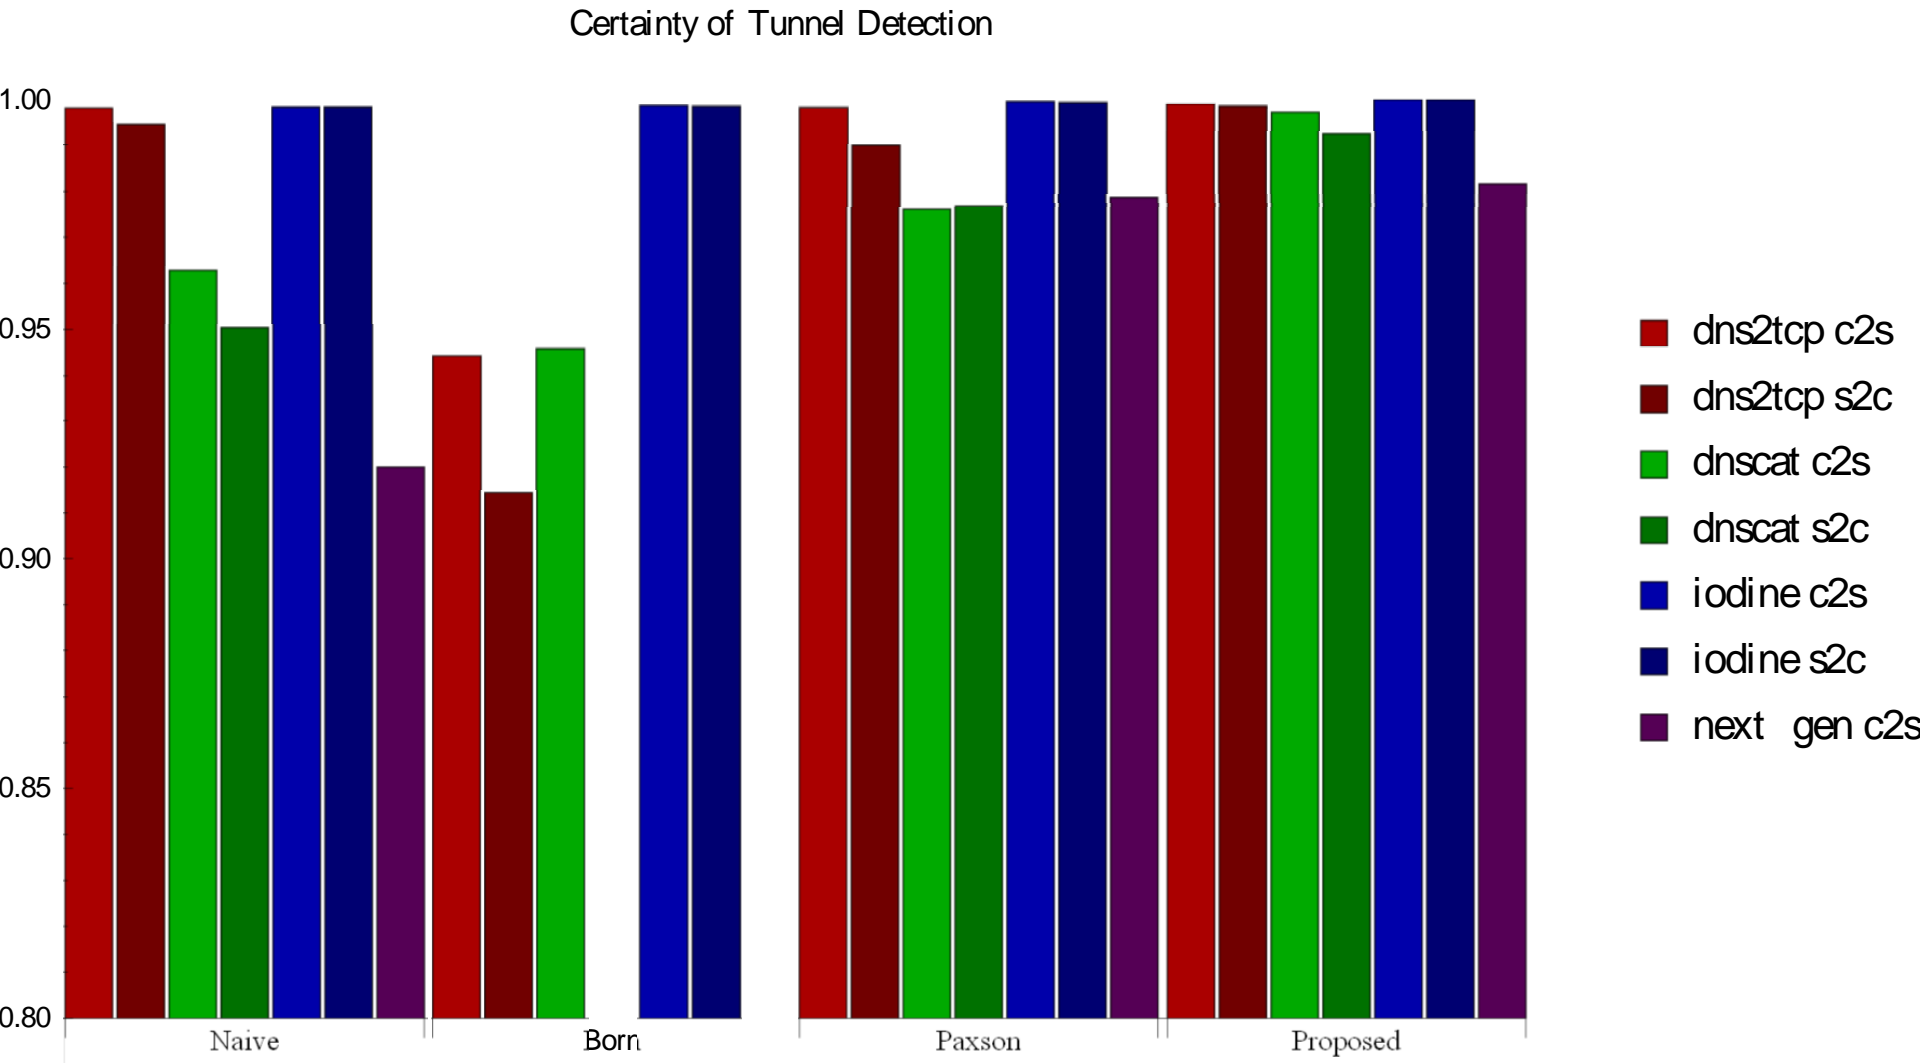


- Normal
- dns2tcp c2s
- dns2tcp s2c
- dnscat c2s
- dnscat s2c
- iodine c2s
- iodine s2c
- next-gen c2s

# Detection Performance: Comparison of False Positive Rates




# Detection Performance: Comparison of False Positive Rates





# Detection Performance: Comparison of False Positive Rates

Certainty of Tunnel Detection



# Detection Performance:

## Conclusions

- The proposed approach achieves categorically lower false positive rates than all other approaches.
- The prototype next-gen tunnel is the most difficult tunnel to detect by far.
  - Born's approach has a false-positive rate little better than random chance.

# Conclusions

- Proposed method:
  - Achieves the best detection performance, and nearly the best processing performance.
  - Represents a notable and novel contribution to the field.
  - *Is already implemented in high-performance C/C++, making deployment possible.*

# Potential Future Work

- Test on more strictly curated data sets to remove confounding factors.
- Identify ways of improving false positive rate
  - Potentially with a more tailored metric
  - Potentially with more temporal knowledge and correlation

# Thank you

- Questions

