

A Novel Approach to Detecting Covert DNS Tunnels Using Throughput Estimation

by

Michael Himbeault

A Thesis

Submitted to the Faculty of Graduate Studies
of the University of Manitoba
in partial fulfilment of the requirements
for the degree of

MASTER OF SCIENCE

Department of Electrical and Computer Engineering
University of Manitoba
Winnipeg, Manitoba, Canada

Copyright ©2013 Michael Himbeault

November 19, 2013

Abstract

In a world that relies heavily on data, protection of that data and of the *motion* of that data is of the utmost importance. Covert communication channels attempt to circumvent established methods of control, such as firewalls and proxies, by utilizing non-standard means of getting messages between two endpoints. The Domain Name System (DNS), the system that translates text-based resource names into machine-readable resource records, is a very common and effective platform upon which covert channels can be built. This work proposes, and demonstrates the effectiveness of, a novel technique that estimates data transmission throughput over DNS in order to identify the existence of a DNS tunnel against the background noise of legitimate network traffic. The proposed technique is robust in the face of the obfuscation techniques that are able to hide tunnels from existing detection methods.

Contents

1	Introduction	1
2	Background	5
2.1	Entropy	5
2.2	Domain Name System (DNS)	6
2.3	Covert Channels	8
2.4	DNS Tunnels	9
3	Review of the State of the Art	13
3.1	General Covert Channel or Anomaly Detection Research	13
3.2	Non-DNS Related Research	15
3.3	DNS Covert Channel Research	15
3.4	DNS Tunnel Detection Landscape	22
4	Problem Statement and Evaluation Criteria	26
4.1	Brief Statement	26
4.2	Detailed Problem Description	26
4.2.1	Theoretical Proof of Satisfaction of Item Four for Higher Dimensions	28
4.2.2	Empirical Demonstration for Smaller Dimensions	30
4.3	Solution Evaluation Criteria	34

5	Proposed Detection Method	36
5.1	Theoretical Basis	36
5.1.1	Assumptions	36
5.1.2	Theory	37
5.1.3	The Effect of DNS Caching on Detection Effectiveness	38
5.2	Implementation	41
5.2.1	General	41
5.2.2	C++	42
5.2.3	Python	43
6	Detailed Testing Methodology and Processing Performance	44
6.1	Situational Performance Goals	45
6.1.1	Determining a Baseline	45
6.1.2	Existing Implementation Detection	46
6.1.3	Demonstration of Existing Weaknesses	47
6.2	Analysis Method Implementations	48
6.2.1	Common Scaffolding	48
6.2.2	Naive	49
6.2.3	Born	49
6.2.4	Paxson	50
6.2.5	Proposed Approach	51
6.2.6	Tunneling Application Throughput	51
6.2.7	Detection Method and Python Interpreter Processing Performance .	52
6.2.8	Processing Performance Conclusion	61
7	Tunnel Detection Evaluation	64
7.1	Detection Performance Against Real World Data	69

7.2	Specificity and Ambiguity of Tunnel Classification	76
7.3	Tunnel Detection Performance Conclusion	78
8	Conclusion	80
8.1	Future Work	82
Appendices		83
A	Properties of the Network Capture	84
B	Probabilistic Encoding	87
B.1	Sample Encoding - English Unigrams	88
B.1.1	Technical Details	92
B.2	Custom DNS Tunnel Endpoint Simulator	96

List of Tables

4.1 Alexa Top One-Million Character Distribution	30
7.1 Relative Standard Deviation of Lowest Throughput Tunnel by Detection Method and Tunnel Application	69
A.1 Packet Capture Statistics from <i>capinfos</i>	84
B.1 English Character Frequency	89
B.2 Encoding Sample - Source	89
B.3 Encoding Sample - Output	89
B.4 English Character Bit-String Mapping - Accurate	94
B.5 English Character Bit-String Mapping - Efficient	95

List of Figures

4.1	Growth of Expected and Mean Norms of Sampled Vectors	31
4.2	Relative Error of μ of Best-Fit Normal Distribution of Sampled Expected Value	33
4.3	Trend of σ of Best-Fit Normal Distribution of Sampled Expected Value . .	33
5.1	Effect of DNS Caching on Query Counts	40
6.1	DNS Tunneling Application Throughput Scaling	52
6.2	Performance of Analysis Method and Python Interpreter on Aggregate Tunnel Data	54
6.3	Performance of Analysis Method and Python Interpreter on Real World Data	55
6.4	Performance of Analysis Method and Python Interpreter on Real World Data - Early Time	56
6.5	Performance of Naive Method on Tunnel Data by Python Interpreter . . .	57
6.6	Performance of Born's Method on Tunnel Data by Python Interpreter . .	58
6.7	Performance of Paxson's Method on Tunnel Data by Python Interpreter . .	59
6.8	Performance of the Proposed Method on Tunnel Data by Python Interpreter	60
6.9	Performance Ratio of the Naive Method to the Proposed Method on Tunnel Data by Python Interpreter	62
7.1	Relative Standard Deviation of Metrics - Naive Metric	65

7.2	Relative Standard Deviation of Metrics - Born Metric	66
7.3	Relative Standard Deviation of Metrics - Paxson Metric	66
7.4	Relative Standard Deviation of Metrics - Proposed Metric	67
7.5	Scaling of Detection Metrics by Tunnel Throughput	68
7.6	Trend of False Positive Rate for Tunnels by Throughput - Naive Metric . .	71
7.7	Tunnel Detection Performance - Naive Metric	71
7.8	Trend of False Positive Rate for Tunnels by Throughput - Born's Metric .	72
7.9	Tunnel Detection Performance - Born's Metric	73
7.10	Trend of False Positive Rate for Tunnels by Throughput - Paxson's Metric	74
7.11	Tunnel Detection Performance - Paxson's Metric	74
7.12	Trend of False Positive Rate for Tunnels by Throughput - Proposed Metric	75
7.13	Tunnel Detection Performance - Proposed Metric	76
7.14	Chart of Specificity of Detection by Tunnel Application and Detection Method	77
7.15	Chart of Certainty of Detection by Tunnel Application and Detection Method - 0.80 to 1.00	78
7.16	Chart of Certainty of Detection by Tunnel Application and Detection Method - 0.95 to 1.00	79
A.1	Throughput in Bytes per Second of Packet Capture	85
A.2	Throughput in Packets per Second of Packet Capture	86
B.1	Encoding Sample - Source (data matrix)	90
B.2	Plot of Character Frequencies	91

Chapter 1

Introduction

Control of the data that moves between hosts on a computer network is vitally important in order to be able to enforce security policies and protect sensitive information. To this end systems such as firewalls, proxies, and content filters are put into place in order to monitor and control network traffic. Covert channels are utilized to circumvent these control mechanisms for purposes that range from benign[31] to malicious. This work focuses specifically on DNS tunnels, however other types of covert channels do exist.

Detecting a DNS tunnel effectively on a busy network link becomes an exercise in discrimination. Since there is such a wide variety of network traffic that is generated on a busy link, there is generally no simple definition of *normal* for a particular class of traffic, including DNS. This tends to either rule out, or decrease the viability of, algorithms that depend on finding a definition of normal and alerting based on deviations.

A highly sensitive and specific detection of DNS tunnels on a busy network link is an important problem in the arena of network security as it enables administrators to block or otherwise control these potential sources of compromise. A non-exhaustive, but informative, list of potential uses for covert channels (including DNS tunnels) is:

- Data exfiltration

- Communication through restrictive firewalls
- Receiving commands or information from a remote source
- Transport layer for complete VPN solutions

In enterprises that deal with sensitive data such as financial, health, personal or intellectual property information, it is of the utmost importance to control the access to this information. Even if an enterprise does not have information that needs protecting, DNS tunnels should still be blocked in order to prevent malware from communicating[20]. Because DNS tunnels can be used for arbitrary communication, they can be used as command-and-control channels for botnets[2] or any other malicious system that relies on data transmission. Preventing botnets from operating is in the best interest for the Internet as a whole and should be a concern of every user of the Internet.

In addition to the malicious uses for DNS tunnels, there are other uses that could be classed as benign. One use of this is the DNS-interface to the NIST National Software Reference Library[31]. Queries to a specific zone can include hashes of files on a system in order to determine the validity or purpose of said file. Such lookups appear, at first glance, to behave very similarly to DNS tunnels due to the fact that they are indeed transmitting arbitrary data over DNS. Security vendors will utilize the fact that DNS queries, and UDP port 53 in general, are loosely controlled and are almost guaranteed to be permitted in any environment to enable deployed devices to perform this type of basic communication with remote intelligence sources unimpeded.

Existing algorithms for detecting DNS tunnels rely primarily on one of two methods, character frequency analysis and signatures, to detect the presence of a DNS tunnel. The signature-based solutions attempt to identify portions of the tunnelling application data (as opposed to the user data that is sent using the tunnelling application) for which a signature can be built. These signature based solutions are subject to many of the same

problems that plague signature based anti-virus solutions including their inability to detect zero-day¹ situations. For example, if the application changes its communication schemes, the signatures may no longer be valid and may no longer trigger when they are intended and expected to. Additionally, signatures cannot be built for applications that are not available for study, or that are not known of. For applications that use a custom communication scheme that has not had a signature built specifically for it, it is very unlikely that an existing signature will be relevant and will successfully detect it.

Because signature based schemes are not effective in detecting DNS tunnels in a zero-day situation, another method is required. Analysis based on character frequency is proposed in[11] by Kenton Born with similar approaches proposed by several other authors (see section 3.3 for more details) which uses character frequency analysis under some assumed properties. This approach, however, is vulnerable to a type of attack that is able to obfuscate tunnel traffic. More details are given in section 3.3. An approach proposed by Paxson[40] examines information similarly to the approach proposed in this work, however incurs computational and performance penalties due to its suggested implementation.

Due to the weaknesses listed above, it is clear that a new approach is necessary that detects DNS tunnels in a zero-day situation and without the known weakness involving character frequency and probabilistic encoding. A demonstration of how a tunnel could implement an evasion that exploits this weakness is given in section B. This work proposes, and demonstrates the effectiveness of, a method of detecting DNS tunnels that meets these requirements.

The proposed method operates under the assumption that DNS tunnels move more data than a normal domain but do not necessarily do it by moving more bytes than a benign domain. This distinction is important since, on a busy network, a large content or service provider such as Google, Amazon, Facebook or Twitter may make up orders of magnitude

¹*Zero-day* refers to a situation where nothing is known about the attack as it has never been seen or analyzed before.

more DNS traffic by byte count than a DNS tunnel. A brief treatment of the effects of DNS caching on the number of bytes observed in repeated query strings is given in section 5.1.3. By leveraging this fact, it is possible to detect any use of DNS to transmit arbitrary data by measuring the amount of data that is transmitted using a particular DNS domain or subdomain. The details of this measurement methodology are contained in chapter 5 and requires estimating the amount of unique data that is transmitted by analyzing the queries themselves and not simply counting characters in the query string.

This proposed detection methodology is shown to detect DNS tunnels in as few as ten packets (as shown in section 6.1.2) and continues to be robust in highly hostile detection environments such as those that contain a great deal of non-tunnel traffic as well as benign uses of DNS for transmitting arbitrary data. Detection performance on commodity hardware is shown to scale to greater than two gigabits of UDP port 53 throughput per second in a performance-oriented C++ implementation as described in section 5.2.2. This indicates that this methodology does not sacrifice performance for detection accuracy and remains practical for monitoring very large networks.

Section 6.2.7 compares the processing performance of various approaches from the literature to the proposed method, demonstrating that the proposed method has better performance in a great many cases. Section 7 compares the detection performance of the proposed approach to its peers, demonstrating that in addition to improved computational performance, the proposed approach also provides superior detection performance.

By improving upon both areas over existing methods from the literature, the proposed method constitutes a worthwhile contribution to the field of DNS tunnel detection.

Chapter 2

Background

2.1 Entropy

Entropy[3][4] is intuitively speaking, a measure of how random a particular collection of items is. In the context of a random variable, entropy is a measure of the uncertainty in the output of the random variable. A random variable that has a distribution that outputs a particular symbol seventy five percent of the time has considerably less uncertainty, and thus less entropy, than one that outputs all symbols with equal frequency. In the context of a collection of symbols in a stream, or data source, entropy can be thought of as a measure of the information content of that collection. If the collection comprises almost entirely a single symbol, then that collection can be thought of as containing less information, and thus having less entropy, than a collection where all symbols occur with equal frequency.

Entropy can be calculated for a collection of symbols $\mathcal{C} = \{c_1, \dots, c_n\}$ with symbol c_i appearing with proportion $0 \leq p_i \leq 1$ where $\sum_{i=1}^n p_i = 1$ as

$$H(\mathcal{C}) = - \sum_{i=1}^n p_i \log p_i$$

The base of the logarithm determines the units of the resulting value. If a logarithm

in base 2 is used, then the entropy has the units of *bits*, if the logarithm with the natural base *e* is used then the entropy has the units of *nats*, and if the logarithm is in the common base 10, then the entropy has the units of *digits*.

2.2 Domain Name System (DNS)

DNS is the service through which names are mapped to resources[35][36]. Typically, this maps a name (such as *www.google.ca*) to an IP address. The value of this service is that names are considerably more flexible, and typically considerably easier to remember, than the resource or record that they point to. For example, *google.com* is considerably easier to remember than one of the IP addresses that it points to, such as 173.194.33.70¹. *google.com* is also considerably more flexible, since it points to not just one but several addresses, and successive responses will receive different records in a round-robin, or random, fashion. This rotation of responses allows for a crude, but natural, form of load balancing and automatic fail over, while retaining its ease of use.

The DNS protocol is assigned both User Datagram Protocol (UDP) and Transmission Control Protocol (TCP) port 53 for communication with most communication operating over UDP as opposed to TCP. The use of TCP depends on the implementation of the resolver, however the specification indicates the TCP should be used if the response data exceeds 512 bytes or during a zone transfer[36]. Domain Name System Security Extensions (DNSSEC), due to the fact that it requires a signature of authenticity for all responses (and thus more data to transmit), will often cause the response to require TCP[5]. Because there are no restrictions on when TCP may be used, some resolvers may be implemented to use TCP for all responses as this does not violate the specifications for DNS.

The experiments related to this work do not consider the situation of DNS over TCP

¹As of November 18, 2013.

since the analysis techniques are identical due to the fact that the formats of the UDP and TCP responses is identical once the TCP stream is reassembled. Modification of the tools developed for this work would require the ability to perform TCP stream reassembly in order to extract the DNS queries and responses from the TCP response².

Because DNS is such an integral component of Internet communication it is not generally reasonable to simply block it while still expecting functional Internet connectivity. A common approach, called *DNS proxying*, which forces all DNS queries to be made to a DNS proxy server that is controlled by the interested entity (Internet Service Provider (ISP), company, etc...). This DNS proxy server is responsible for handling all DNS queries for the internal network, and any DNS queries that are destined for the Internet (as opposed to the proxy) are typically dropped by the firewall in this type of configuration. The DNS proxy server operates in *recursive mode*, which means that if a question is asked of it to which it does not know the answer, the proxy server will then query for the answer (by issuing its own query to the global DNS network) and then respond to the initial request using the response from the global network.

DNS is a heavily cached protocol due to how often data can be reused between queries. Consider how often a desktop Internet user causes a request for *google.com*. If a request had to traverse the entire DNS network every time, this would represent a very considerable amount of traffic being generated. To avoid this, every DNS record has extra information included with it that includes, among other things, how long it can be cached for. Standard caching lengths can be around one hour which means that a DNS server will only recursively pass on a query for such a record once an hour. This caching period is not constant and can be set differently, depending on the information the record contains. Some records require a considerably lower Time To Live (TTL), as low as one minute, while for others a considerably longer duration (months) may be appropriate.

²A tool that performs this reassembly is included in the same source distribution as the code for this work.

This proxy architecture removes some naive operation modes for DNS tunnels that will be discussed in more detail in 2.4, but does not offer any protection against the more sophisticated forms of DNS tunnels.

2.3 Covert Channels

Covert channels are methods of communication that use non-standard means of communication for the purpose of evading detection and/or blocking by the existing security infrastructure. Covert channels may utilize portions of an existing protocol[10] or communication channel, or they may find ways of transporting information utilizing a completely new medium. An example of the latter is called a *timing channel*[48], which can utilize the timing between packets to convey information. A timing channel carefully controls the timing between packets sent to a remote server to encode information, thereby utilizing a method of communication that is not utilized by any standardized protocol or communication method.

Covert channels come in many forms and not all types support the properties that are normally associated with a communication channel. Because they are built on unorthodox, or unreliable, transport platforms and are subject to the effects of intermediate routing and networking devices they cannot always offer all of the same functionality as a legitimate channel. For example, covert channels need not support bi-directionality due to either the constraints of the underlying medium, or the effect of intermediate devices. Such a covert channel that is only useful unidirectionally is one that utilizes a third-party image hosting service. It is possible to embed arbitrary information into the header portion of an otherwise completely benign JPEG image file[1] which could then be posted to Facebook, Flickr, or any other publicly accessible image hosting service. This image file can then be checked by the remote hosts to pull the information however, due to the nature of the

image services, the remote hosts may have no way of posting information back to the other endpoint, thus making the communication channel unidirectional.

Real time data transfer refers to the ability for a communication channel to send data immediately. UDP, by its very nature, supports this and TCP supports this via the PUSH flag which indicates that data is being sent before a full window has been accumulated. The TCP PUSH flag is used, for example, during a Secure Shell (SSH) connection in order to provide interactivity when typing and viewing output. Timing channels, or any channel that relies on modifying normal system traffic instead of generating their own traffic, by their nature, are unable to support real time data transfer. This is because they need to wait for a system packet in order to send their data, and if the system goes for a period of time without sending data then the covert channel must wait as well.

2.4 DNS Tunnels

DNS tunnelling is the method by which arbitrary data is transferred over the same channels as DNS. DNS tunnels come in one of two primary types: raw, or conforming.

Raw DNS Tunnels

Raw DNS tunnels do not attempt to mimic or conform to the DNS specifications, and simply attempt to utilize the fact that UDP port 53 is often left relatively uncontrolled in firewalls. Raw tunnels attempt to exploit this by transmitting arbitrary traffic using UDP port 53 packets with arbitrary payload ³. This is the most efficient exploitation of the ubiquity of DNS as it incurs the lowest amount of overhead, both computationally and in terms of network throughput. The trade off for this efficiency is that it is the least conforming and the most likely to get stopped by either a firewall or a proxy. In

³Iodine demonstrates this behaviour when operating in its raw transport mode

the situation where all DNS queries are forced to be proxied through a dedicated DNS server, raw DNS tunnels will fail to operate as expected. This is because when the UDP port 53 traffic is redirected to the proxy, the DNS server will attempt to interpret the arbitrary payload as a DNS packet and will likely fail. When it fails, it will drop the packet thereby preventing all raw UDP port 53 communication. Because these types of tunnels are effectively blocked by standard firewall and proxy practises, detection of these tunnels is not considered in this work.

Conforming DNS Tunnels

Conforming DNS tunnels produce DNS packets that conform to all appropriate specification and RFC documents and, as far as any DNS server is concerned, the traffic generated is valid DNS traffic. These tunnels incur the highest computational and throughput overhead, but have the advantage that detecting and blocking them is a very difficult process. The detection of this type of DNS tunnels is the topic of this work. This type of tunnel is capable of operating in almost any environment, even those with very strict firewall and proxy policies. Because this type of tunnel operates in very hostile (to the operation of the tunnel) environments, detection of this type of DNS tunnel is of interest to all levels of government and industry.

Conforming DNS tunnels operate by embedding the data for transmission into the query string and response, requiring a modified, non-conforming, server on one end of the connection and a piece of software on the client end. Typically these types of DNS tunnels have one endpoint that is controlled by the tunnel user, with that controlled endpoint running dedicated server software. The client and server software are responsible for transforming arbitrary information to and from DNS queries and responses. The precise details of how the translation is done between DNS and the raw data depends entirely on the implementation.

DNS Tunnel Software

Some existing DNS tunneling software currently available is OzymanDNS[42], Iodine[19], Dns2tcp[18], DNScat[41] and DeNiSe[21], and PSUDP[9]. Each of these have slightly different operational characteristics, but they all aim to do the same thing, which is transmission of arbitrary data over DNS.

Iodine supports raw tunnels, as well as moving information back from the server in IPv4 addresses (A), mail server (MX), arbitrary text (TXT) and many other supported DNS record types. TXT records are rarely used by consumers or end-user applications, and so a blanket block policy of TXT records for end-user devices would have very little impact on end-user applications. TXT records are as close to a raw tunnel that a conforming tunnel can get to in terms of throughput since they allow large blocks of largely uncontrolled content.

DNScat utilizes a type of DNS record that is an alias to another record (a CNAME record) and a supplementary A record, when appropriate⁴. OzymanDNS and DeNiSe utilize solely the TXT record, which is as close to the raw tunnel as possible, however can be easily blocked by simply blocking TXT records. Because these tools only use TXT records it is possible that they are the least flexible and deployable out of those listed above given a hostile environment. Dns2tcp utilizes either TXT or KEY records⁵ which makes it as flexible as OzymanDNS and DeNiSe. The KEY DNS record was designated for specific uses[23], but has been deprecated now[34] in favour of the DNSKEY record for use with DNSSEC[53] and IPSECKEY for use with IPSEC[45]. Because of this deprecation, use of the KEY record is subject to strict filtering which greatly reduces the effectiveness of this solution.

⁴The A record is not actually used for throughput but rather to give a 'termination' point for a sequence of CNAME records

⁵Historically KEY records were used to transmit encryption related keys, however their use has been phased out by other DNS record types

All of the above tools utilize encoding and decoding mechanisms that would successfully propagate through a proxy server, at the cost of the fact that the tunnel applications have to generate their own traffic. PSUDP, proposed by Kenton Born, aims to remove the latter requirement by creating slack space within existing DNS packets at the UDP transport layer. He proposes two ways of creating this slack space: naively placing it at the end of the packet, or rearranging the DNS query string to utilize pointers to create this slack space in the middle of the packet. Pointers allow for the re-use of DNS query strings within a packet to save on space. They are an optional component of the specification, but allow for considerable space savings. A pointer in a DNS packet is a special sequence of bytes that indicates where in the packet the processing of the query string should jump to. When processing a query string, only a single pointer can be followed, according to spec, which prevents multiple redirection and infinite loops (where a pointer points to itself). By having a pointer point forward in the packet, it is possible to cause the parsing of a query string to skip a number of bytes, creating slack space.

This method relies on this slack space, which is not parsed by normal servers or clients, but can contain arbitrary data that is extracted by special clients. However, because this slack space is not processed by DNS servers, in an environment where all DNS queries must go through a proxy, this method is incapable of producing a covert channel that is able to penetrate a strictly proxied DNS environment.

Chapter 3

Review of the State of the Art

The solutions that exist to date to detect DNS tunnels generally make very little use of complex and static signatures, but rather attempt to exploit a characteristic trait or property that the DNS tunnel will exhibit. If a tunnel can be crafted to not exhibit that feature, then those detection strategies will normally fail in their detection. This section summarizes the known detection methods, their commonalities, advantages, and disadvantages. It will then take the body of detection methods as a whole and establish the existence of any gaps, or weaknesses, that could be exploited by an application to circumvent an IDS that made use of *every one of the mentioned detection methods*. This will essentially identify gaps in the current state of the art that could be filled by a new technique, or an adaptation of an existing technique.

3.1 General Covert Channel or Anomaly Detection Research

All of the work in this section is aimed at detection of general covert channels, and does not specifically focus on DNS tunnel detection. Because of the non-specificity of these

approaches, a direct realtion to the method proposed in chapter 5 is not appropriate.

(Browne, 1994)[13] establishes an entropy conservation based approach for testing the completeness of general (that is, not specific to DNS) covert channel analysis and detection methodologies. (Shaffer, 2008)[49] Proposes a Security Domain model for assessing the surface of a piece of software for exploitable covert channels.

(Ray, 2008)[44] proposes a protocol for use in a covert channel that incorporates stealth, low overhead, data integrity, data confidentiality, and data reliability. The protocol can be used on top of any other covert channel transport method (ICMP, IP, HTTP, DNS, etc...)

(Horenbeeck, 2006)[29] discusses, briefly, DNS tunnels and their implications. A short mention of proxying DNS requests is given as a potential solution but without examining the multitude of ways that a DNS tunnel could still operate in such an environment. The rest of the paper discusses the risk management and policy based mitigations that can be applied to covert channels in general.

(Moskowitz, 2003)[37] investigates the link between anonymity and covert channels. It identifies several linking factors such as covert channel capacity and the properties of anonymizing networks, and investigates how these affect the anonymity of participants in a communication.

(Newman, 2007)[38] discusses covert channels in a broad sense, examining the various types of covert channels along with the relationship between covert communication, cryptography, steganography and secrets.

(Okamura, 2010)[39] discusses a fascinating type of covert channel for communication between virtual machines that share a physical host. The paper proposes that virtual machines can manipulate their CPU core load, which is peripherally visible to other virtual machines on the host, in order to send and receive information.

Tunnel Hunter[22] is an application that aims for general covert channel detection over a variety of tunnelling communication channels.

3.2 Non-DNS Related Research

Since these approaches do not target DNS tunnels specifically, comparisons to the method given in chapter 5 do not apply.

(Bauer, 2003)[6] discusses a new type of HTTP-based covert channel that adds the unwitting web browser application to the anonymity set.

(Borders, 2004)[7] discusses a method of detecting data egress using HTTP-based covert channels.

(Cabuk, 2004)[15] and (Cabuk, 2009)[16] discuss the design and detection of IP (Internet Protocol) based covert timing channels. (Gianvecchio, 2007)[26] discusses an entropy-based approach to detecting covert timing channels on the Internet based on their effect on the original process' entropy properties.

3.3 DNS Covert Channel Research

The SANS Institutes's InfoSec Reading Room published a report on the design and detection of DNS tunnels[24]. The report covers a very wide variety of topics including background information, tunnel-specific information, technical information, existing applications, detection techniques, detection implementations, and a sample detection scenario. This report is exceptionally good reading as a primer on the topic.

The sample detection scenario employs an analysis technique very similar to the technique that will be outlined in chapter 5.

(Karasaridis, 2006)[33] proposes and evaluates mechanisms that use network flow data¹ to detect DNS anomalies including cache poisoning and tunnels. Their detection of DNS tunnels involves estimating average packet size distributions over a given time interval (in

¹Flow data is a way of digesting network packet data into information per communication, stream, or (in the case of UDP since there is no inherent concept of a stream of interrelation of packets) temporally contiguous collection of packets.

the paper hourly distributions were produced), and then comparing the actual distributions with a baseline distribution using cross-distribution entropy computation. The authors are able to observe considerable changes in their cross-distribution entropy measurement during the onset of the Sinit virus in their real-world data. This approach is discussed in additional detail in (Roolvink, 2008)[47].

Because this technique makes use of flow-level data instead of packet level data, it is inherently less able to make distinctions based on information contained in the packets. Information that typically gets removed in the process of digesting the packets into a flow includes query/response information as well as burst and timing information from within the interval over which the flow spans.

(Born, 2010)[8] discusses a way of using javascript in a web browser to exfiltrate data from a network, while [10] discusses a novel way of crafting a DNS tunnel that exploits the nature of a DNS packet and the ability to create unused space in the packet in which arbitrary data can be stored. [11] discusses a method of detecting DNS tunnels by examining character and n -gram frequencies in the names that are being queried for. [12] demonstrates the effectiveness of data visualization when attempting to detect a DNS tunnel using a custom visualization engine using the character frequency analysis proposed in [12]. If a DNS tunnel can be crafted such that its character frequencies are distributed sufficiently close to those of legitimate DNS names, then it is possible to hide a DNS tunnel from this type of analysis.

A weakness of the approach proposed by Born is that it relies on the assumption that DNS tunnel traffic *necessarily* has a character frequency distribution that is different from that of normal DNS queries. This assumption is not necessarily true, and a proof-of-concept application was developed that demonstrates this fact. The tool performs a *probabilistic encoding* that takes an arbitrary data source and encodes it into a stream of characters that conforms to a given distribution. This tool can be used to modify the output of

any DNS tunnel application so that their output conforms to the distribution that Born found for normal DNS traffic (or any other distribution, for that matter). By utilizing this transformation, Born’s approach fails to detect the DNS traffic as it becomes, from the viewpoint of his algorithm, indistinguishable from normal DNS traffic. Details of this proof-of-concept tool are given in appendix B and a demonstration of its effectiveness in this situation, including the robustness of the method proposed in chapter 5 against this evasion, is given in chapter 7.

(Butler, 2011)[14] demonstrates a way of quantitatively analyzing covert communication channels with particular focus on DNS covert channels. It proposes a *codeword mode* of communication over DNS where a specific lexicon is chosen that allows the two endpoints to communicate with each other. Each word in the dictionary has a particular meaning² that is understood by both endpoints. This lexicon must be chosen *a priori* and must be common to all endpoints wishing to communicate using this method. Butler also discusses the concept of perfect stealth of a covert channel based on DNS, and proposes a deep packet inspection based countermeasure that utilizes the *Jensen-Shannon divergence* measure.

This method relies on an assumption very similar to that in Born’s character frequency analysis where the tunneled traffic uses DNS names with a measurably different character distribution than that of legitimate DNS traffic. Because of this similarity of assumption this approach suffers from the same vulnerability as Born’s.

(Romana, 2007)[46] discusses their analysis of DNS data on a large campus network. They use the output of a DNS resolver’s query logging as their input, but the process works just as well on other inputs even though they are not discussed. Digestion of the large query log file is done with standard Unix utilities and logic available on almost all Unix-based systems. The authors estimate the entropy of the source IP address (of the DNS query) and the queries themselves, and perform analysis based on that output. The scalability

²The words can represent binary information, or can represent higher level constructs such as commands in the context of a botnet or piece of malware.

of this approach is not discussed in detail, nor is an adaptation of it proposed to consider more refined sets of queries since this approach only takes all parts of all queries together for analysis. Adaptation of this approach for real-time analysis is not discussed either. Because this approach does not discriminate with respect to finely grained slices of time, nor to domain or subdomain information, it necessitates that all of the query data be kept around for postmortem analysis. The approach is able to alert to the fact that something was detected, but it cannot indicate precisely when, alert in a timely fashion, nor can it give any indication as to what caused the alert. These details must be ascertained from the raw query data after the proposed approach throws an alert which increases response time and the manpower required to investigate an alert.

In essence the approach given by Romana is very similar to that given in chapter 5 albeit far more coarsely applied. The implementation details, however, are vastly different and represent the most significant drawbacks of Romana's approach. The method in chapter 5 is able to alert within seconds of the triggering event as well as include contextual information that facilitates fast response and low manpower requirements.

(Thomas, 2011)[52] proposes and evaluates the efficacy of a Field Programmable Gate Array (FPGA) based solution for detecting malicious DNS packets on a high throughput network link. The work extends prior work to include more flexible detection and to support more current-generation network infrastructure (the previous work was limited to 100Mbit network connections, with the new work operating on 1000Mbit network connections). The fact that this approach makes use of specialized hardware makes it prohibitively complex for smaller companies to implement and use. The analysis performed on the DNS packets in order to determine their validity is done via a signature-based system where the DNS query is hashed, and the hash is compared to a blacklist of domains that are disallowed based on the network policies. Because this is signature and blacklist based, this approach suffers from the standard problems such as weakness against zero-day situations and the

inability to be agile in the face of an adaptive attacker³. The tests of this system use highly synthetic testing methods that do not involve real-world data or synthetic data that attempts to resemble real-world data. Due to the lack of real-world data in the tests, it is possible that the results do not apply directly to situations in which it may be deployed.

(Dietrich, 2011)[20] examines the use of DNS for command and control of botnets based on the reverse engineering of the *Feederbot* botnet application. Based on the lessons learnt from Feederbot, the authors applied their methods to other real-world traffic and detected other botnets that also use DNS as their command and control medium. The authors make use of two different approaches for classifying malicious DNS traffic from benign and legitimate traffic. The first approach makes use of entropy calculated over the DNS queries, very similar in theory to the character frequency analysis proposed by Born[11] and is vulnerable to similar methods of circumvention⁴. This portion of their approach operates on a single packet at a time, and does not consider aggregate information. The authors also propose the use of behavioural analysis on data and statistics gathered from the aggregate of several packets to estimate the persistence of DNS queries as well as the amount of data moved over DNS by each host on the network. The persistence of the connection is estimated by considering the maximum time between DNS packets whereas the throughput over DNS is measured by counting the bytes in all of the data segments of response packets and summing over time. The persistence analysis can be countered by a botnet that models its communication with the command-and-control server based on a Poisson distribution with inter-packet times imitating that of legitimate traffic. Similarly, the throughput analysis can be countered by employing codewords and rate limiting to utilize a high-level protocol compression to reduce the amount of data that needs to be

³In this case, if the attacker chooses to use a new domain the tunnel will succeed since the system does not have the new domain on its blacklist yet. Similarly, if an attacker is using a publicly available domain for their tunnel, the blacklist can affect other legitimate traffic on that domain.

⁴That is, if the botnet architected the data in the examined fields to conform to the author's model of benign traffic, then the botnet could effectively masquerade as benign traffic and become invisible to this method of detection.

sent, and to reduce the amount of data actually transmitted in a given time period to reduce the footprint of the bot on the network.

This approach makes use of techniques that rely on having host-level identification on the network. Applying these techniques at an upstream location that may only be looking at network traffic through a Network Address Translation (NAT) would be difficult due to the compression of many different hosts behind a much smaller number of addresses (typically one). It may still be possible for this to operate effectively, however this was not investigated by the authors. The approach proposed in chapter 5 does not suffer in this situation.

(Paxson, 2011)[40] is a slide deck that discusses the author’s searches through large campus networks for DNS tunnels in the wild. The author proposes an approach for detecting DNS tunnels that is very similar to the method proposed in this work in that it examines the approximate amount of data transferred per domain and/or subdomain. The author, instead of utilizing entropy measures, makes use of the utility *gzip*⁵ to estimate the amount of data moved under a domain in a given collection of queries. The author’s assumption is that *gzip* is already optimized for compressing data, which can be thought of as measuring the amount of unique data that is contained in the input stream. The author also mentions the codebook/codeword method of embedding information into DNS queries, similarly as to what was discussed in [14]. The author successfully applies this approach to real-world data and identifies DNS tunnels that were previously unknown.

The author defines absolute thresholds for use in detecting which domains classify as being a tunnel as opposed to comparing the data moved from each domain to its peers and performing a more relative and context-aware analysis. The utilization of the *gzip* algorithm reduces the scalability of the approach due to the computational overhead incurred, limiting the applicability to smaller networks or postmortem (as opposed to real-time) analysis.

⁵*gzip* is a compression utility that is used to compress input streams such as archives or other files.

jhind[32] gave a presentation at DefCon 17 that discusses the use of artificial neural networks to identify DNS tunnel traffic. The author proposed that the neural network operate on the euclidean distance between the various queries to a particular subdomain, treating the queries as vectors in higher dimensional Euclidean space. The author successfully detected DNS tunnels as produced by several software packages (Iodine, Ozymandns and Dns2tcp) using the described approach. The approach outlined by the author suffers from the normal training problems associated with neural networks, such as over or under fitting to the training data, which may or may not prove to be problematic in the real world. The author does not go into detail about the accuracy and precision of the neural network which is very important for real-world applications where false alarms and false negatives are costly errors. Additionally, it is theoretically possible for a DNS tunnel to encode its outputs using a codebook where every word has a distance from every other code word that is within a desired range. Such a codebook could be constructed by including all words that, when treated as vectors in n dimensional Euclidean space, lie within a ball of radius r where r is the maximum desired distance. By doing this, it is possible for a tunnel to fit within the neural network's definition of normal and to pass by undetected. It is demonstrated in section 4.2 that by bounding the output distribution of the queries both in character frequency and in length, the queries span a bounded ball in higher dimensional Euclidean space.

Jeffrey Guy[27] blogged in 2009 about visualization as an aid for detecting DNS tunnels by looking at frequency and request/DNS name length plotted together. In the concluding portions of the article the author mentions briefly, and in passing, that the count of the number of different host names per domain could be of value - this piece of information is precisely the foundation of the solution proposed in chapter 5. The author provides no further discussion of this topic, however, and leaves it as an anecdote to the article.

Static signatures exist for at least three common network anomaly detection engines

(Snort[17], Proventia[43], and TippingPoint⁶) engines, with others likely offering similar functionality. It is important to note that these filters and rules do not trigger on all DNS tunnelling applications and may not be robust in the face of a zero-day situation or version update.

Due to limitations of various platforms, the existence of a signature for one does not imply the ability for a signature on all platforms. For example, HP's TippingPoint platform utilizes regular expression matching to perform its operations, and so any signatures that rely on stateful analysis will not be implementable on a TippingPoint.

3.4 DNS Tunnel Detection Landscape

Taken together as a collective body of work, the detection approaches for DNS tunnels can be summarized as follows, with the weaknesses and strengths of each general approach outlined.

- Signature based approaches exist for several popular detection platforms.

Strength: The fact that the platforms are common and already deployed makes it very easy to deploy these signatures to a large number of existing networks.

Weakness: The static nature of the signatures means that they are not flexible enough to effectively identify more than a small portion of the available tunnelling tools.

- A detection method based on flow data, which offers a more scalable approach due to the reduced amount of information that needs to be processed, is proposed which examines average packet length and statistical deviations thereof compared to a normal

⁶TippingPoint does not make information about its filters available publicly, however a personal correspondence with a TippingPoint user revealed that filters 9932 and 9938 trigger on the application data contained in DNS packets generated by Ozymandns.

baseline.

Strength: This approach is flexible in that it is not limited to looking at characteristics of particular applications, but rather at patterns of behaviour that may be exhibited by any DNS tunnel.

Weakness: This approach assumes that DNS tunnel software will exhibit longer packet and query lengths than normal traffic which is not necessarily true. DNS tunnels can use carefully constructed encodings to ensure that their queries stay small enough so as not to stand out against benign and legitimate traffic. Simply limiting the size of their queries will not suffice, since the proposed detection algorithm relies on comparing the distribution to a known normal distribution, however carefully choosing the length of the queries such that they satisfy the normal distribution will allow the tunnel to remain undetected. Further, since this approach relies on identifying a baseline, it is not necessarily suitable for links with a high variability of traffic patterns (perhaps due to time-of-day variability, or where it is not feasible to determine if the chosen normal baseline contains malicious traffic or not) where false alarms and false negatives may become common.

- The use of artificially created slack space in a packet is a novel approach with a great deal of flexibility for creating a DNS tunnel.

Strength: The slack space requires application aware inspection that performs deep packet inspection to determine the existence of, and then the contents of, the slack space.

Weakness: This type of DNS tunnel has a crucial weakness in that this slack space is not processed by recursive resolving DNS servers, and as such will not persist past the first resolver in a chain in such an environment. If these packets are not sent directly to the DNS tunnel server endpoint, the payload will not survive and the

tunnel will not operate. Because of this, no special detection or analysis mechanisms are required, and a simple DNS proxy will suffice in preventing these types of tunnels.

- A form of character frequency analysis is used in several approaches to detect the existence of DNS tunnels.

Strength: This approach makes use of the assumption that DNS tunnels produce queries and/or responses with a measurably different character distribution than that of benign traffic. Since this assumption is quite general, it applies to any DNS tunnelling application.

Weakness: Because this approach relies on the assumption that the distributions are measurably different, if a DNS tunnel were able to construct its queries such that its character distribution matched the expected distribution, then it would be able to evade this type of detection. A proof-of-concept approach and software application are presented in appendix B that is able to perform a loss-less two-way coding from a high entropy source (such as compressed or encrypted data) to a stream whose character frequency matches any⁷ given distribution.

- Hashes and blacklists are used along with an FPGA based implementation for analyzing DNS traffic and blocking packets deemed to be malicious.

Strength: This approach, due to its fast hashing algorithm and FPGA based implementation, scales to very high throughput.

Weakness: Due to the blacklist nature of this approach, it suffers from the same vulnerabilities as other signature based methods; inability to react intelligently to a zero-day situation or clever adversary. Further, since it is built on highly custom hardware requirements, it is not always practical for smaller network operators to

⁷There are some small caveats that are explained in detail along with the rest of the algorithm in appendix B.

deploy.

- A few approaches examine the behaviour of DNS tunnels and their effects on the statistical properties of the queries themselves over time. These approaches consider very similar techniques to the one given in this work, explained in detail in section 5.

Strength: These approaches are considering the most fundamental source of information for a DNS tunnel; the queries themselves. Because DNS tunnels use the queries as their communication, it makes the most sense to attempt to examine these queries for the keys to detecting the tunnels.

Weakness: The weaknesses of the techniques proposed in the existing literature include cleverly constructed queries (such that they sit within a ball of a desired radius in n dimensional Euclidean space), they are not suitable for real-time analysis (such as the use of higher overhead measuring mechanisms like *gzip*), they do not discriminate between different domains or subdomains, or they do not offer temporal resolution that enables adequate response times.

Chapter 4

Problem Statement and Evaluation Criteria

4.1 Brief Statement

The purpose of this work is to investigate the feasibility of real-time DNS tunnel detection that does not suffer the common weaknesses of existing techniques as outlined in section 3.4

4.2 Detailed Problem Description

DNS tunnel detection is a complicated task made more difficult by the fact that DNS tunnel traffic can appear to be completely legitimate network traffic that conforms to all standards and restrictions. It need not violate any established standards or conventions, which makes it difficult to detect against the background of normal DNS traffic based on testing for violations.

This property of DNS tunnels makes them a particularly effective transport mechanism

when data exfiltration or network control circumvention is the end goal. For this reason an efficient method of detecting DNS tunnels is required that can effectively detect a DNS tunnel against normal DNS traffic with a low false-positive rate and that must not be susceptible to existing methods of circumvention.

From chapter 3 it is evident that there are currently several approaches to detecting DNS tunnels that are not signature based as well as signature based approaches of varying flexibility. There is only one mention of performing real-time analysis at the domain/subdomain level, and it is anecdotal in nature with no clear analysis of its merits or validity. The only other similar approach involves aggregating all domains together and taking their queries together for analysis which obliterates any per-domain statistics that could have been gathered.

Looking at this landscape, it becomes evident that there is a highly advanced theoretical DNS tunnel that could evade all of the proposed real-time detection techniques. Any detection that may occur postmortem would not be able to alert to the threat in an adequate time frame to stop the attack in progress. This tunnel would have the following traits:

1. All of its DNS packets would conform to all appropriate DNS RFCs.
2. Its queries would be chosen such that the character frequency distribution matches benign DNS queries (to evade [11] and similar approaches).
3. Its queries would be chosen such that they have a distribution of lengths that matches benign DNS queries (to evade [33] and [24]).
4. Its queries are chosen such that they do not span too great a space when taken as vectors in higher dimensional Euclidean space (to evade [32]).

Item one is already demonstrated in practise by most of the tunnelling applications available, and item two is shown to be possible in appendix B. Item three is easily accom-

plished by splitting queries based on a statistical model of the desired lengths, and item four can be shown to approximately follow from items two and three.

4.2.1 Theoretical Proof of Satisfaction of Item Four for Higher Dimensions

Part I Consider two n -dimensional Euclidean vectors, \mathbf{p} and \mathbf{q} , such that every component is non-negative and \mathbf{p} and \mathbf{q} are chosen from a set of vectors with a finite maximum norm of N . Since every component of $\mathbf{p} = \{p_i\}$ and $\mathbf{q} = \{q_i\}$ is non-negative $|p_i - q_i| \leq \max(p_i, q_i) \leq |p_i + q_i|$. This together with the triangle inequality show that

$$N \geq \|\mathbf{p}\|, N \geq \|\mathbf{q}\| \Rightarrow 2N \geq \|\mathbf{p}\| + \|\mathbf{q}\| \geq \|\mathbf{p} + \mathbf{q}\| \geq \|\mathbf{p} - \mathbf{q}\| \quad (4.1)$$

This result indicates that if vectors with nonnegative components have an upper bound on their norm, then there is an upper bound on the diameter of the space spanned by those vectors. Since DNS queries can be thought of as vectors satisfying the above conditions (when considering the characters as having values according to the ASCII character code table[30]), the result applies.

Part II By further restricting that the queries follow a prescribed character distribution, there are statistical properties that the norms of the queries treated as vectors in \mathbb{E}^n will have.

For this discussion, let χ be the random variable that chooses symbols from the set $\mathcal{C} = \{c_1, \dots, c_N\}$ with probability distribution $\forall c_i \in \mathcal{C}, 0 \leq P(c_i) = p_i \leq 1$ satisfying $\sum p_i = 1$.

By observing that queries can be considered Independent Identically Distributed (IID) observations of χ , the notion of a typical set and the Central Limit Theorem (CLT) can

be applied. Note that for this discussion, the dimension of the vector and the length of the query have the same value and which one is being discussed depends how the object is being interpreted. If the object is being treated as a vector in \mathbb{E}^n , then its components are the ASCII character codes of the characters of the DNS query string.

As the length of the query grows, a typical sequence can be assumed to have each symbol $c \in \mathcal{C}$ appearing approximately $nP(c)$ times. If $v \in \mathcal{C}^n$ is a typical vector with the Euclidean n -norm and $E[\chi]$ being the expected value of the random variable, then asymptotically

$$\exists k \in \mathbb{R} \ni \|v\| \approx k\sqrt{n}E[\chi] \quad (4.2)$$

Further, since the aggregate probability of the typical set approaches unity as the length increases, the probability distribution of norms of vectors in \mathcal{C}^n becomes very tightly packed (having a small variance) around the norm of typical vector as given above.

The above relationships are consequences of the CLT since the components of the vectors (characters in the DNS query) are IID with a finite variance. By the CLT, the variance decreases proportionally to $\frac{1}{\sqrt{n}}$ and the expected value increases proportionally to $\sqrt{n}E[\chi]$.

Part III By combining part I and part II, it is possible to make statements about the properties of IID observations of χ for large numbers of observations, n .

Part II shows that for large n , it is reasonable to assume that the queries fall into the typical set, and thus will have a norm very close to that of a typical vector which is approximately proportional to $\sqrt{n}E[\chi]$. Part I indicates that if there is an upper bound on the norm of a collection of vectors, then in the case of DNS queries there is an upper bound on the diameter of the ball containing all such vectors. In this case, it is reasonable to assume that vectors will have a norm of approximately that of a typical vector, and thus the ball will have a diameter of approximately twice that, or proportional to $2\sqrt{n}E[\chi]$.

The consequences of these results indicate that by enforcing that DNS queries from a tunnel follow an appropriate character and length distribution, a maximum is placed on the diameter of the space spanned by the resulting queries. Thus, by controlling the target character distribution and query length distributions, it is possible to control the diameter of the resulting query space, thus satisfying item four.

4.2.2 Empirical Demonstration for Smaller Dimensions

Because the results in section 4.2.1 rely on typical sets and the CLT, they only hold for very large query-lengths (dimensionality). It is important to demonstrate that the results apply for more reasonable query lengths and observation counts.

The empirical results were obtained by selecting a probability distribution that matches common DNS domain names. Such a distribution was built from the Alexa top one-million domain names as retrieved in June 2013 and is given in table 4.1.

Character	Frequency	Character	Frequency	Character	Frequency
-	0.0117991	C	0.0366638	P	0.0276835
0	0.0023847	D	0.0321575	Q	0.00204425
1	0.00317953	E	0.0986792	R	0.0635874
2	0.00292606	F	0.0166497	S	0.0654764
3	0.00190129	G	0.0243943	T	0.0611875
4	0.0018249	H	0.0253057	U	0.0328546
5	0.00136287	I	0.0731983	V	0.0129516
6	0.0012152	J	0.0055877	W	0.0127842
7	0.00113832	K	0.0184679	X	0.00655439
8	0.00151475	L	0.0468207	Y	0.0178645
9	0.00124893	M	0.0332358	Z	0.00695053
A	0.091269	N	0.0606199	-	0.000025004
B	0.0240761	O	0.0724148		

Table 4.1: Probability distribution used for DNS character frequency simulation when sampling vectors.

The distribution in table 4.1, when considering the ASCII character code values of the given characters [30], has an expected value of 74.8797 and a mean value (that is, not

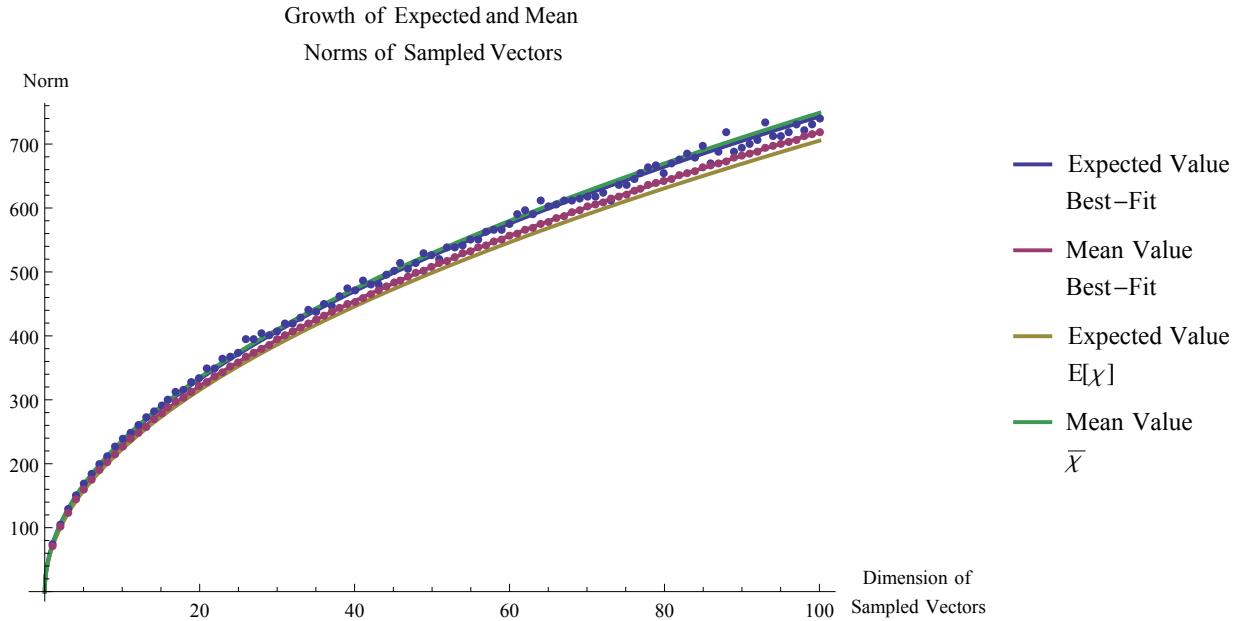


Figure 4.1: The computed expected and mean values of the norms of the vectors from higher-dimensional spaces formed by IID observations from the distribution in table 4.1 are shown as points. The variation of the expected value from the best-fit curve is due to the relatively small sample taken from the enormous spaces.

considering the probability of a character) of 70.5263.

For dimensions up to $n = 100$, vectors were sampled from the spaces with components chosen IID from the distribution given in table 4.1. From the sampled vectors the expected and mean values were computed and a function of the form $a\sqrt{n}$ was fit to both. The growth of the expected and mean values is shown in figure 4.1 as well as the best-fit curves and curves with coefficients taken from the seed distribution.

It is important to observe how the various functions compare to the sampled data. Observe that the best-fit curve for the mean value of the norms of the sampled data almost perfectly describes the sampled data. The best-fit curve for the mean values is given by $71.8527\sqrt{n}$ while the yellow curve is given by $\bar{\chi}\sqrt{n} = 70.5263\sqrt{n}$. The marked difference is indicative that k (as in equation 4.2) is non-unitary. Similarly, the best-fit curve for the expected value of the norms is given by $74.3039\sqrt{n}$ while the green curve is given by

$E[\chi]\sqrt{n} = 74.8797\sqrt{n}$. Note that the best-fit curve in this situation is subject to much higher uncertainty which may be responsible for some of the deviation.

In addition to the above, a cumulative density function (CDF) was built from the sampled vectors and then a function of form given by the CDF of the normal distribution with mean μ and standard deviation σ

$$CDF[N(\mu, \sigma)] = \frac{1}{2} \operatorname{erfc}\left(\frac{\mu - x}{\sqrt{2}\sigma}\right) \quad (4.3)$$

was fit to the data where

$$\operatorname{erfc}(n) = \frac{2}{\sqrt{\pi}} \int_n^{\infty} e^{-t^2} dt \quad (4.4)$$

is the complementary error function (used in the CDF of the normal distribution). Fitting a function of this form to the data allows for an estimation of the mean and variance of the distribution of the sampled vectors as predicted by the CLT.

As is shown in figures 4.2 and 4.3, empirical results show that the theoretical results given in section 4.2.1 still hold and are valid for query lengths/dimensionality within the practical ranges.

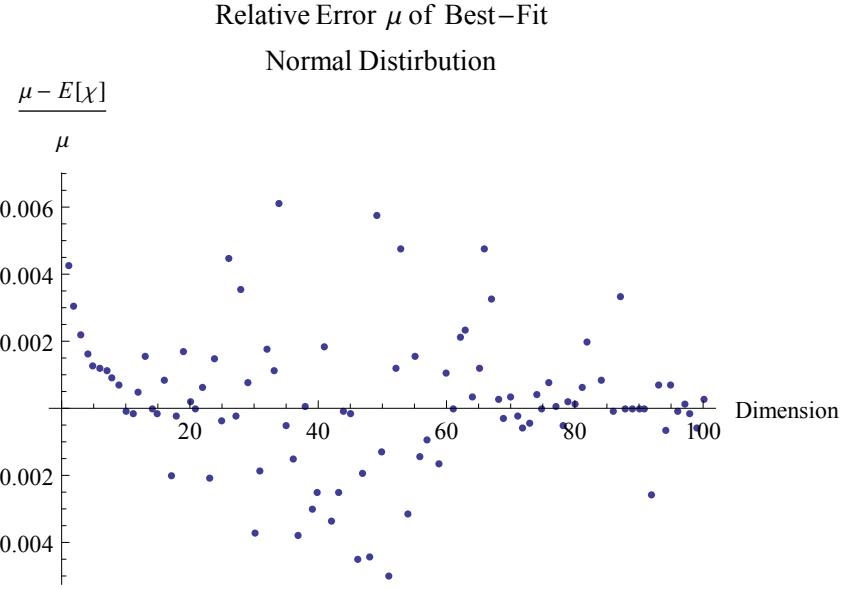


Figure 4.2: For each n , vectors were sampled from χ^n where χ is given by the distribution in table 4.1 and a normal distribution $N(\mu, \sigma)$ was fit to them. This plot shows the relative error of the best-fit μ compared to $E[\chi]$. The trend toward zero for small dimensions becomes dominated by error due to sampling as n grows.




Figure 4.3: For each n , vectors were sampled from χ^n where χ is given by the distribution in table 4.1 and a normal distribution $N(\mu, \sigma)$ was fit to their CDF. This plot shows the trend of the best-fit σ as n increases. The trend toward zero for small dimensions becomes dominated by error due to sampling as n grows.

4.3 Solution Evaluation Criteria

The objectives that must be met for an approach to have successfully solved the problem posed in section 4.1 are as follows:

- Successfully discern tunnel traffic generated from existing tunnel applications and theoretical tunnel traffic (built using additional parameters to attempt to hide from known detection methods) from a baseline of normal traffic.
- Be resistant to known obfuscation methods compared to existing detection methods.
- Be able to operate at high speed on general purpose, easily obtainable hardware.

The proposed approach will be evaluated against these criteria to determine whether or not it can be considered an improvement on the state of the art for this type of detection.

Item 1 will be validated by comparing the chosen approaches against the proposed approach in a relative scoring fashion. Methods will be compared to their peers for relative detection performance, and improvement therein, in the various test scenarios. Methods will be scored based on false positive rate, with lower rates being more desirable. Successful approaches must be highly specific with a very low false positive rate in order to prevent overloading alerting systems with unhelpful information. The standard requirement of a high detection rate with low false alarm rate is present here and by making relative comparisons, it is possible to make statements about improvement on the current state of the art.

Because the implementation of the approaches is built on a Python framework and is not tuned for high performance, it is not reasonable to set absolute performance metrics to measure the success of a detection method. Instead, in order to satisfy item 1, a relative ranking will be applied that examines the performance of the methods in the context of their peers, all implemented using the same data source and common Python framework.

The framework is described in section 6.2.1 and the processing performance of the methods is examined in section 6.2.7.

Item 2 will be tested using a next-generation tunnel, described in section B.2 and referred to as `next-gen`, that simulates what DNS tunnelling applications may look like in the future. Due to the implementation details of the next-gen tunnel, there is no server-to-client transfer direction for that tunnelling application.

Chapter 5

Proposed Detection Method

The method proposed in this work examines the information theoretical properties of the DNS queries to each domain, thus retaining the flexibility to filter and alert per domain as opposed to more generally on the set of all DNS queries. The tools developed to test this approach utilize full packet data for its analysis, but can be modified to use name server query logs (as were used in [46]) or other sources of query information. The prototype C++ software described in section 5.2.2 is easily capable of running at greater than gigabit speed on inexpensive, off the shelf hardware making this approach practical and uncomplicated to deploy on smaller networks or in resource constrained situations. Deployment in large environments is similarly straight forward.

5.1 Theoretical Basis

5.1.1 Assumptions

The detection approach proposed in this work makes certain assumptions about the nature of DNS tunnels in order to effectively detect them. The primary assumption made is that DNS tunnels move more data than a normal DNS subdomain, with a very particular

meaning of *data* that goes beyond simply counting bytes or the number of queries. The concept of the amount of data moved under a DNS domain involves considering the entropy of the queries as a whole, and not the characters that make up a query. The list of assumptions follows:

- DNS tunnel applications use the queries themselves to transport data from the client to server.
- This mechanism will force there to be more unique queries per domain (or subdomain) proportional to the amount of unique data transferred from the client to the server.
- In a server-to-client transfer of data, there will still need to be acknowledgements sent from the client to the server, with the acknowledgement data encoded in the query string.

The primary assumption, in the language of DNS queries, is that DNS tunnels will cause more unique DNS queries to a domain (or subdomain) than benign traffic. If a DNS tunnel is able to construct its network traffic in such a way that this assumption is no longer true, then the proposed approach will be ineffective in detecting it.

5.1.2 Theory

In a large Internet provider network, it is possible that there could be many copies of the same DNS query - say *google.com* - each of which would count towards the total number of bytes or queries transferred to/from that domain. If a naive approach to detection is used, such as counting bytes or queries to/from a domain, then this kind of repetition will have a detrimental effect on the metric calculated for some popular domains.

In order to work around this, the proposed approach takes a different stance on unique data and instead uses entropy to measure the amount of data moved in the queries to a

domain or subdomain. By considering queries as atomic objects, and maintaining a tally of the queries to a domain, and their counts over an interval, a probability distribution function (PDF) is generated. By computing the entropy of this PDF, a basic measure of throughput is achieved. However, since there is value in the capturing the length of the queries that were sent (since longer queries are moving more bytes than shorter ones), the entropy is multiplied by the average query length (in bytes) over that interval. This metric, which will be referred as the Domain Length-Weighted Entropy (DLWE), is the primary mechanism through which the proposed approach detects DNS tunnels.

With this new measure of data, the approach considers intervals of time and computes the amount of data estimated to be moved by each domain over that interval. By sorting all domains by their data throughput, the heavy-hitters can be examined in each time interval with white-listing preventing many of the top benign contenders from causing alerts.

As will be shown in chapter 6.2.7, this approach is capable of processing packets nearly as fast as a naive approach with equivalent or better detection performance which is shown in chapter 7.

5.1.3 The Effect of DNS Caching on Detection Effectiveness

Because the packet capture was done in an environment where a large portion of the clients use one of only a few different DNS servers, the effects of caching will cause the naive approach to have better detection performance than if this were not the case. If this were an environment such as a large ISP or Internet backbone where such DNS caching were not present then the naive approach would have far different detection performance and characteristics. Unfortunately, due to the nature of the network capture obtained for this work, it is not possible to demonstrate how the naive approach performs in a uncached DNS environment.

In order to grant some possible context to the effects that DNS caching has had on the

naive method’s performance, a simple comparison is offered. Data was taken from a home network serving five computers and smartphones, with DNS traffic logged over a twenty-two day period to match the time frame of the capture for real world data. Over that twenty two day period, *www.google.com* was queried 5645 times compared to the 268842 times the same query was seen in the real world traffic capture. It is important to modulate these values by the number of hosts that the real world traffic represents, which is on the order of approximately thirty thousand, or six thousand times the number of hosts the home network was supporting.

Approximately scaling the home network by a factor of six thousand results in an estimated three million queries to *google.com* occurring in the real world traffic, of which only a twelfth actually appeared in the capture due to DNS caching. This is only one common domain name, so applying similar logic to other domains, it is easy to see that in these scenarios, the normal curves would take on a very different shape, easily obscuring tunnel traffic for low throughputs.

A small sample of data was collected from Merlin’s caching DNS servers which represents *every* DNS query made of them, regardless of whether those queries were served from the cache or not. This type of DNS sampling presents a high load on their instrumentation infrastructure, and so is only reasonable for this type of comparison. The logs obtained from their servers span four hours from 1200 to 1600 on a Thursday afternoon. Figure 5.1 shows the DNS query logs from Merlin’s servers, the home network’s DNS query logs, and the queries from the packet capture and their counts.

Figure 5.1 shows the effects of query caching on the repetition counts of queries in networks. The home network, mentioned above, as well as Merlin’s network are represented in order to demonstrate scale. The horizontal axes represents the count of queries, normalized as a proportion of the maximum count. Because the compared networks and captures have vastly different scales, with the packet capture spanning weeks and the query log capture

only spanning hours, it was important to normalize the data sets in order for a direct comparison to be possible.

Each plot was built by tallying the DNS queries in each capture, sorting by count, and then dividing by the largest count. The y -value on the plot then represents the proportion of unique queries that had a count greater than xM_d where x is the horizontal value and M_d is the maximum count for that particular dataset.

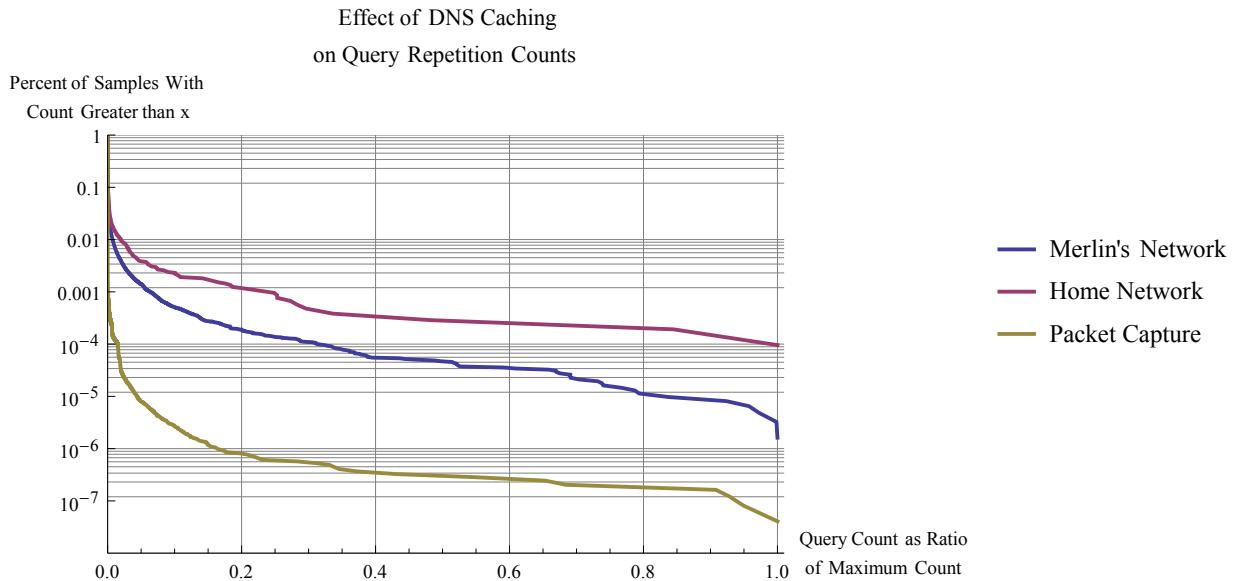


Figure 5.1: Shows the trends for DNS query counts in networks with and without caching. The axes are normalized to account for the fact that the real world data has much more data than the home network. By normalizing the query counts, it is possible to perform a direct comparison.

As can be seen from the plot, which is logarithmic in the vertical axis, it is evident that, on average, the queries in an uncached environment occur over two orders of magnitude more frequently than in a cached environment. This increased occurrence in an uncached environment would have a strongly visible impact on the naive metric's ability accurately detect DNS tunnels. The additional increase in the home network over the query logs from Merlin's network is due to the fact that there may be on-premise customer caching DNS servers on Merlin's network that are not, and can not be, accounted for in the logging

obtained from Merlin.

As can be seen from the proposed method’s underlying mechanisms, the DNS caching actually provides a *more pessimistic* detection environment compared to the uncached environment. Unlike the naive method, whose detection performance results will not be applicable in an uncached environment, the proposed method can be expected to perform better in an uncached environment than in the testing in section 7.1.

Born’s method will also improve in detection performance in an uncached network, since it will cause normal traffic to have a character distribution more heavily skewed away from uniform, resulting in larger metrics and better certainty. Paxson’s metric, however, will suffer in an uncached environment due to the fact that additional data, even if it is highly repetitive, will still increase the metrics of normal traffic and may obscure some low throughput DNS tunnels in the process.

5.2 Implementation

5.2.1 General

The implementation specifics of this approach differ slightly between the C++ and Python versions. The differences come primarily in efficiency and performance of the approach, with the C++ implementaiton designed for very high throughput applications. Both implementations, however, share a common architecture:

- The input is DNS queries and a timestamp at which the query was seen.
- The DNS query is broken down into a top-level domain (TLD) - such as *google.com*, *yahoo.ca* or similar - and the rest of the query - such as *www* in *www.google.com* and *plus* in *plus.google.com*.

- For each TLD, a data structure is created that maps queries to an integer, which represents the number of times that query was seen in the current interval.
- At the end of each time interval, the DLWE of each TLD is calculated, and the collection of TLD+DLWE pairs is output.

5.2.2 C++

The C++ implementation ingests raw packet data in the PCAP file format, relying on a purpose designed network protocol dissector for the extraction of the DNS query string, TLD, and any response information. The query timestamps are obtained from the packet header which is part of the PCAP format.

The data structure used for storing the query-count mappings for each TLD is a custom high-performance, low memory usage, red-black tree written as part of libodb[25]. Computation of the DLWE for each TLD at the end of each interval is handled asynchronously in a separate thread so as not to block the ingestion of packets on the main thread. This asynchronous behaviour allows the processing to operate on very busy networks and allows for inherent and elegant non-blocking buffering of bursts of traffic beyond the processing rate of the host.

The C++ implementation is capable of processing in excess of four hundred thousand packets per second on commodity quad-core CPUs from 2011. The source code is included in the libodb source code distribution[25]. Because there is an existing reference implementation of the proposed approach suitable for deployment on even larger networks, it is possible for this detection approach to be rapidly integrated into existing security infrastructure. Unlike other approaches in the literature where industry members looking to deploy the approach need to find a way of developing and deploying the method themselves. When a reference implementation that is production ready is made available, the time to

practical use can be drastically reduced.

Because of the advantages of supplying an existing reference implementation, it is possible for other research and development to take place in a much shorter time frame.

5.2.3 Python

The Python implementation details are described in section 6.2.5.

Chapter 6

Detailed Testing Methodology and Processing Performance

A collection of tests were run on several detection methods, demonstrating the performance of this work's proposed method when compared to existing methods from the literature. The detection methods chosen for comparison are

- The n -gram detection proposed by Born[11] because it is well defined and was the most prevalent approach found during the literature search. Additionally, the approaches built on this technique claim reasonable success in detecting DNS tunnels.
- The use of *gzip* on domain and subdomain packet data as proposed by Paxson[40] because it involves looking at data that is very similar to the approach outlined in section 5.2, but makes use of different methods for measuring the data throughput.
- A naive approach that simply measures the volume (in number of characters in the query strings) of packets per domain/subdomain in an attempt to illustrate that simple volume of queries is a highly inadequate approach, and that more sophisticated approaches can perform considerably better.

The collection of methods was put through several tests in an attempt to demonstrate their performance in average (using existing implementations) and worst case scenarios. All of the tests involving traffic generation were performed in a virtual environment of two linux-based virtual guests directly connected via a virtual network on a single physical host. The tunnelling applications were communicating between the virtual hosts, transmitting content from the high entropy source `/dev/urandom` under linux.

Due to the nature of the tunnelling applications, specifically the use of encryption and compression (or the necessary requirement of a high entropy source for the next-gen tunnel) on the data that is transmitted, alternative sources of input data would produce very similar results. Because of this fact, only the one source is considered here and the results should be generally applicable.

6.1 Situational Performance Goals

6.1.1 Determining a Baseline

Through cooperation with Merlin, an educational Internet Service Provider (ISP) in Manitoba, DNS traffic was collected over a period from Thursday November 4 2010 until Friday November 26 2010. The hosts responsible for the DNS traffic observed include several dozen school divisions totalling tens of thousands of individual computers. The capture includes just over one billion packets destined to, or sourced from, UDP port 53 (the standard DNS port). Not all packets are valid DNS with more detailed information on the properties of the sample available in appendix A.

This captured traffic will be used to determine a baseline distribution to which the metrics produced on isolated tunnel traffic can be compared. This baseline will provide context in order to determine if a method is able to detect a tunnel with sufficiently high certainty.

It is assumed that the incidence of tunnels in this baseline traffic is sufficiently low that it can be discounted. This assumption may not be perfectly accurate due to reasons indicated in the introduction. Because some security vendors use DNS to transmit some of their information, these transmissions are in essence a DNS tunnel and so will represent a certain portion of the real world traffic. The effect of tunnels present in the real world traffic given the assumption that there are none will result in a more pessimistic environment for testing. Since some of the traffic that lies further out than the synthesized traffic may in fact be a tunnel, a portion of the false positive rate that the detection approaches will suffer may actually be due to the classification of existing traffic as a tunnel, and not due to misclassification. By making the assumption that no tunnels exist in the real world, however, the false-positive rates given here represent an lower bound and so should transfer well to other real-world scenarios.

6.1.2 Existing Implementation Detection

This test will involve the two hosts communicating at varying throughput rates using the chosen existing DNS tunnel implementations. The throughput rates will scale from as little as several bytes per second, to several megabytes (or as high as the tunnel applications can support) per second. The wide range of throughputs used is done to give an indication of how the detection methods scale with tunneled throughput.

The existing implementations chosen for testing in this section are Iodine[19], DNScat[41], and DNS2TCP[18]. Iodine is chosen due to the fact that it provides a full VPN solution without additional work by the user. DNScat is chosen due to being written in Java and so runs on multiple platforms¹ without the need for a compiler or other complex dependencies that the user must obtain. DNS2TCP is chosen since it does not require root access, and is written in C indicating potentially better throughput than other mechanisms.

¹DNScat will run on any platform that supports Java 1.4 or later[41]

The detection results of each of the approaches on the existing tunnelling applications is of value to those wishing to identify uses of DNS tunnelling currently in the wild. Instead of generating a large number of distinct events that will be detected (or not) resulting in a ROC plt, rather a much smaller number of prototypical events were produced. Details are given in chapter 7 regarding the statistics of the metrics from these tests and how the detection methods will be evaluated. In essence, detection methods will be scored for false-positive rate and ranked against each other in order to determine a relative rating and ranking. The relative performance comparisons allow for a more contextual performance analysis to be done between two methods that perform very similarly.

6.1.3 Demonstration of Existing Weaknesses

The tests in this section will demonstrate that the weaknesses listed in section 3.4 are in fact exploitable by software and are not simply theoretical in nature. For the purposes of this section, a new type of tunnel was simulated using a tool that generates queries according to potentially more complicated rules. Inclusion of the same types of logic into existing tools is possible but would be non-trivial for an ad-hoc evaluation and as such, the inclusion of features into these applications is outside of the scope of this work. The implementation details for this tool are given in section B.2.

As in section 6.1.2, the certainty indicator will be used to evaluate the efficacy of the detection methods on this more advanced type of traffic.

6.2 Analysis Method Implementations

6.2.1 Common Scaffolding

All implementations of the analysis methods are built on the same basic framework written in Python. The framework has the following basic properties:

- Unlike the C++ implementation of the proposed approach described in section 5.2.2, which ingests raw PCAP formatted data, the Python scaffolding is unable to do so. Instead, the Python scaffolding relies on another tool to process raw packet captures (or other data sources such as DNS server logs) into a comma-separated value (CSV) file with a timestamp and a query string on each line.
- The scaffolding takes in two arguments: a file name to read the timestamp/query pairs from, and how long the intervals should be, in seconds.
- At the start of each interval, the main loop initializes an array to hold all of the queries for the upcoming interval. For the duration of the interval, the queries are added to the end of the array, which is passed to the analysis routine (where the specific detection method details are) at the end of the interval.
- When the queries from the interval are handed off for processing, each query has its top-level domain (TLD) identified, and then the rest of the query is passed to the specific method. The returned value is then assigned to the TLD that the query came from. At the end of each interval after processing each TLD, the digested values are output to the console.
- Once processing of the interval is complete and the information has been printed, the array is cleared in preparation for the beginning of the next interval.

There are some very important differences between the C++ implementation and this Python scaffolding that will cause the performance metrics to not be directly comparable.

Unlike the C++ implementation which is designed for maximum performance and non-blocking ingestion of traffic, the Python implementations are poorly behaved around the ends of intervals where ingestion stops while processing occurs. Additionally, there is a reliance on another tool to produce the appropriate input data. This removes the complexity of parsing and verifying the packet protocols from the performance figures for these implementations.

6.2.2 Naive

The metric returned for each query (stripped of its TLD) is simply the number of characters in the query. For each TLD, after all of the packets are processed, the total of all of the metrics for each query are summed together resulting in a single value that counts all of the characters that appeared in queries to that TLD.

6.2.3 Born

Since Born's[11] assumption is that normal DNS queries approximately follow Zipf's Law, and that tunnels produce character distributions that are much closer to uniform, this was used to determine a metric. At the end of each interval, for each TLD all of the queries to that TLD are taken together and joined into a single long string. The character distribution of that string is then computed, and its standard deviation calculated.

Since it is expected that tunnel traffic is close to uniform in its character distribution, it is then assumed that tunnels will have a very low standard deviation that continues to decrease as throughput (and thus amount of data per interval) increases while normal queries will have a standard deviation that increases very quickly by comparison.

Because the next-gen tunnel (as described in section B.2) used in the tests is designed to circumvent Born’s detection method (any any other method that relies on character frequency analysis), it is expected that this approach will have poor detection performance on the next-gen tunnel.

6.2.4 Paxson

Paxson’s[40] approach uses the *gzip* utility to compress the queries to domains under the assumption that it will approximately measure the amount of unique data contained in those queries. Because *gzip* is a command line interface to the underlying *zlib* library, in Python this approach was implemented by making calls to Python’s *zlib* library and its `zlib.compress()` method.

As with the previous implementations, the queries for each TLD over an interval are accumulated and joined together into a single long string. This string is then passed to the `zlib.compress()` method resulting in a compressed string whose length becomes the metric for that TLD in the current interval.

Because Paxson’s approach uses very similar assumptions, that DNS tunnels can be detected by measuring the amount of unique data transferred to each domain and/or sub-domain, it is expected that Paxon’s approach and the proposed approach will have very similar detection performance. However, due to the fact that Paxson’s approach relies on a general-purpose compression library whereas the proposed approach uses a tailored algorithm and optimized data structures, it is expected that the proposed approach will have superior processing performance.

6.2.5 Proposed Approach

In place of the high-performance red-black tree used in the C++ implementation, a Python `dict` was used in the Python implementation. The `dict` is used as a mapping from query strings (with the TLD removed) to integers indicating how many times that string appeared in the interval. At the end of each interval, such a mapping is built for every TLD from the packets seen in the interval, and for each TLD, the DLWE is computed from the obtained distribution. The computed DLWE is then used as the metric for the TLD for that interval.

6.2.6 Tunneling Application Throughput

The tunnelling applications used during the evaluation were subjected to different rates of traffic in both client-to-server and server-to-client directions. For each tunnelling application, sixteen captures were performed at each of the following target throughput rates (in bytes per second) in each direction (client-to-server and server-to-client where applicable):

10, 25, 50, 100, 250, 500, 1000, 2500, 5000, 10000, 25000, 50000, 100000, 250000, 500000, 1000000

The rates stop at ten bytes per second for practical reasons. It is considered reasonable to assume that tunnels with a throughput rate lower than this are ineffective at transmitting sufficient data to be practically useful in many situations. As will be seen in chapter 7, throughput rates lower than ten bytes per second would quickly become lost in normal DNS traffic for even the most discerning detection methods.

Due to implementation details of the applications, and of the next-gen tunnel, not all applications were able to transmit traffic at the target rate. Figure 6.1 shows how the various tunnelling applications responded to various input rates, plotting their actual rate of ingestion of input and the observed number of characters of query output they generated on the network. Note that for the lighter coloured graphs showing the observed output,

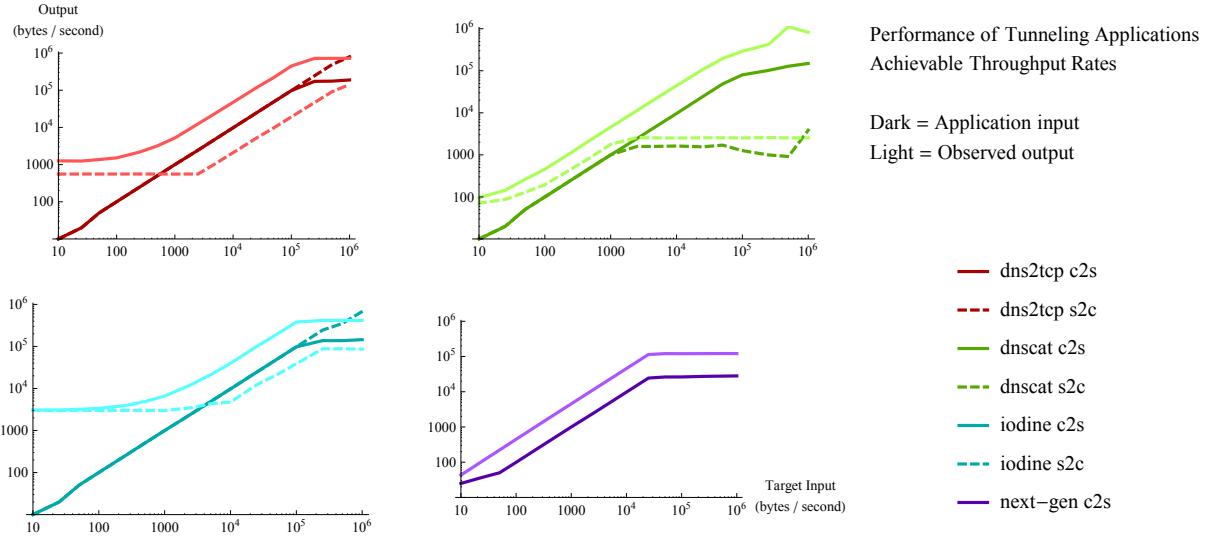


Figure 6.1: Shows the scaling behaviour of DNS tunnelling applications as the input rate is scaled up. Note that not all applications are capable of transmitting data at the rate they are given data, which is visible as a plateau on the right-hand-side of each plot.

the vertical axis is representing a value equivalent to the naive metric. The next-gen tunnel was not tested in a server-to-client direction since it does not implement that functionality, with more information given on the implementation in appendix B.

6.2.7 Detection Method and Python Interpreter Processing Performance

Python has several interpreters available freely in addition to the standard interpreter (for this discussion, the standard interpreter will be referred to as Cython). A notable alternative, called PyPy, is a Python interpreter written in Python itself that contains just-in-time compilation (JIT) mechanisms that Cython does not have. The advantages of JIT compilers and interpreters include being able to dynamically optimize heavily used code paths based on run-time statistics and information in ways that would not otherwise be possible statically at compile-time. Because of this advantage, PyPy can offer an order of magnitude or better speedup[51] in some workloads.

Because of this potential advantage, PyPy was investigated as an alternative to Cython when performing the analysis with the Python-based implementations of the detection algorithms. Figures 6.2, 6.3, and 6.4 show the performance of the various detection methods over aggregate tunnel and real-world data.

Figure 6.2 shows that, on tunnel data, the analysis methods under Cython all suffer, to varying degrees, as the amount of data to process per interval increases. The naive and proposed methods suffer the least, while Paxson’s method suffers by far the most, dropping to approximately half of its original processing rate. Note that Born’s method and Paxson’s method trade places as the slowest performer at a throughput of approximately five kilobytes of data per second.

When looking at the PyPy performance, with the exception of a dip around one kilobyte of data per second, performance increases as the amount of data increases in direct contrast to the behaviour of Cython. This increase in performance can be due to the JIT components having enough time to achieve some measurable optimizations of common code-paths. In addition, the general ranking of the algorithms by performance is maintained (with the exception of the swap of Born and Paxson seen under Cython) from Cython to PyPy. Despite this improvement, the average performance is still nearly an order of magnitude below that of Cython in many of the cases.

Figure 6.3 shows the performance over time of the different detection methods and Python interpreters as more packets are ingested from real world data. Observe that as time progresses, the methods get progressively slower, likely due to inefficiencies in the interpreter and/or method. Again, the naive method is the fastest, with the proposed method performing at approximately two thirds of the speed of the naive method. Unlike with the aggregate tunnel data, no swapping of ranks between Born’s and Paxson’s methods is observed.

Also unlike the aggregate tunnel detection performance, the methods when run under




Figure 6.2: This plot shows the performance of the detection methods and Python interpreters over the aggregate tunnel data. The output is packet processing rate as a function of the target input rate (rate at which the tunnels are transmitting traffic). As the target input rate increases, the structures involved in the methods get larger in a non-linear way, resulting in longer operation times and slower performance.

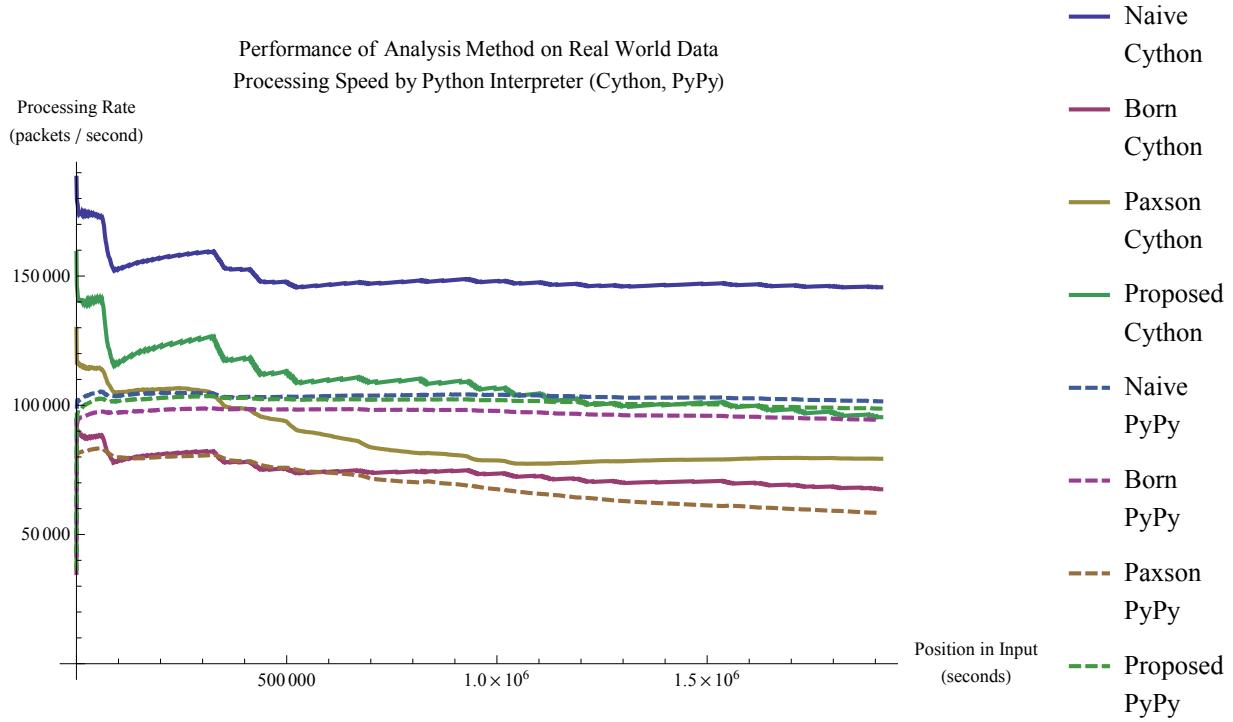


Figure 6.3: This plot shows the performance of the detection methods and Python interpreters on real world DNS traffic. The output is packet processing performance as time progresses and more packets are processed by the script. As more packets are fed into the script, inefficiencies in the methods and/or the Python interpreters themselves become observable in the degradation of performance.

PyPy perform *far* better, even surpassing the Cython counterparts in at least one case. Note that the Paxson's approach when run under PyPy shows decreasing performance over time while no such decrease for larger times is observed when run under Cython. PyPy out-performs Cython when performing Born's approach by a considerable margin beginning very early on. A much smaller performance difference is witnessed between the two interpreters for the proposed method, with PyPy out-performing Cython by a small margin near the end of the sample data.

Figure 6.4 shows the same data, but with a limited scale allowing early-time behaviour to be examined. Note that the ramp-up of PyPy's JIT components is observed in the very early time followed by a very consistent processing rate. PyPy's processing rate, as shown




Figure 6.4: This plot is identical to figure 6.3 but restricts the time displayed to the first one hundred thousand seconds. The 'spool up' of the JIT portion of PyPy is noticeable in the very early time-scales.

in 6.3 can be seen to be far more consistent and not subject to the same degradation and spikes as when run under Cython. This difference in behaviour indicates that the issues observed under Cython are likely related to Cython and not the implementation, since the same source code files defining the detection methods are used under both the Cython and PyPy interpreters.

Figures 6.6, 6.5, 6.7, and 6.8 attempt to represent the performance of the various detection methods and Python interpreters as more tunnel data is moved through them per interval. Their horizontal axes are the data input rate (actual data input rate is used, instead of target input rate, due to the performance characteristics discussed in 6.2.6), and the vertical axes indicate the processing rate (in packets per second) of the approach on the two Python interpreters (Cython and PyPy).

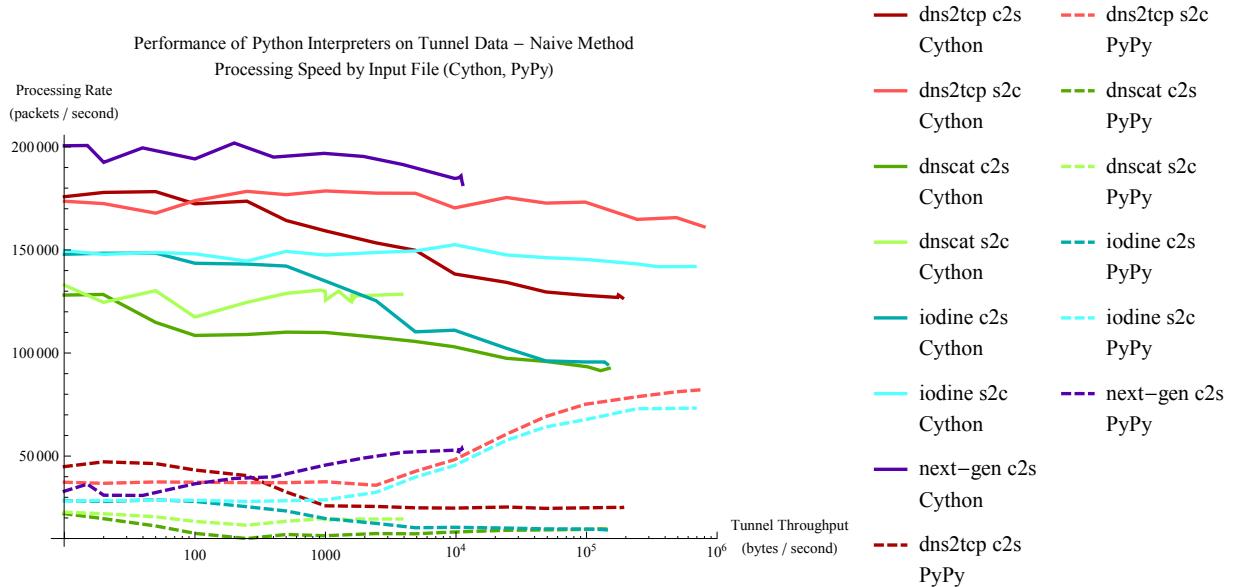


Figure 6.5: Performance of the naive method on separated tunnelling application data, showing processing rate as a function of input rate.

In the above mentioned figures the legend requires some additional context. The plot legends contain labels of the form *dns2tcp c2s Cython* which contains three distinct pieces of information. The first word indicates which tunnelling application being one of DNS2TCP, DNSCat, Iodine, or the application described in 6.1.3 which is indicated by a name of **next-gen**. The second word indicates whether the data being moved over the tunnel is being transferred from the client to the server (**c2s**) or from the server to the client (**s2c**). The final word indicates which Python interpreter is being used.

There are fourteen lines on each figure, each corresponding to a Python interpreter, tunnel application, and data transfer direction triple. The solid lines correspond to runs made under the Cython interpreter and dashed lines indicate the use of the PyPy interpreter.

The performance of the naive method has very little dependence on the input rate, suffering minimally as the amount of data per interval increases. This is expected behaviour since Python’s string objects allow for efficient computation of length, being of $O(1)$ computational complexity[50]. Because the naive method’s implementation must calculate the

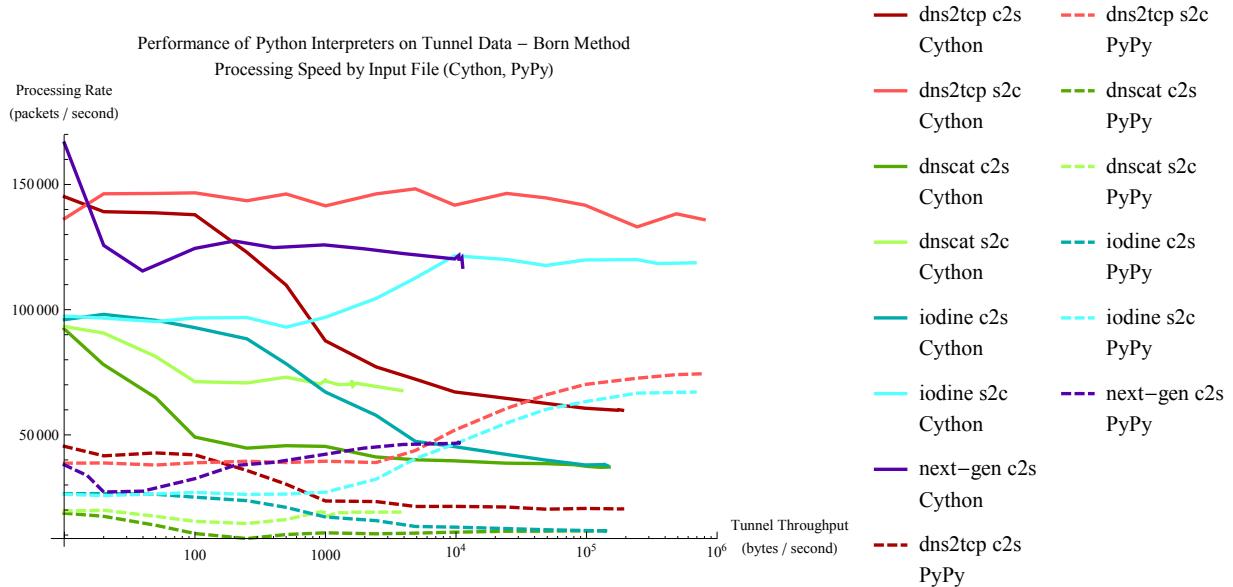


Figure 6.6: Performance of Born’s method on separated tunnelling application data, showing processing rate as a function of input rate.

length of each query, additional queries increase the time complexity of the method linearly with the number of queries that must be processed. The Iodine and DNS2TCP client-to-server transfers show marked drops in performance, potentially due to longer queries being used which would increase the time taken to read the files into the script.

The PyPy performance figures show similar clustering to the Cython plots, with the notable exception that the DNS2TCP and Iodine server-to-client performance improves drastically for higher throughputs. This is again likely due to the JIT components of PyPy being able to optimize for runtime conditions. Despite this improvement, they still perform at approximately half the rate of their Cython counterparts at the most favourable throughputs.

Born’s approach shows a very high level of variation in performance, with marked decreases in processing rate as the throughput is increased for many of the the tunnel applications. Notable exceptions are the Iodine and DNS2TCP server-to-client transfers which suffer minimal degradation or improve as throughput increases respectively. Again,

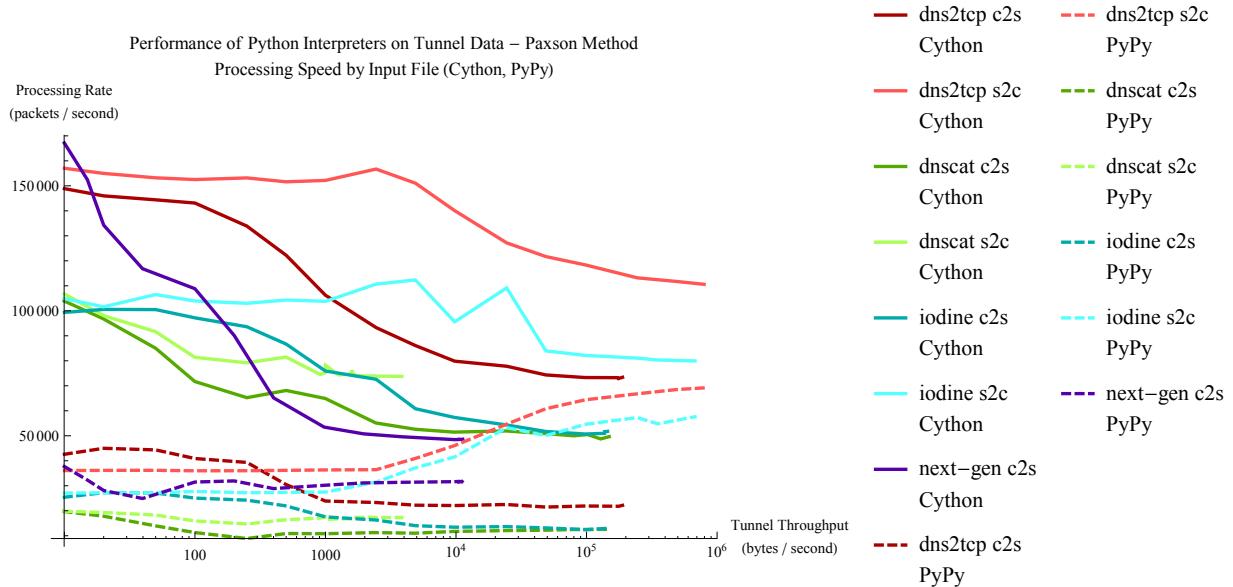


Figure 6.7: Performance of Paxson’s method on separated tunnelling application data, showing processing rate as a function of input rate.

the same patterns as with the naive approach are seen under the PyPy runs, however relative performance overall is much better from PyPy. This is in part due to the much larger performance hit taken by the Cython implementations. This performance is likely due to the fact that the Python implementation of Born’s approach makes use of a collection of loops and arithmetic. This type of computation is not well optimized under Cython but can be subject to excellent run-time optimization under PyPy’s JIT components. PyPy is still, however, unable to improve upon, or match, Cython’s performance in these workloads.

Very similar behaviour is again seen in the performance of Paxson’s approach as compared to Born’s, with Cython out-performing PyPy and a general trend of performance degradation with additional throughput.

The proposed approach shows performance characteristics and trends that match the naive approach far more closely than either of the other two approaches. There is minimal degradation in performance for most of the tunnelling applications, and almost all of the samples under the Cython interpreter are above one hundred thousand packets per second.




Figure 6.8: Performance of the proposed method on separated tunnelling application data, showing processing rate as a function of input rate.

It is noteworthy that only DNS2TCP’s server-to-client remained above that threshold for throughputs beyond five kilobytes per second for Paxson’s approach. For Born’s approach, only DNS2TCP and Iodine’s server-to-client transfers and the next-gen tunnel remained above that threshold for throughputs above five kilobytes per second.

It is instructive to observe that PyPy’s performance overall is considerably lower than Cython on tunnel data, but as is shown in figure 6.3, this is not the case on real-world data. The reasons for this may involve intricacies of the dataset used, the JIT components of PyPy, and the effects of large amounts of data on Cython’s internal structures in this workload. PyPy is a viable choice for some of the methods on real-world data despite being considerably less optimal than Cython in all cases when working with tunnel application data.

6.2.8 Processing Performance Conclusion

As has been shown, the detection methods were tested on both real world and purpose-generated tunnel application traffic as well as on two different Python interpreters and were instrumented for their processing performance.

In the real world data test, the Cython interpreter performed well with the PyPy interpreter matching or exceeding its performance for Born's and the proposed detection methods. When operating on tunnelling application data, the Cython interpreter is categorically faster in every case than PyPy with performance factors ranging from two to two hundred depending on the application, throughput, and method.

When evaluating the performance of the detection methods on real-world data, in all cases the naive method is the fastest, followed by the proposed approach, Born's method, and Paxson's method in order. Under the Cython interpreter, the naive method is approximately fifty percent faster than the proposed method. The proposed method is approximately twenty five to thirty percent faster than Born and Paxson's approaches respectively. Under the PyPy interpreter, however, the performance of the proposed method, and the naive method are very close with a difference dropping to less than ten percent and Born's approach a close third place. Paxson's method falls to the slowest position under the PyPy interpreter.

When operating on tunnelling application traffic, the average performance of the methods (averaged over all tunnels and cases) can be seen in figure 6.2 where the naive and proposed methods both perform well, maintaining processing rates in excess of one hundred twenty thousands packets per second. Born's and Paxson's approaches both suffer severe degradation of performance as throughput increases, resulting in final processing rates well below one hundred thousand packets per second. As was mentioned above, when looking at the tunnelling applications individually (and not as an aggregate), only the naive and proposed approaches are able to maintain processing rates above one hundred

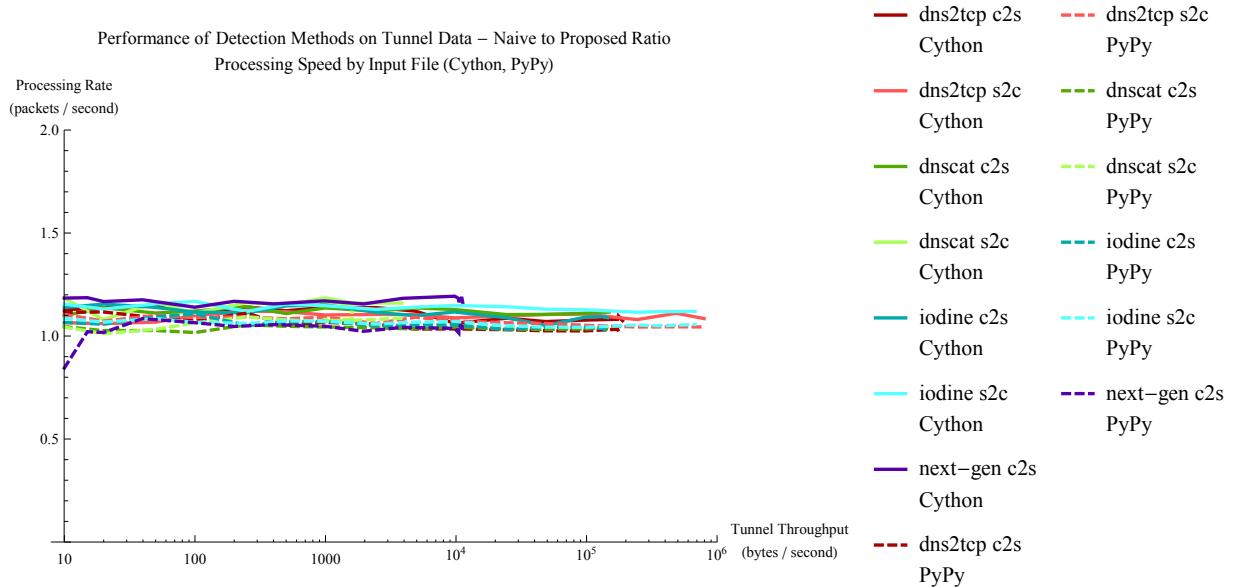


Figure 6.9: The ratio of the performance of the naive method to the proposed method on separated tunnelling application data, with the vertical axis showing the speedup of the naive method over the proposed method.

thousand packets per second for all (or most, in the case of the proposed approach) of the tunnelling applications and throughputs. Additionally, only the naive and proposed approaches are able to prevent significant performance degradation, and in fact show very similar trends when their plots are compared directly.

When the ratio of the performance of the naive method to the proposed method is examined, a clustering very close to a fixed value is observed as in figure 6.9. This indicates that much of the performance degradation, and potentially other performance characteristics, of the two methods are dominated by the common scaffolding and/or the Python interpreter as opposed to by the underlying methods or their implementation.

Through this examination, it has been shown that the proposed method out-performs both methods from the literature by a considerable margin and comes very close to matching the naive method in performance in many cases. Performance relative to Paxson’s method approaches a factor of two in favour of the proposed method in many of the test scenarios.

When compared to Born's method, similar performance improvements are seen in favour of the proposed method. This by itself demonstrates a useful contribution of the proposed method, provided that it is able to detect the tunnelled applications with a sufficiently high certainty.

Additional quantified details of the performance of the methods are given in the final conclusion (chapter 8).

Chapter 7

Tunnel Detection Evaluation

The processing performance of teh detection methods was examined in the previous chapter which demonstrated that the proposed method is able to process packets at a considerably higher rate than the methods proposed by both Born and Paxson. This chapter will examine and compare the ability of the methods to detect tunnelling applications against the background of normal DNS traffic.

In order to obtain the metrics used in this chapter, tunnel application and real world data was separated into adjacent ten second windows for processing. The distribution of metrics across these windows is computed and used to produce the plots for real world data as well as indicative representative values for tunnel applications.

When examining the distribution of metrics produced by tunnel applications, it was discovered that the metrics were clustered extremely tightly around the mean. The relative standard deviation for the various tunnel applications, transfer directions (server-to-client or client-to-server) are given as a function of input rate for each of the detection methods in plots 7.1, 7.2, 7.3, and 7.4. These plots show that the clustering around the mean for these metrics is so tight that a standard receiver operating characteristic (ROC) plot would be of little additional value since the true and false positive rates are dependant on the

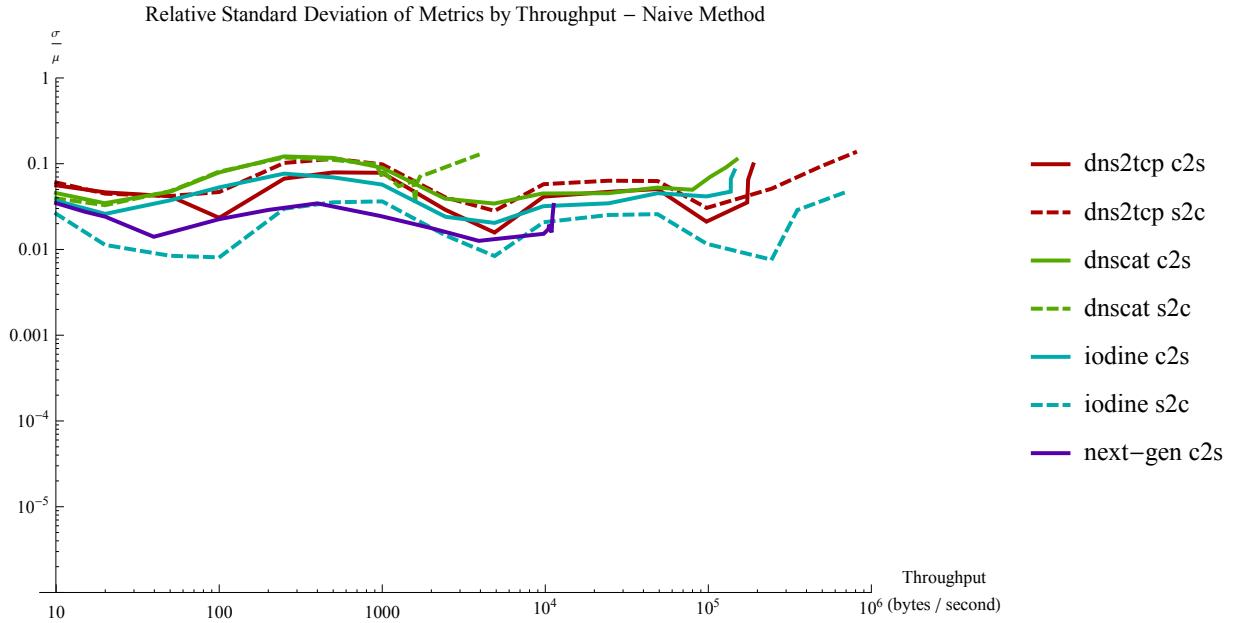


Figure 7.1: Relative Standard Deviation of Metrics - Naive Metric

value ranges that the metrics take. Since the range of values that the metrics take on is so limited, the differences between a low and high percentile in true and false positive rates would be minimal.

Because of the extremely tight distributions and the insensitivity on input data distribution (described in section 6), the mean value will be taken as representative of the tunnel metric for a given throughput, direction, detection method, and tunnel application. This choice of a single representative value simplifies discussion and makes presentation of the salient characteristics of the detection methods more straight forward.

Figure 7.5 shows several log-log plots (in order to be able to provide adequate resolution for both very small and very large throughputs), one for each detection method, that demonstrates how the mean metric generated by the methods scale as the throughput of the tunnel is increased. The expected trends of the detection methods is as follows:

- The naive metric is expected to increase with throughput with a slope dependant on the implementation. In practice, all of the implementations have a slope of ap-

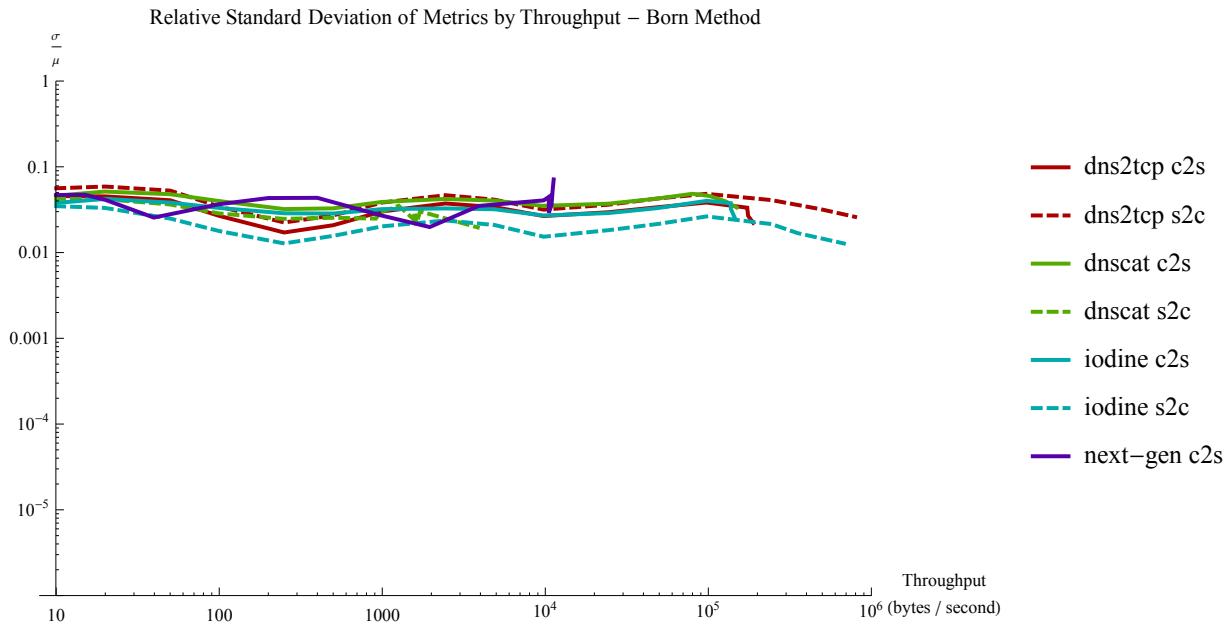


Figure 7.2: Relative Standard Deviation of Metrics - Born Metric

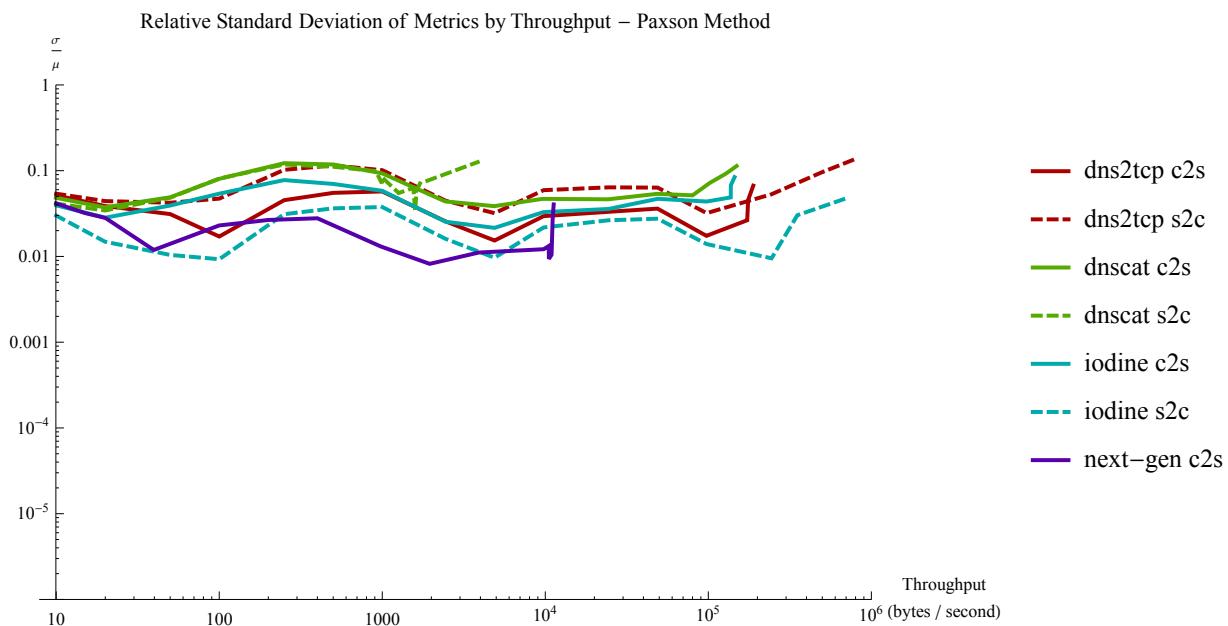


Figure 7.3: Relative Standard Deviation of Metrics - Paxson Metric

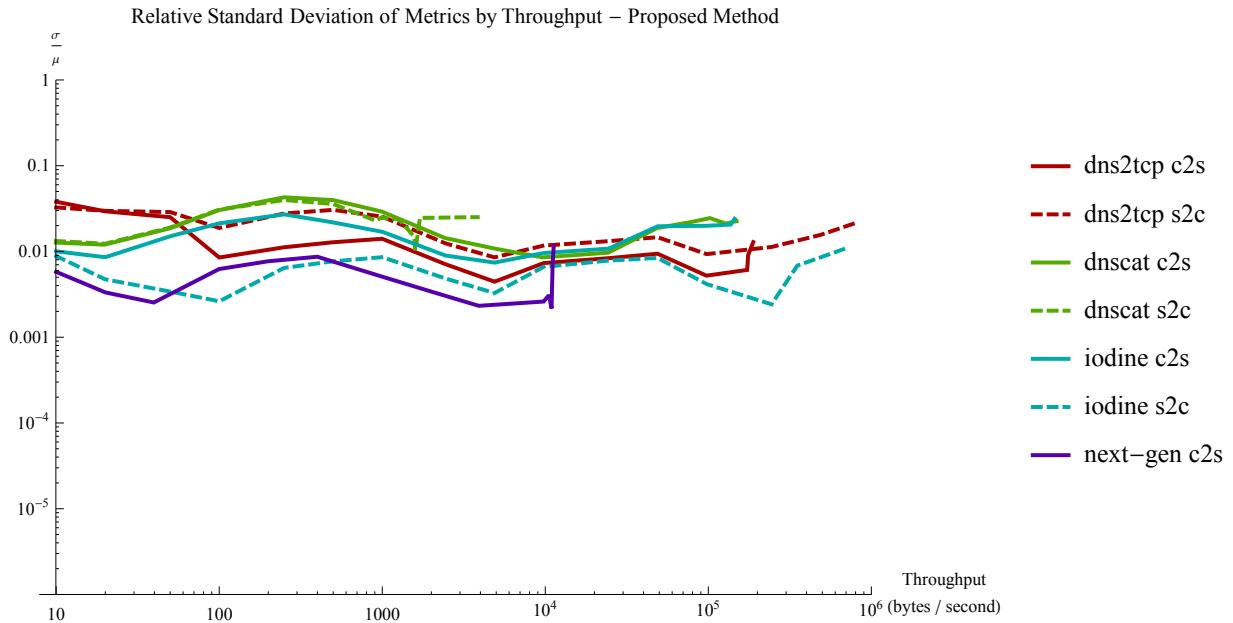


Figure 7.4: Relative Standard Deviation of Metrics - Proposed Metric

proximately one indicating that no significant inflation of the input data takes place (beyond that incurred by encoding to base 32 or 64 as the implementation may choose).

- Born’s metric is expected to decrease to zero for tunnels as their character distribution approaches uniform. The character distribution of tunnel traffic is not directly tied to throughput, but rather has an implicit dependency on it due to the law of large numbers. Since the distribution is approximately uniform in the limit, there needs to be enough sample data before the distribution begins to converge and the effects of small-scale variations begin to average out.

This is the behaviour seen for most tunnels with the exception of the Iodine and DNS2TCP server-to-client transfers as well as the next-gen tunnel. The next-gen tunnel’s behaviour is precisely as expected due to its construction specifically to evade character frequency analysis such as is used in Born’s approach.

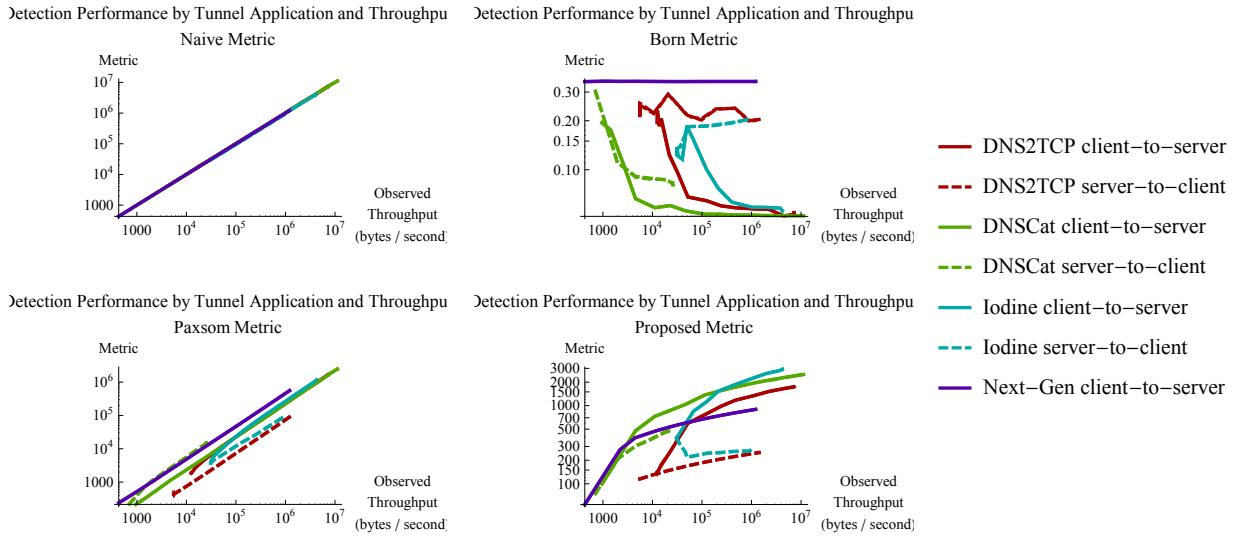


Figure 7.5: These plots show how the metrics computed by the various detection methods scale per tunnelling application as a function of the tunnel throughput rates. Note that since not all tunnels are capable of a full range of throughputs (some generate a minimum amount of traffic, regardless of how low the throughput is), resulting in some of the shorter plots.

- Paxson's metric is expected to increase approximately linearly with a slope that depends on the implementation and on the compressibility of the transferred data. Because the transferred data is pseudorandom, it is nearly incompressible and thus the observed slopes of the lines is derived from the implementation details of the various tunnelling applications.
- The proposed metric is expected to increase approximately logarithmically with additional throughput due to its reliance on entropy as a multiplicative factor in its final metric. This is the behaviour seen for all except Iodine's server-to-client transfer, with varying scaling factors.

As is visible in figure 7.5, tunnels with lower throughput produce categorically smaller (or in the case of Born's approach, larger) metrics. This property of the lower throughput tunnels makes them necessarily harder to detect when laid over top of normal traffic.

	Naive	Born	Paxson	Proposed
DNS2TCP c2s	0.0559644	0.0497584	0.0450118	0.0371504
DNS2TCP s2c	0.0477585	0.0675144	0.0495964	0.0392231
DNSCat c2s	0.0155828	0.0526164	0.018919	0.00980975
DNSCat s2c	0.0152077	0.0452746	0.0172972	0.0101241
Iodine c2s	0.00720872	0.0439898	0.0106345	0.00500239
Iodine s2c	0.00687431	0.0377363	0.0119256	0.00375623
Next-Gen	0.0309484	0.0581522	0.0399604	0.00414985

Table 7.1: The relative standard deviation given by $\frac{\sigma}{\mu}$ of the multiple samples for each tunnel application and detection method pairing for the lowest throughput rate.

Because of this, the methods will be tested on their ability to detect the most hidden tunnel of each application, which in practice is one of the tunnels that transmits only ten bytes per second, against background normal DNS traffic. The assumption, which proves true in practice, is that tunnels that move more data are necessarily easier to detect, and the most dangerous tunnels are those that go undetected. By testing the methods to ensure they detect low-rate tunnels, their ability to detect high-rate tunnels is implicitly demonstrated.

It is instructive to observe the relative standard deviations for the various tunnel applications and detection methods at the lowest throughput level, ten bytes per second. This is given in table 7.1. As can be seen, the relative standard deviation is extremely low, well under five percent in almost every case, with many samples for the proposed method measuring only a few tenths of a percent.

7.1 Detection Performance Against Real World Data

Since tunnel metrics are represented by a single value, the mean of their samples for a given scenario, the detection methods will be ranked based on how they partition the metrics produced by real world data. For simplicity a thresholding approach will be considered as the classification mechanism, with anything below the tunnel’s metric classified as legit-

mate, and anything above the tunnel’s metric classified as a tunnel. Detection methods will be scored based on the number of false positives that could be expected to occur given the distribution obtained for normal traffic. It is important to note that Born’s metric, due to its construction, produces smaller metrics for tunnels instead of larger as is the case for the other approaches. This is important to know since the classification is reversed in Born’s case, with lower metrics being classed as a tunnel and higher metrics being classed as normal.

For each detection method, two plots are given that show how the tunnel applications compare to real world data. The first plot demonstrates how the false-positive rate behaves as a function of throughput rate, direction, and tunnel application. The second plot shows the distribution of metrics of real world data with indicators represented by vertical coloured bars placed to mark the mean metrics of various tunnel application and direction pairs as given in the corresponding legend. For each marker, the false-positive rate is the y value of the normal traffic curve at intersection of the normal curve with the vertical marker. The second plot only shows the tunnelling combinations that produce the highest false-positive rates, since those are the scenarios of greatest interest.

Figure 7.7 shows the performance of the naive metric, with markers visible for the next-gen and DNSCat tunnels. The lowest marker belongs to the next-gen tunnel, and has a false positive rate of 0.0799815, followed by a DNSCat server-to-client transfer with a rate of 0.0496253. Iodine’s client-to-server transfers are the most easily detectable with its lowest false positive rate at 0.00164529 indicating a very low ambiguity when classifying its traffic.

The naive method, due to the nature of the the capture involving (relatively) very few duplicate queries, performs quite well overall even on the next-gen tunnel traffic. Figure 7.6 shows the trends of the false positive rate for the tunnelling applications as a function of the data throughput. The lack of duplication in the DNS queries and its impacts were




Figure 7.6: Trend of False Positive Rate for Tunnels by Throughput - Naive Metric

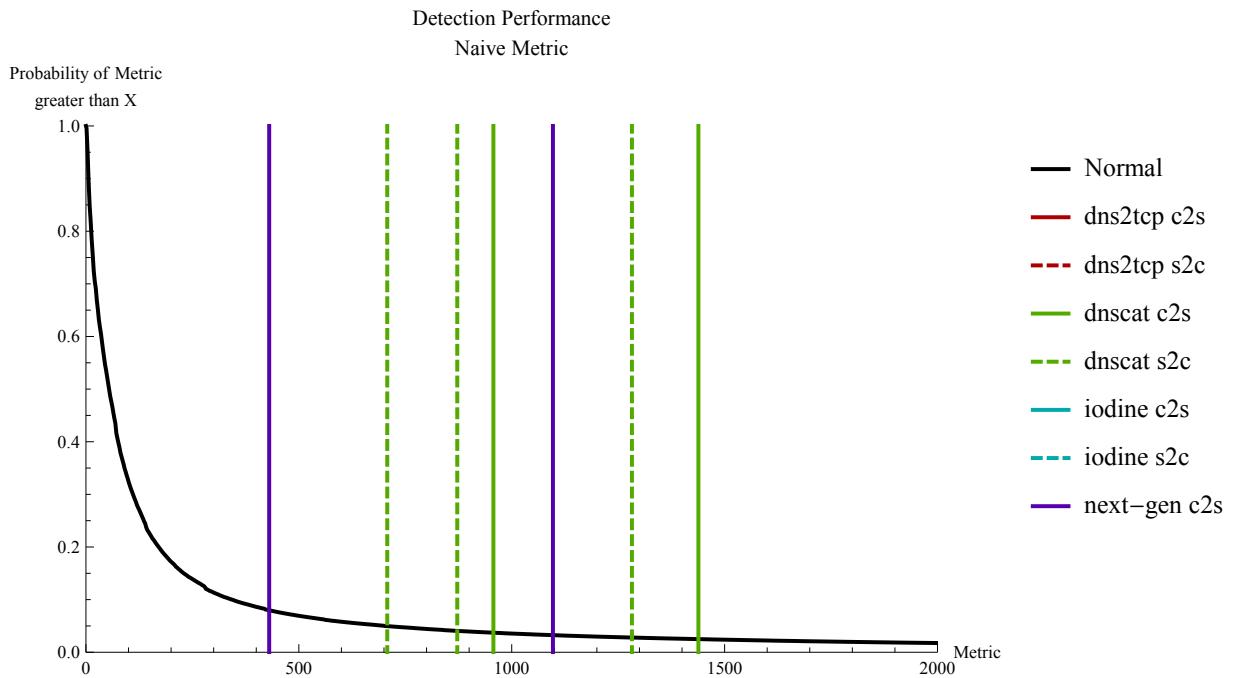


Figure 7.7: Tunnel Detection Performance - Naive Metric

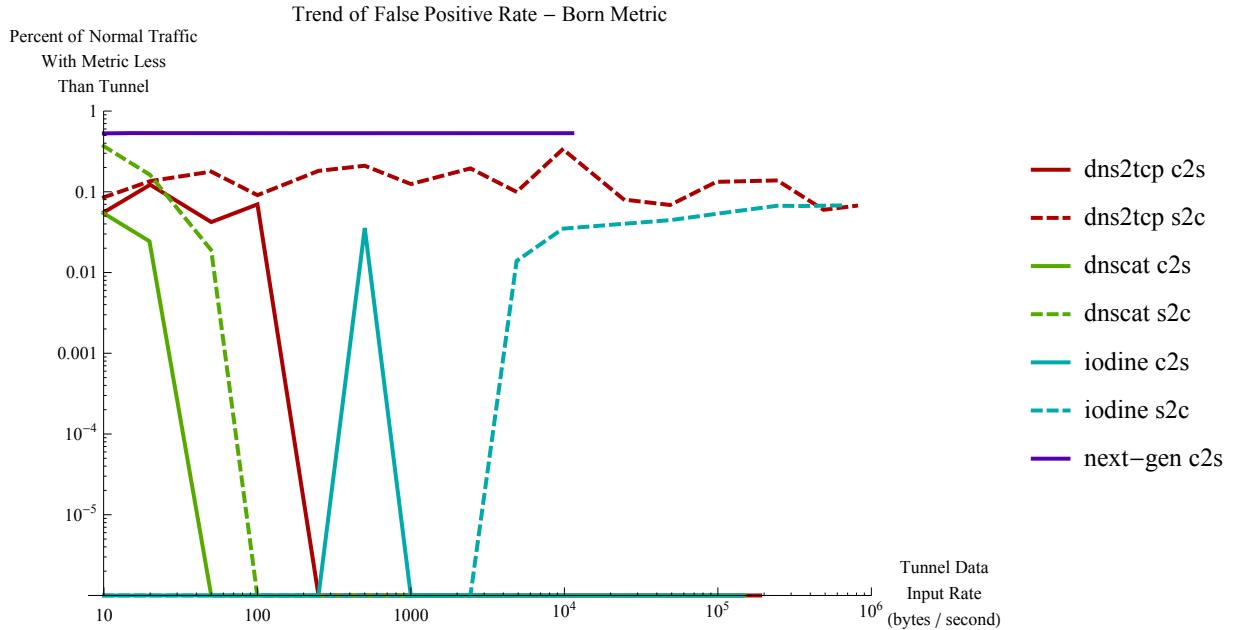


Figure 7.8: Trend of False Positive Rate for Tunnels by Throughput - Born's Metric

discussed in section 5.1.3 in greater detail.

Figure 7.9 shows the performance of Born's metric with almost all markers visible clustered near the smaller metric values. Since the implementation of Born's metric produces smaller metrics for tunnels, the highest values are the ones of interest. The highest false positive rate of 0.529347 occurs for the next-gen tunnels. It is not easily seen in the plot, but all sixteen (one for each target throughput) of the next-gen tunnel markers are superimposed on each other at that value. This is due to how the tunnel was implemented for the proof-of-concept, as all traffic generated very closely conforms to a specified distribution. This distribution produces a particular metric under the implementation of Born's approach, independent of the amount of throughput. The next-gen tunnel markers are followed by the DNSCat server-to-client marker with a false positive rate of 0.367036. Similarly to the naive metric, Iodine's client-to-server tunnel was the most easily detected tunnel with the lowest throughput samples achieving a rate of no more than 0.00133572.

It is easily seen that Born's metric will not be able to identify the next-gen tunnel with

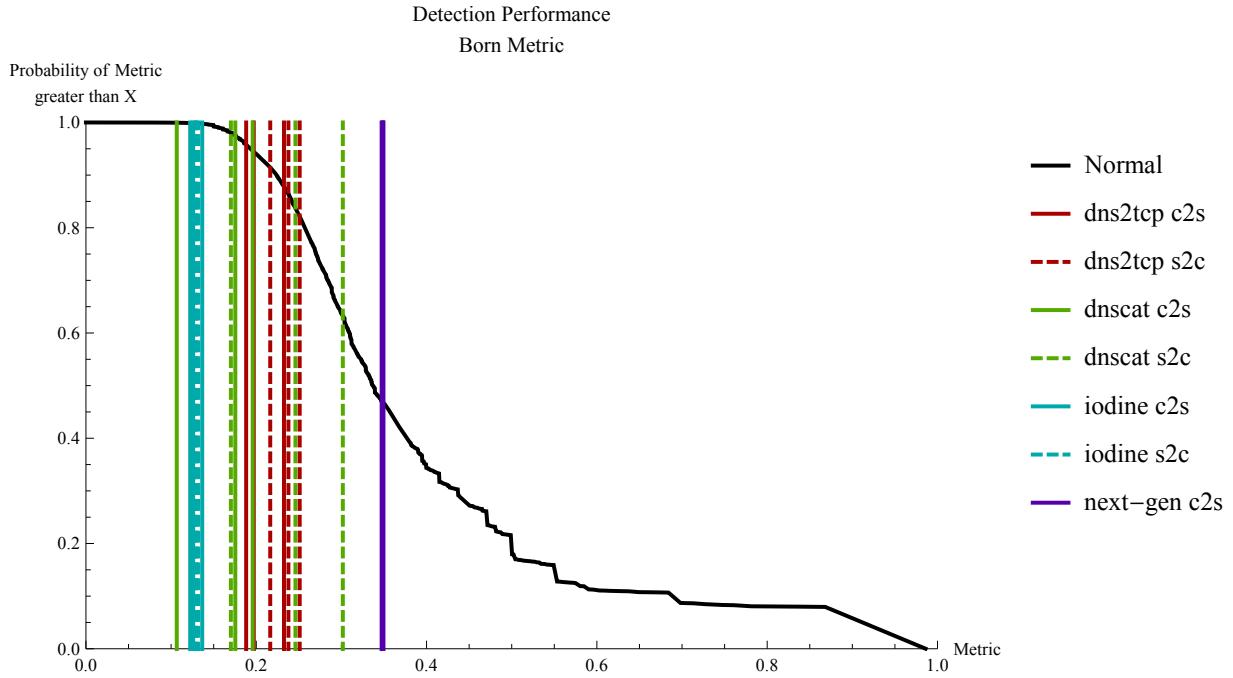


Figure 7.9: Tunnel Detection Performance - Born’s Metric

any reasonable certainty due to the ambiguity caused by the tunnel’s traffic generating metrics very near to the median of real-world traffic. It is possible to augment the next-gen tunnel to adhere less tightly to a given distribution in order to spread its metrics over a wider range surrounding the median, further increasing the difficulty of detection. Figure 7.8 shows the trends of the crossing values for the tunnelling applications as a function of the data throughput.

Figure 7.11 shows the performance of Paxson’s metric similarly to the plot shown for the naive method. The lowest metric is shown by the DNSCat client-to-server transfer with a false positive rate of 0.0239753 followed very closely by DNSCat’s server-to-client transfer with a rate of 0.0232046. The next-gen tunnel produces a false positive rate of 0.0212873 and again Iodine is the most easily detected with its client-to-server transfers producing false positives at a proportion no larger than 0.00392552.

As is expected, Paxson’s method shows considerable improvement over both Born’s




Figure 7.10: Trend of False Positive Rate for Tunnels by Throughput - Paxson’s Metric




Figure 7.11: Tunnel Detection Performance - Paxson’s Metric

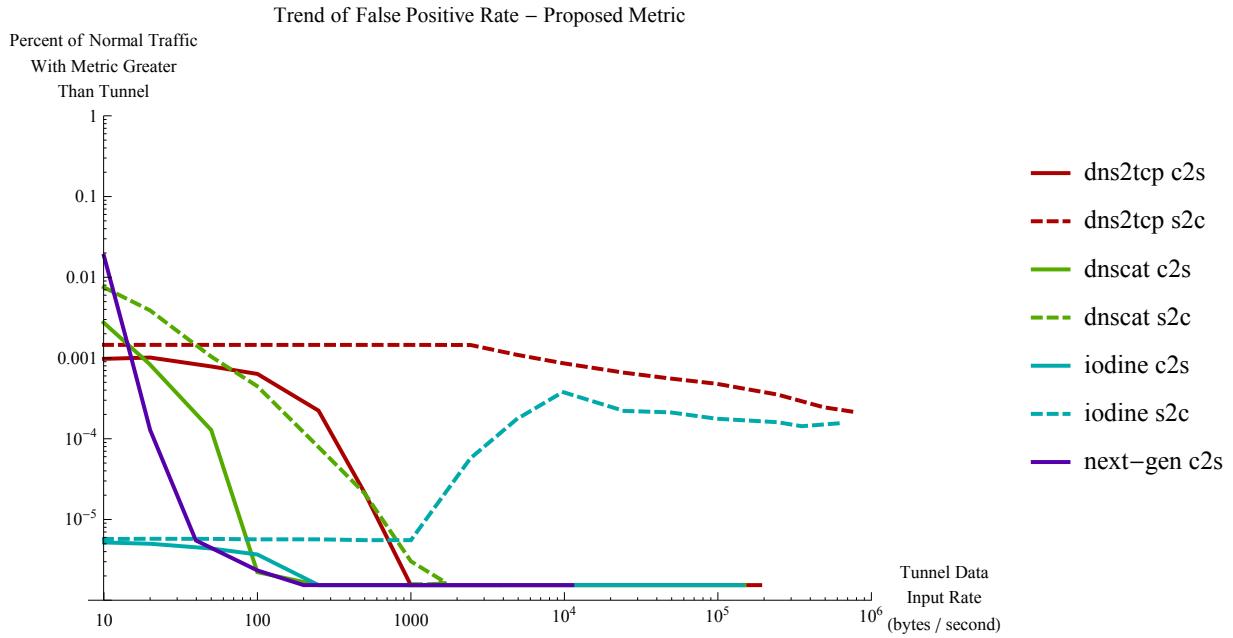


Figure 7.12: Trend of False Positive Rate for Tunnels by Throughput - Proposed Metric

approach and the naive method in terms of ambiguity of tunnel detection. Figure 7.10 shows the trends of the false positive rates for the tunnelling applications as a function of the data throughput.

Figure 7.13 shows the performance of the proposed metric in a fashion similar to the naive method and for Paxson’s method. The highest false positive rate encountered is from the next-gen tunnel at 0.0184927 . This is followed by two of the samples generated by the DNSCat server-to-client transfers, the first of which produced a false positive rate of 0.00751152 . As has been the case for all metrics thus far, Iodine is the most easily detected tunnel generating proportions of false positives for its lowest throughput client-to-server and server-to-client transmissions of no more than 5.19129×10^{-6} and 5.73392×10^{-6} respectively.

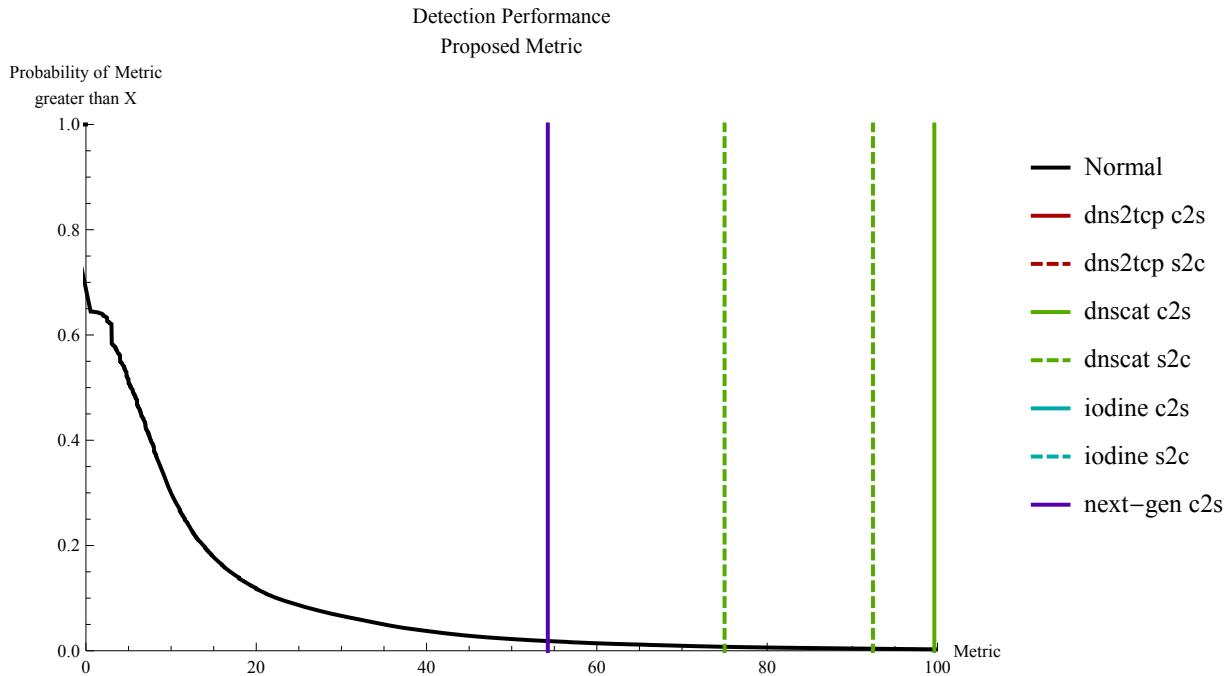


Figure 7.13: Tunnel Detection Performance - Proposed Metric

7.2 Specificity and Ambiguity of Tunnel Classification

It is possible to simplify the above plots and figures into a single chart that plots the minimum detection specificity observed for each method and each tunnelling application. The following charts only consider the certainty of detecting the tunnel in which the method is least certain. By comparing the methods in their most hostile scenarios more substantial distinctions can be observed with clearer separation between the best two methods. In all cases, this corresponded to one of the tunnelling applications at the ten bytes per second throughput level, which demonstrates the effectiveness of these methods to solve the *needle-in-a-haystack* problem of low throughput DNS tunnels on a very busy network link.

Figure 7.14 shows the specificities of each method for each tunnelling application, in the interval $[0, 1]$. Observe the extremely low certainty of Born's metric on two tunnelling applications (next-gen and DNSCat's server-to-client), with lacklustre performances for

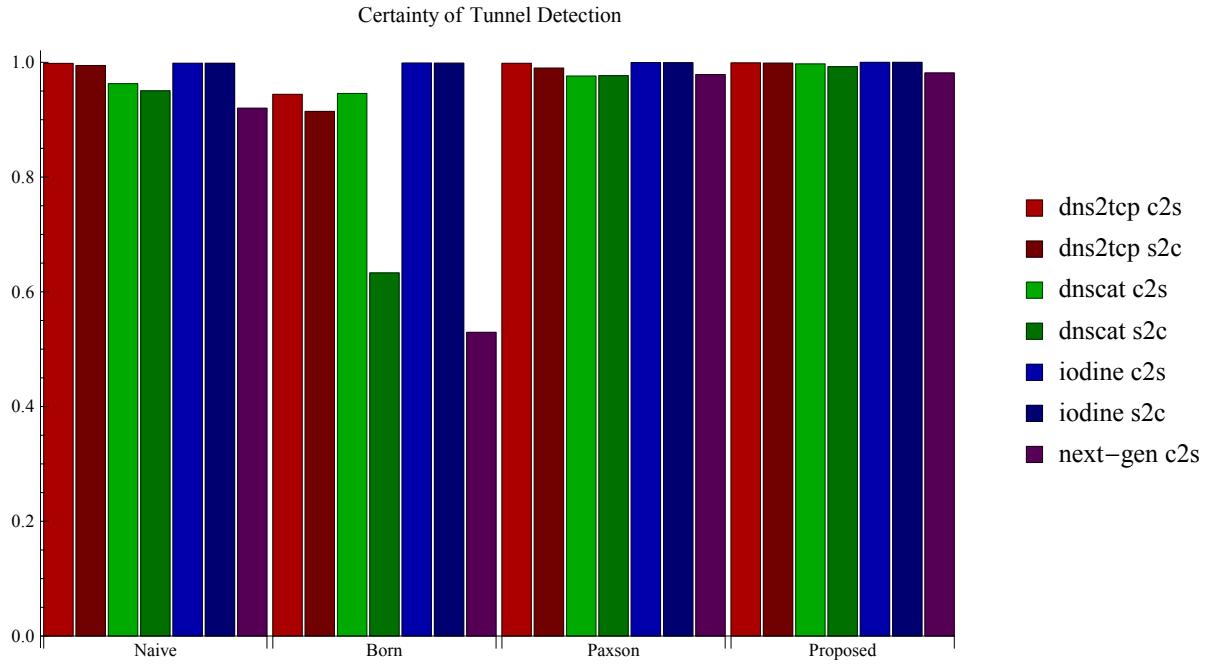


Figure 7.14: Comparison of the specificity of classification of a tunnel against real-world traffic for the least certain tunnel in each detection scenario (method/application pair).

DNSCat’s client-to-server and both DNS2TCP transfer directions. The only tunnelling application that Born’s method reliably picks out is Iodine, however as was shown earlier Iodine is by far the most easily distinguished tunnel by a large margin for all detection methods. The naive and proposed methods as well as Paxson’s method perform similarly with differences that are not easily distinguished in this chart.

Figure 7.15 shows the same data as in figure 7.14 with a restricted range spanning the interval $[0.80, 1.00]$ as opposed to $[0, 1]$. This restricted charting range makes the the differences between the top performing methods more easily visible. In this chart, it is easily seen that the naive method outperforms Born’s method in all except the Iodine case. The Iodine performance differences are not easily visible due to the very close values. The naive method has a certainty of 0.998355 for both transfer directions while Born’s method has certainties of 0.998664 and 0.998512 for the client-to-server and server-to-client directions respectively. The differences, on the order of 0.0003, are not visible in the charts. The

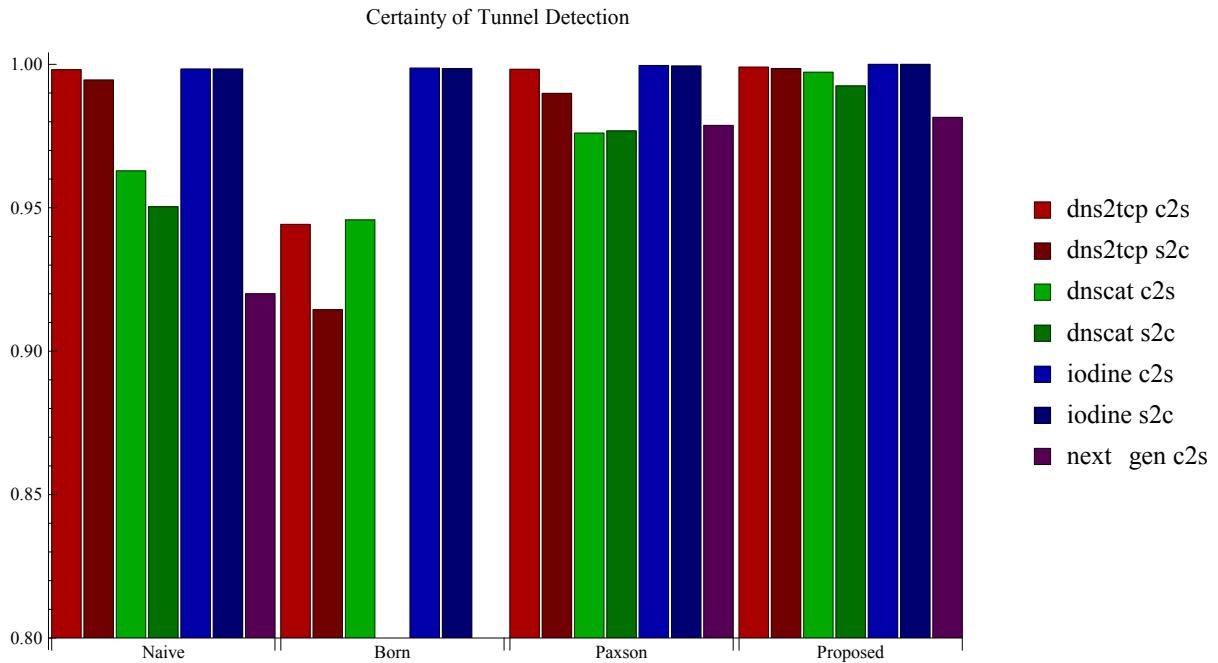


Figure 7.15: Figure 7.14 replotted with a reduced range to accentuate small differences.

performance of the proposed method and Paxson’s method relative to the naive method is clearly visible in this chart.

The differences between Paxson’s method and the proposed method are visible, but an additional chart further accentuating them is instructive and is shown in figure 7.16. In this final chart which shows a range of certainties in the interval [0.95, 1.00], the differences between the two methods are clearly visible, with the proposed approach achieving a higher certainty in every detection scenario.

7.3 Tunnel Detection Performance Conclusion

It is visible from the detailed plots in sections 7.1 and 7.2 that the proposed method is superior to its peers in its ability to detect tunnels with certainty in excess of ninety eight percent. This extremely high detection rate is achieved at a very short time scale and with very low tunnel throughput. Depending on the tunnelling application, detection

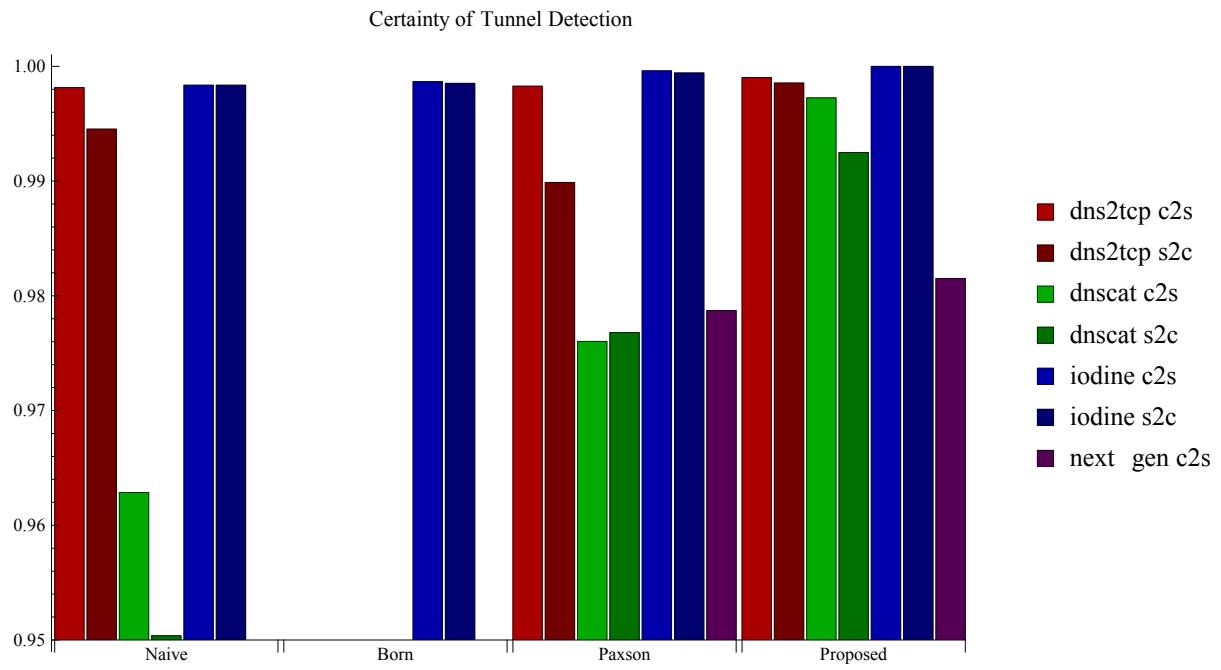


Figure 7.16: Figure 7.14 replotted with a further reduced range to accentuate *very* small differences.

performance rises above ninety nine percent for all except the next-gen application, and above 99.999% for Iodine.

Additional quantified details are given in the final conclusion (chapter 8)..

Chapter 8

Conclusion

In this work, a new method of detecting DNS tunnels was proposed, described, and evaluated.

The existing landscape of detection methods was summarized (section 3) and a gap identified indicating a need for a new method (section 4.3) that has both high processing performance on commodity hardware, and robust tunnel detection. A prototypical next-gen tunnel application was postulated (section 4.2) and simulated (details in appendix B) in order to present a more difficult detection task during evaluation alongside existing tunnel applications (section 6.1.2). Several detection methods from the literature as well as the proposed method were selected (section 6) and implemented on a common framework (section 6.2). The implemented methods were tested for processing performance (section 6.2.7) and tunnel detection performance on a large sample of real world DNS traffic as well as existing and next-gen tunnelling application data (section 7).

As is shown in the relevant sections, the proposed method outperforms its peers in both processing performance and tunnel detection in almost every situation by a measurable and often considerable margin.

When compared to its peers from literature, the performance improvement over Pax-

son’s approach on real-world data is approximately 25% while a speedup of almost 100% is observed when processing tunnel data under the Cython interpreter. The margins are similar under PyPy, with the primary difference being that the proposed method is nearly double the performance of Paxson’s approach on both real world and tunnel application data. Born’s approach is able to narrow the gap on tunnel data to only approximately a 50% improvement, while the gap on real-world data increases to about 30%. Under PyPy however, the performance of the proposed method and Born’s method become very similar with gaps less than 10%.

Examining tunnel detection performance shows additional benefits to the proposed method. Born’s method shows very poor performance, as expected, when faced with the next-gen tunnel, with a less drastic performance hit shown on other tunnelling software. In all cases the proposed, as well as Paxson’s, approach out-perform Born’s method by a considerable margin in false-positive rates. The proposed approach reduces the false positive rate by almost 98% when compared to Born’s approach. When comparing Paxson’s approach with the proposed approach, the proposed approach reduces the number of false positives by up to nearly 90% in the best case, with an average reduction of 70%.

The exception is the naive method which outperforms the proposed method in processing performance by approximately fifty percent at the expense of detection performance, with the proposed approach reducing the false positive rate by 82%. As was mentioned in section 5.1.3 however, the naive method’s detection performance is only as good as it is in this case due to the lack of duplication of DNS queries.

Recall that the improvements demonstrated by the proposed approach are on measured in a highly pessimistic detection scenario, with false-positive rates rapidly dropping below 10^{-6} as throughput increases.

The end product of these results is a contribution to the field comprising a new detection method with superior processing and detection performance. The new detection method

falls short of matching the naive method’s processing performance in many situations with the trade off of far superior detection performance in the general case.

8.1 Future Work

The proposed method was shown to be highly effective in detecting the targeted network traffic with superior performance, outperforming existing methods in almost every scenario. In order for this to be of value, however, an adoption of this method into an existing commercial products or an implementation as a plugin for an existing security framework would be necessary.

Future work could include implementing the proposed method for Bro, Snort, Suricata or other existing intrusion detection systems in order to improve the ability of organizations to observe DNS tunnels in their network. Partnership with, and adoption by, an existing industry partner would aid in the spread and deployment of this technique in enterprise and corporate environments. Due to the nature of computer security research, regardless of how valuable a contribution may be, unless there is deployment of the approach into real world environments it is unable to aid in securing the Internet.

Appendices

Appendix A

Properties of the Network Capture

The packet capture used to represent real world data was performed in cooperation with Merlin, an ISP in Manitoba for educational institutions.

File name:	data1.udp53.pcap
File type:	Wireshark/tcpdump/... - libpcap
File encapsulation:	Ethernet
Packet size limit:	file hdr: 65535 bytes
Number of packets:	1035425650
File size:	156484681395 bytes
Data size:	139917870971 bytes
Capture duration:	1915055 seconds
Start time:	Thu Nov 4 15:05:58 2010
End time:	Fri Nov 26 18:03:33 2010
Data byte rate:	73062.06 bytes/sec
Data bit rate:	584496.51 bits/sec
Average packet size:	135.13 bytes
Average packet rate:	540.68 packets/sec
SHA1:	6fb043f37e659d9a65a45b5c8074d55eae386600
RIPEMD160:	4aba72d43f5726f0a8576f35d60845f895080a43
MD5:	32956ef8efe5c3873226386e1608561a
Strict time order:	True

Table A.1: Statistics and information as reported by Wireshark's utility *capinfos* on the real world packet capture..

While the capture file contains over one billion total packets, only 339,321,911 (or 32.7%)




Figure A.1: Throughput in bytes per second of the packet capture, with each sample being the average throughput over ten-second windows.

of the packets represent valid DNS traffic.




Figure A.2: Throughput in packets per second of the packet capture, with each sample being the average throughput over ten-second windows.

Appendix B

Probabilistic Encoding

Several detection methods in the literature involve examining the character distribution in the queries, or responses, of DNS packets. These approaches compare the distributions obtained from the traffic being analyzed, and compare them to a known distribution that can be considered normal. If there is a sufficiently significant measurable difference, then the packet is flagged as anomalous.

If it were possible for a DNS tunnel to encode its output in order to match the distribution that these detection methods consider normal, then it would be possible for it to evade detection by masquerading as benign DNS traffic. These detection methods assume that the output of a DNS tunnel will have high entropy (due to compression and/or encryption) and/or span a wide range of character values, whereas normal traffic does not have these features. The solution to this problem becomes one of converting a high entropy source into a lower entropy encoding in a way that the high entropy stream can be recovered from the low entropy encoding, where the low entropy encoding has a specific distribution of characters.

A proof-of-concept tool was written in C that performs precisely this task. The tool

takes, as input, what it assumes to be a high entropy source¹ and a configuration file that describes the desired output distribution.

B.1 Sample Encoding - English Unigrams

A sample encoding is given using the english character frequency distribution with the character distribution table used given in table B.1. In order to ensure that interested parties can verify this output, the source high-entropy data is contained in table B.2. A data matrix image of the same base64 encoding is also given in figure B.1. This can be converted back to the original source via the standard Unix command line tool *base64*. The output of the encoder using english character frequencies and the given input data is shown in table B.3. It is important to observe that the output produced does not resemble typical English text and can be easily detected by any software that performs analysis on higher level characteristics. Such distinguishing characters include, but are note limited to, digrams, trigrams, and the existence of English words.

Analyzing the frequencies of the output shows that it closely approximates english character frequencies. A larger sample of random data (one hundred thousand bytes) was also encoded, and its output analyzed as well. The small sample, the large sample, and actual english character frequencies are displayed together in figure B.2. As is evident from the combined plots of the various distributions, there does not exist a clear distinction between them and thus it becomes very difficult to perform automated classification when information is encoded using this tool.

¹Since any source can be used to produce a high entropy source through encryption or compression, this is considered a safe assumption and limitation.

Character	Count	Character	Count
<space>	0.11965	u	0.0242803
e	0.111823	m	0.0211814
t	0.0797253	w	0.0207765
a	0.0718989	f	0.0196144
o	0.0660886	g	0.0177392
i	0.0613258	y	0.0173783
n	0.0594154	p	0.0169821
s	0.0557003	b	0.013135
h	0.0536491	v	0.00860991
r	0.0527071	k	0.00679637
d	0.0374417	j	0.00134695
l	0.0354345	x	0.00132054
c	0.0244916	q	0.000836341
		z	0.000651466

Table B.1: Probability distribution used for english character frequency in the sample encoding and decoding.

HEzQcP9uxPOzeOB1SfRP+DQ7x3X5dTA1WF3vpcn05kgLQQEP7xnUrnUy7U8ISNbNQd+8da64+Ci
nNBRvu3TR1GOKJvLzb6QDBWqxVfa4VeAU+cSvj+uzoajp2F0jd5/csHxJQ1guoeqT76pq8oStoLG
3k48PZWuebgweZx5KdxAgUfN7/kvzwBzG70+9B7J/08y6BfGz6In0H+LuDCOsFqM1j1jX3SkgJq/
yPka52SDxa3D4GXecj9d

Table B.2: The base64 encoding of the binary data used for the sample encoding.

" mfeo hnhaeeeewsaos eewk es fsemroh eeeewthesdre s sol sreese ea hrzltse
wesrnyemnpaoonef ht hsheawnfdemue stntleewth home esaiu s feesgealne seaoa
hlhyp hmb ewaaeesdubotongtnploaanewefochd htugeside s oeesi fptsmr
hweeeessoltsnahao esde hgmuo weeeeewyl haoesroo frmehstieest ewad svikefui
hncestphstestidau eeahl hrosoan wrthaa ft heabpeeses eeeeeewrsn hrlucne ewr
esesooe eeeeeewsqe eewefgeesng est f eeeeewlhiae hgao hiuggnrrancewgmob
hdreeesgopesuoamyufiolisx esrrsloaeewtse"

Table B.3: The encoded form of the binary data according to English character frequencies.

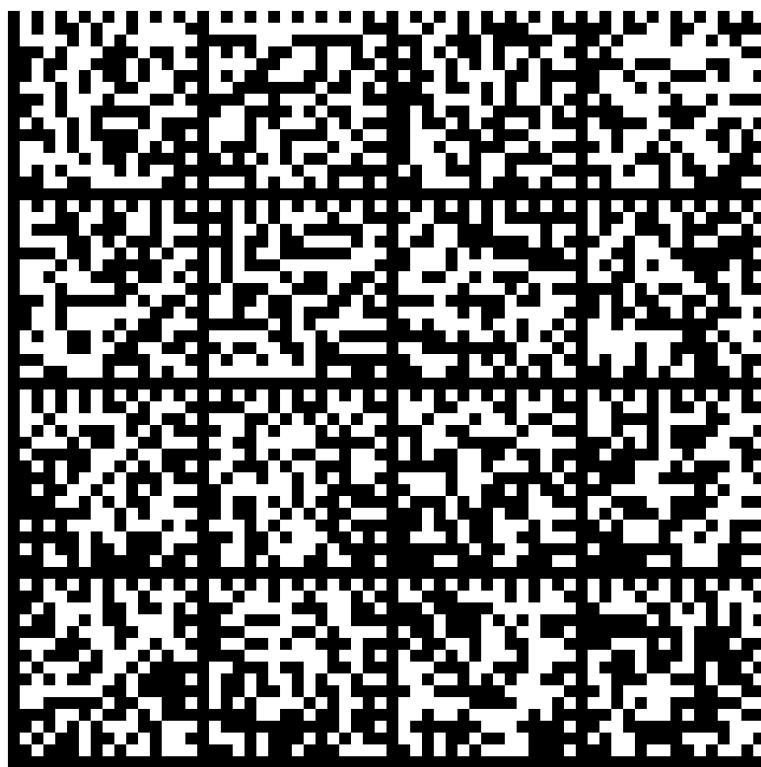


Figure B.1: The base64 encoding of the binary data used in the sample encoding, represented in a data matrix image.

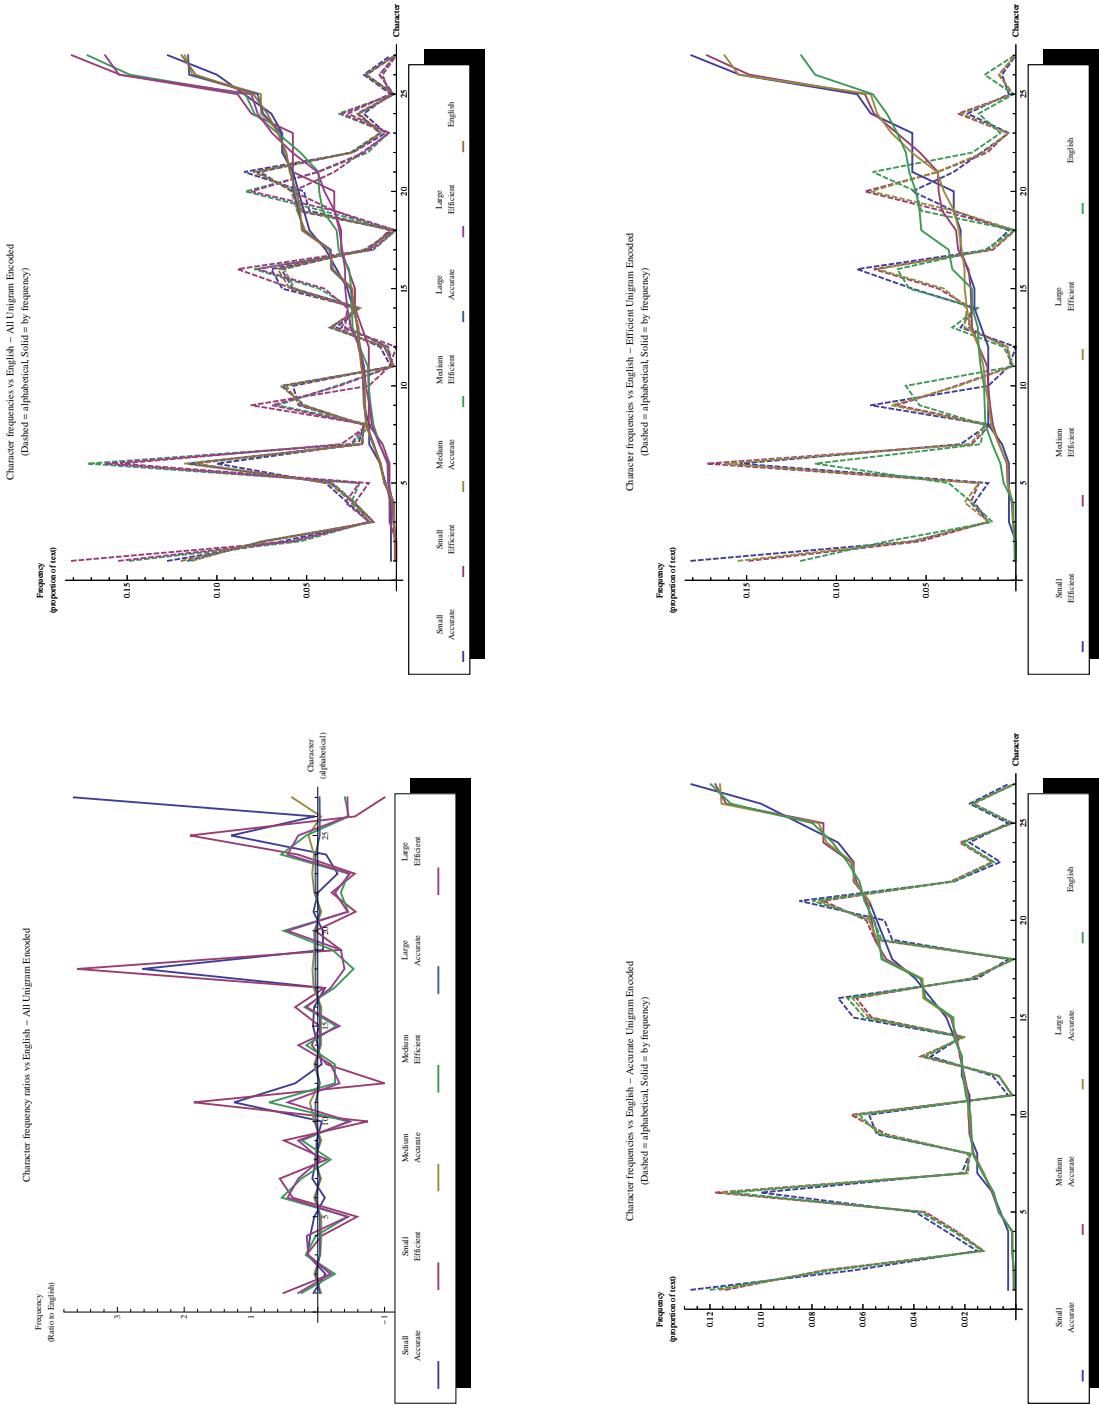


Figure B.2: The character frequencies of the small and large encoded samples using both accurate and efficient encoding, as well as the English target distribution sorted by character.

B.1.1 Technical Details

The probabilistic encoder is open-source and freely available, licensed under the Mozilla Public License version 2 which is GPL compatible. The source code is available at [28] for download and use under the MPL2 license terms. The source code compiles on Unix and other POSIX based systems such as Cygwin on Windows as well as non-POSIX platforms such as Visual Studio on Windows.

The encoding and decoding processes can be broken up into two steps; probability distribution analysis followed by the encoding or decoding. The analysis of the distribution is the same for both the encoding and decoding operations, however the encoding operation is somewhat more involved than the decoding operation from an algorithmic point of view.

The probability distribution is analyzed in order to map each element of the distribution (character, word, etc...) to a string of bits. An assumption is made that all bit strings of a given length are equally likely to appear in the high-entropy source, and this places a constraint in the output PDF; the most frequent symbol cannot be more frequent than the most frequent symbol in the high-entropy source. That is, no symbol in the output distribution can have a frequency above 50%. This is considered a fair assumption since in English, the most common letter has a frequency of approximately 11%.

The mapping from symbol to bit string is done by sorting the distribution from most to least frequent, and then selecting a shortest available bit string of a length that has an expected occurrence rate not less than the symbol's occurrence rate in the output distribution. For example, the bit string to symbol mapping for the English unigram encoding sample is given in Table ?? along with the approximate frequency of the letters in English text.

Observe that it is possible to make a tradeoff between encoding accuracy (that is, how accurately the output matches the desired distribution) and encoding efficiency (that is, the ratio between output and input size with higher ratios indicating a less efficient encoding)

by tuning the bit string mapping. More accurate mappings result from choosing shorter bit strings (table B.4 chooses the shortest possible bit string mappings and achieves an efficiency factor of approximately 3.3) with more efficient mappings sacrificing accuracy for efficiency. Figure B.2 shows the frequencies of the English target distribution along with the most efficient encoding and the most accurate encoding. The efficient coding (shown in table B.5) in that figure achieves a factor of 2.4 compared to the accurate encoding's factor of 3.3. The efficiency factors represent the ratio of bytes in the output compared to bytes in the input.

Once the mapping, and various lookup tables built, have been built the encoding or decoding process can begin. For encoding, a table is maintained that keeps track of how many times each symbol in the desired distribution has been output. The purpose of this table is to facilitate the greedy selection of symbols in order to ensure that one symbol is not chosen more often than is necessary and that those symbols that are farthest from meeting their quota in the output stream get preferential selection when the input bits match their mapping. Decoding is a much more straight forward process that involves reading input characters and outputting the appropriate bit string based on the mapping generated from the distribution file.

The examples given above show the unigram mapping process, however since the theory and implementation are symbol agnostic, the symbol set can be as large as is desired, with each symbol being arbitrarily long. One could potentially use this tool to build a mapping on all words in a language paired with the nominal frequencies if one so desired, however the Markov chain coding portion of the tool may be a better choice.

Character	Bit-String	Count
<space>	1	0.11965
e	0	0.111823
t	11	0.0797253
a	01	0.0718989
o	10	0.0660886
i	00	0.0613258
n	111	0.0594154
s	011	0.0557003
h	101	0.0536491
r	001	0.0527071
d	110	0.0374417
l	010	0.0354345
c	100	0.0244916
u	000	0.0242803
m	0011	0.0211814
w	1101	0.0207765
f	0101	0.0196144
g	1001	0.0177392
y	0001	0.0173783
p	1110	0.0169821
b	0110	0.013135
v	1010	0.00860991
k	0010	0.00679637
j	1100	0.00134695
x	0100	0.00132054
q	1000	0.000836341
z	0000	0.000651466

Table B.4: Accurate Bit-string mapping for English characters and their frequency for comparison.

Character	Bit-String	Count
<space>	0	0.11965
e	1	0.111823
t	010	0.0797253
a	110	0.0718989
o	001	0.0660886
i	1010	0.0613258
n	0110	0.0594154
s	1110	0.0557003
h	0001	0.0536491
r	1001	0.0527071
d	0101	0.0374417
l	1101	0.0354345
c	00110	0.0244916
u	10110	0.0242803
m	01110	0.0211814
w	11110	0.0207765
f	00001	0.0196144
g	10001	0.0177392
y	01001	0.0173783
p	11001	0.0169821
b	001010	0.013135
v	101010	0.00860991
k	0110100	0.00679637
j	111010000	0.00134695
x	000110000	0.00132054
q	1001100000	0.000836341
z	0101100000	0.000651466

Table B.5: Efficient Bit-string mapping for English characters and their frequency for comparison.

B.2 Custom DNS Tunnel Endpoint Simulator

The next-gen tunnel that was used alongside the existing implementations was built by using standard Unix utilities and the application described in the previous section. The utility was only capable of unidirectional transfer, from server to client, and relied on three standard Unix utilities: *gzip*, *fold* and *nslookup*. The basic flow was as follows:

1. The data to be transferred was piped through *gzip* in order to compress a potentially low entropy stream, reducing its size and guaranteeing a high entropy stream as input to the next stage of the pipeline.
2. The high entropy stream is used as input to the tool described in the previous section which was targeting a character distribution that matched the Alexa top one million domains, as shown in table 4.1. The output of this step is a stream of characters that asymptotically converges to the given domain.
3. Because DNS queries have a maximum length of a token (a string of characters in between periods) at sixty three characters, the Unix utility *fold* is used to insert newline characters in the stream no less often than every sixty-fourth character. Depending on the targeted input rate and desired output parameters, folding could produce queries of any length from one to sixty three characters. If one was looking to target a query length distribution, the use of *fold* could be replaced by another tool that used a more intelligent method of deciding when to insert newline characters. The output of this stage is a stream of lines of equal length, each of which will represent a query that is sent to the server.
4. The Unix utility *nslookup* allows for command-line issuing of DNS queries to a server. The output of fold was piped into *nslookup* with the queries directed at the localhost address. This resulted in a relatively high throughput as *nslookup* read a line from its

input, issued the query to the localhost, was met with a negative response (specifically, an NXDOMAIN response indicating that the address was not something that the localhost knew about), at which point it would grab a new line of input and repeat.

All traffic used for the analysis of the next-gen tunnelling application involved this flow of information. By combining existing applications with the application in the previous section, it was possible to very quickly, and simply, prototype a unidirectional DNS tunnelling application that runs on any linux, Unix, Mac or Windows based host. Windows hosts using Cygwin are able to use common Unix utilities such as *fold* and *gzip*. Windows has its own version of *nslookup* bundled with it that has sufficiently similar behaviour to the Unix style version as to support this use case.

Bibliography

- [1] *Information technology – Digital compression and coding of continuous-tone still images: Requirements and guidelines*, 1994, ISO/IEC 10918-1:1994.
- [2] Wolfram Alpha, *botnet - Wolfram—Alpha*, <http://www.wolframalpha.com/input/?i=botnet>, aug 2013.
- [3] _____, *entropy - Wolfram—Alpha*, http://www.wolframalpha.com/input/?i=entropy&a=*C.entropy-*Word-, aug 2013.
- [4] _____, *entropy - Wolfram—Alpha*, http://www.wolframalpha.com/input/?i=entropy&a=*C.entropy-*MathWorld-, aug 2013.
- [5] R. Arends, R. Austein, M. Larson, D. Massey, and S. Rose, *Resource Records for the DNS Security Extensions*, <http://www.ietf.org/rfc/rfc4034.txt>, mar 2005, RFC 4034 (Proposed Standard).
- [6] Matthias Bauer, *New covert channels in HTTP: adding unwitting Web browsers to anonymity sets.*, WPES (Sushil Jajodia, Pierangela Samarati, and Paul F. Syverson, eds.), ACM, 2003, pp. 72–78.
- [7] Kevin Borders and Atul Prakash, *Web tap: detecting covert web traffic.*, ACM Conference on Computer and Communications Security (Vijayalakshmi Atluri, Birgit Pfitzmann, and Patrick Drew McDaniel, eds.), ACM, 2004, pp. 110–120.

- [8] Kenton Born, *Browser-Based Covert Data Exfiltration*, CoRR **abs/1004.4357** (2010).
- [9] _____, *PSUDP* — Kenton Born, <http://www.kentonborn.com/psudp>, jul 2010, PSUDP source code and implementation.
- [10] _____, *PSUDP: A Passive Approach to Network-Wide Covert Communication*, (2010).
- [11] Kenton Born and David Gustafson, *Detecting DNS Tunnels Using Character Frequency Analysis*, CoRR **abs/1004.4358** (2010).
- [12] _____, *NgViz: Detecting DNS Tunnels through N-Gram Visualization and Quantitative Analysis*, CoRR **abs/1004.4359** (2010).
- [13] Randy Browne, *An Entropy Conservation Law for Testing the Completeness of Covert Channel Analysis.*, ACM Conference on Computer and Communications Security (Dorothy E. Denning, Raymond Pyle, Ravi Ganesan, and Ravi S. Sandhu, eds.), ACM, 1994, pp. 270–281.
- [14] Patrick Butler, Kui Xu, and Danfeng (Daphne) Yao, *Quantitatively Analyzing Stealthy Communication Channels.*, ACNS (Javier Lopez and Gene Tsudik, eds.), Lecture Notes in Computer Science, vol. 6715, 2011, pp. 238–254.
- [15] Serdar Cabuk, Carla E. Brodley, and Clay Shields, *IP covert timing channels: design and detection.*, ACM Conference on Computer and Communications Security (Vijayalakshmi Atluri, Birgit Pfitzmann, and Patrick Drew McDaniel, eds.), ACM, 2004, pp. 178–187.
- [16] _____, *IP Covert Channel Detection.*, ACM Trans. Inf. Syst. Secur. **12** (2009), no. 4.

- [17] Michel Chamberland, *Snort rules for Iodine Covert DNS Tunnel Detection*, <http://www.securitywire.com/2009/07/snort-rules-for-iodine-covert-dns-tunnel-detection>, jul 2009.
- [18] Hervé Schauer Consultants, *HSC - Tools - Dns2tcp*, <http://hsc.fr/ressources/outils/dns2tcp/index.html.en>, may 2012.
- [19] L. Peter Deutsch, *kryo.se: iodine (IP-over-DNS, IPv4 over DNS tunnel)*, <http://code.kryo.se/iodine>, feb 2010.
- [20] Christian J. Dietrich, Christian Rossow, Felix C. Freiling, Herbert Bos, Maarten van Steen, and Norbert Pohlmann, *On Botnets That Use DNS for Command and Control*, Proceedings of the 2011 Seventh European Conference on Computer Network Defense (Washington, DC, USA), EC2ND '11, IEEE Computer Society, 2011, pp. 9–16.
- [21] Maximillian Dornseif, *mdornseif/DeNiSe GitHub*, <https://github.com/mdornseif/DeNiSe>, jan 2006.
- [22] Maurizio Dusi, Manuel Crotti, Francesco Gringoli, and Luca Salgarelli, *Tunnel Hunter: Detecting application-layer tunnels with statistical fingerprinting.*, Computer Networks **53** (2009), no. 1, 81–97.
- [23] D. Eastlake 3rd, *DNS Request and Transaction Signatures (SIG(0)s)*, <http://www.ietf.org/rfc/rfc2931.txt>, sep 2000, RFC 2931 (Proposed Standard).
- [24] Greg Farnham, *Detecting DNS Tunneling*, InfoSec Reading Room (2013).
- [25] Travis Friesen, *An in-memory database for prototyping anomaly detection algorithms at gigabit speeds*, 2013.
- [26] Steven Gianvecchio and Haining Wang, *Detecting covert timing channels: an entropy-based approach.*, ACM Conference on Computer and Communications Security (Peng

Ning, Sabrina De Capitani di Vimercati, and Paul F. Syverson, eds.), ACM, 2007, pp. 307–316.

- [27] Jeffrey J. Guy, *dns part ii: visualization*, <http://armatum.com/blog/2009/dns-part-ii>, feb 2009.
- [28] Michael Himbeault, *Probabilistic Encoder: Summary*, <https://www.riebart.ca/hg/probcode>, jun 2013.
- [29] Maarten Van Horenbeeck, *Deception on the network: thinking differently about covert channels*, 2006.
- [30] <http://www.lookuptables.com>, *Ascii Table - ASCII character codes and html, octal, hex and decimal chart conversion*, <http://www.asciiitable.com/>, oct 2013.
- [31] Internet Storm Centre - SANS Institute, *Hash Database — SANS Internet Storm Center; Cooperative Network Security Community - Internet Security*, https://isc.sans.edu/tools/hashsearch.html#dns_interface, oct 2013.
- [32] jhind, *Catching DNS tunnels with A.I.*, <http://www.meanypants.com/meanypants>, jul 2009.
- [33] Anestis Karasaridis, Kathleen S. Meier-Hellstern, and David A. Hoeflin, *Detection of DNS Anomalies using Flow Data Analysis.*, GLOBECOM, IEEE, 2006.
- [34] D. Massey and S. Rose, *Limiting the Scope of the KEY Resource Record (RR)*, <http://www.ietf.org/rfc/rfc3445.txt>, dec 2002, RFC 3445 (Proposed Standard).
- [35] P.V. Mockapetris, *Domain names - concepts and facilities*, RFC 1034 (INTERNET STANDARD), November 1987, Updated by RFCs 1101, 1183, 1348, 1876, 1982, 2065, 2181, 2308, 2535, 4033, 4034, 4035, 4343, 4035, 4592, 5936.

- [36] P.V. Mockapetris, *Domain names - implementation and specification*, <http://www.ietf.org/rfc/rfc1035.txt>, nov 1987, RFC 1035 (INTERNET STANDARD).
- [37] Ira S. Moskowitz, Richard E. Newman, Daniel P. Crepeau, and Allen R. Miller, *Covert channels and anonymizing networks.*, WPES (Sushil Jajodia, Pierangela Samarati, and Paul F. Syverson, eds.), ACM, 2003, pp. 79–88.
- [38] Robert C. Newman, *Covert computer and network communications*, Proceedings of the 4th annual conference on Information security curriculum development (New York, NY, USA), InfoSecCD '07, ACM, 2007, pp. 12:1–12:8.
- [39] Keisuke Okamura and Yoshihiro Oyama, *Load-based covert channels between Xen virtual machines.*, SAC (Sung Y. Shin, Sascha Ossowski, Michael Schumacher, Mathew J. Palakal, and Chih-Cheng Hung, eds.), ACM, 2010, pp. 173–180.
- [40] Vern Paxson, *Behavioral Detection of Stealthy Intruders*, <https://seclab.cs.ucsb.edu/academic/projects/projects/cybaware/2011>, sep 2011.
- [41] Tadeusz Pietraszek, *DNScat*, <http://tadek.pietraszek.org/projects/DNScat>, sep 2005.
- [42] Julius Plenz, *DNSTunnel.de - free DNS tunneling service*, <http://dnstunnel.de>, jun 2011.
- [43] Proventia, *Proventia Server IPS - DNS tunnel traffic detected*, http://www.iss.net/security_center/reference/vuln/DNS_Tunnel_Detected.htm, jan 2013.
- [44] Baishakhi Ray and Shivakant Mishra, *A Protocol for Building Secure and Reliable Covert Channel*, Proceedings of the 2008 Sixth Annual Conference on Privacy, Security and Trust (Washington, DC, USA), PST '08, IEEE Computer Society, 2008, pp. 246–253.

- [45] M. Richardson, *A Method for Storing IPsec Keying Material in DNS*, <http://www.ietf.org/rfc/rfc4025.txt>, mar 2005, RFC 4025 (Proposed Standard).
- [46] D. Romaña and Y. Musashi, *Entropy Based Analysis of DNS Query Traffic in the Campus Network*, (2007).
- [47] Stephan Roolvink, *Detecting attacks involving DNS servers*, (2008).
- [48] Sarah H Sellke, Chih-Chun Wang, Saurabh Bagchi, and Ness Shroff, *Tcp/ip timing channels: Theory to implementation*, INFOCOM 2009, IEEE, IEEE, 2009, pp. 2204–2212.
- [49] Alan B. Shaffer, Mikhail Auguston, Cynthia E. Irvine, and Timothy E. Levin, *A security domain model to assess software for exploitable covert channels.*, PLAS (Åslfar Erlingsson and Marco Pistoia, eds.), ACM, 2008, pp. 45–56.
- [50] CPython Team, *TimeComplexity - Python Wiki*, <https://wiki.python.org/moin/TimeComplexity>, oct 2013.
- [51] PyPy Team, *PyPy Status Blog: PyPy is faster than C, again: string formatting*, <http://morepypy.blogspot.com/2011/08/pypy-is-faster-than-c-again-string.html>, oct 2013.
- [52] Brennon Thomas, Barry E. Mullins, Gilbert L. Peterson, and Robert F. Mills, *An FPGA System for Detecting Malicious DNS Network Traffic.*, IFIP Int. Conf. Digital Forensics (Gilbert L. Peterson and Sujeet Shenoi, eds.), IFIP Advances in Information and Communication Technology, vol. 361, Springer, 2011, pp. 195–207.
- [53] S. Weiler, *Legacy Resolver Compatibility for Delegation Signer (DS)*, <http://www.ietf.org/rfc/rfc3755.txt>, may 2004, RFC 3755 (Proposed Standard).